

Abstract

Credit risk evaluation is a serious challenge for the modern finance with loan defaults posing significant risk to not just the profitability of single institution but also stability of a whole economy. This study investigates the application of machine learning models for loan default prediction, focusing on both model performance and the treatment of class imbalance, which is a common issue for analysing credit datasets. Six models, Logistic Regression, Decision Tree, Random Forest, Support Vector Classifier, XGBoost, and LightGBM, were evaluated under four imbalance-handling techniques, Random Undersampling, Synthetic Minority Oversampling Technique (SMOTE), Threshold Optimization, and Cost-Sensitive Learning keeping original imbalanced data as a baseline for comparison. Performance was measured using a range of statistical and financial metrics including precision, recall, F1-score, and ROC-AUC, features were selected using Recursive Feature Elimination with Cross Validation (RFECV) and models were interpreted using SHapley Additive exPlanations (SHAP).

Results stated that gradient boosting ensemble models such as XGBoost and LightGBM consistently outperformed baseline models, especially when combined with Threshold Optimization and Cost Sensitive Analysis, they demonstrate notable improvements in balancing sensitivity(recall) and precision, aligning with the findings from recent literature. Additionally, the study highlights trade-offs between model perfection and interpretability, highlighting the importance of adopting methods that ensure regulatory compliance and practical deployment in financial institutions. The research contributes to existing literature by comparing most used industry wide models coupled with class imbalance handling strategies, prioritizing not only predictive ability but also computational efficiency, interpretability, and financial relevance.

Acknowledgement

I would like to express my deepest gratitude to my supervisor, Ms Anh Luong, for her invaluable guidance, encouragement, and constructive feedback throughout the course of this research. Her insights into credit risk modelling and machine learning greatly shaped the direction and quality of this dissertation.

I am also grateful to the faculty and staff at MSc of Business Analytics, Warwick Business School, for providing a supportive academic environment and access to the necessary resources for completing this work.

Special thanks are extended to my family and friends, whose constant encouragement, patience, and belief in my abilities helped me stay motivated during this journey. I would also like to acknowledge my peers, who provided helpful discussions, feedback, and moral support.

Finally, I am indebted to the developers and contributors of open-source libraries and models such as numpy, pandas, columntransfer, RFECV, matplotlib, seaborn, scikit-learn, Logistic Regression, Random Forest, Decision Trees, XGBoost, LightGBM, and SHAP, without which this research would not have been possible.

Table of Contents

Chapter 1	7
Introduction	7
1.1 Motivation, Objective and Aim:	7
1.2 Purpose:	9
1.3 Section Overview:	9
Chapter 2	11
Literature Review	11
2.1 Historical Background of lending:	11
2.2 How bad lending affected economy:	12
2.3 From Paper Trails to Algorithms: The Evolution of Loan Default Prediction Research	12
2.3.1 Traditional Statistical Models (Pre-2000 era)	12
2.3.2 Rise of Machine Learning (2000–2010)	13
2.3.3 Big Data and Ensemble Models (2010–2020)	13
2.3.4 From the Age of Explainable AI to Ethical Fairness (2020 and counting)	13
2.4 Industry Adoption and Human Oversight:	13
2.5 Barriers to Precision: Challenges and Limitations:	14
2.6 Contribution of this study to literature:	15
Chapter 3	17
Methodology	17
3.1 Data Source and Sampling:	17
3.2 Data Preprocessing:	18
3.3 Feature Engineering:	18
3.4 Feature Selection:	19
3.5 Class Imbalance Handling Techniques:	20
3.6 Models Selection:	21
3.7 Evaluation Metrics:	23
3.8 Model Interpretation:	25
Chapter 4	26
Results	26
4.1 F1 Score:	27
4.2 ROC-AUC Curve:	28
4.3 Precision, Recall and Confusion Matrix:	28
4.4 Feature Importance for model interpretability:	29
Chapter 5	31
Discussion	31

5.1 Findings for Research Questions:.....	31
5.2 Limitations:	32
5.3 Future Recommendations:.....	33
Chapter 6	34
Conclusion	34
Reference List:	35
6 Appendix	40
Appendix 1: Lending Club Dataset:	40
1.1 Data Dictionary:.....	40
1.2 Original Data Summary Statistics:	44
Appendix 2: Top 10 features for models using RFECV:	45
Appendix 3: Classification Reports	47
3.1 Models on Original Imbalanced dataset:	47
3.2 Models under Random Undersampling:	49
3.3 Models under SMOTE:.....	51
3.4 Models under Threshold Optimization:	53
3.5 Models under Cost Sensitive Analysis:.....	55
Appendix 4: ROC-AUC Curves:	57
Appendix 5: Shap figures:.....	58
Appendix 6: Script of Python Code:.....	60

Table of Figures

FIGURE 1 PROCESS FLOW CHART FOR METHODOLOGY	17
FIGURE 2 ROC-AUC SCORES BY TECHNIQUES AND MODEL	28
FIGURE 3 FEATURE IMPORTANCE OF LIGHTGBM WITH THRESHOLD OPTIMIZATION	29
FIGURE 4 FEATURE IMPORTANCE OF XGBOOST WITH THRESHOLD OPTIMIZATION.....	30
FIGURE 5 AREA UNDER CURVES FOR EACH CLASS IMBALANCE HANDLING TECHNIQUE	57
FIGURE 6 FEATURE IMPORTANCE FOR XGBOOST ON SMOTE	58
FIGURE 7 FEATURE IMPORTANCE FOR COST SENSITIVE XGBOOST	58
FIGURE 8 FEATURE IMPORTANCE FOR LIGHTGBM ON SMOTE	59
FIGURE 9 FEATURE IMPORTANCE FOR COST SENSITIVE LIGHTGBM	59

Table of Tables

TABLE 1 BRIEF MODEL DEFINITIONS AND THEIR EQUATIONS	23
TABLE 2 TP, TN, FP AND FN FROM CONFUSION MATRIX.....	23
TABLE 3 EVALUATION METRICS DEFINITION AND FORMULAS.....	24
TABLE 4 COMPARATIVE ANALYSIS RESULT TABLE SORTED BASED ON ROC-AUC VALUES	26
TABLE 5 OPTIMAL THRESHOLD VALUES FOR EVERY MODEL.....	27
TABLE 6 CONFUSION MATRIX IN CONTEXT OF LOAN DEFAULT PREDICTION.....	29
TABLE 7 DESCRIPTIVE SUMMARY STATISTICS FOR ORIGINAL IMBALANCED DATA.....	44
TABLE 8 TOP 10 FEATURES OF LR.....	45
TABLE 9 TOP 10 FEATURES OF DT	45
TABLE 10 TOP 10 FEATURES OF RF	45
TABLE 11 TOP 10 FEATURES OF SVC.....	46
TABLE 12 TOP 10 FEATURES OF XGBOOST	46
TABLE 13 TOP 10 FEATURES OF LIGHTGBM.....	46

Chapter 1

Introduction

In today's day and age, loan default prediction plays a pivotal role for various stakeholders in financial ecosystem ranging from Banks, Insurance corporations to fintech companies including Government to evaluate borrower's ability to meet debt obligations, therefore it is an absolute necessity to do a thorough and firm credit scoring for applicants. Credit risk scoring used to evaluate borrower's ability to pay back the loan amount using loan officer's intuition and traditional score methods such as Logistic Regression, however these methods perform on a suboptimal level when it comes to dealing with modern complex high dimensional data with non-linear relationships (Hand & Henley, 1997).

With development of technology, time consuming traditional methods are becoming obsolete, instead, robust and automatic machine learning models are being deployed to predict probability of default across almost all the financial institutions. These advanced machine learning models including Decision Trees, Random Forest, Support Vector Machines and Gradient Boosting, generalize better for unseen data (Lessmann et al., 2015). By utilising large datasets and sophisticated algorithms, ML models have demonstrated superior predictive power compared to traditional methods in a wide range of financial applications, including credit scoring, fraud detection, and loan default prediction (Galindo & Tamayo, 2000; Baesens et al., 2003).

1.1 Motivation, Objective and Aim:

The loan default prediction has emerged as a key area of research for financial institutions due to its direct impacts on profitability, risk management, and regulatory compliance. In recent years, the global rise in consumer credit availability, linked with economic volatility, has increased the weight of powerful default prediction models. Typical credit scoring systems, such as those based on logistic regression, often struggle to capture complex nonlinear relationships within high-dimensional financial datasets (Lessmann et al., 2015; Abellán & Mantas, 2014). Meanwhile, the Basel III, a set of international banking regulations developed by the Basel Committee on Banking Supervision (BCBS), 2019, which is based at the Bank for International Settlements (BIS) in Basel, Switzerland, emphasizes not only predictive power but also transparency in risk assessment, requiring banks to justify their models' decisions to regulators and stakeholders (Kou et al., 2021). It was introduced after the 2008

global financial crisis to strengthen the regulation, supervision, and risk management of banks worldwide.

Recent advances in machine learning, including ensemble models like XGBoost and LightGBM, have demonstrated superior predictive performance over customary approaches in credit risk modelling (Xia et al., 2022; Louzada et al., 2016). However, the trade-off between accuracy and interpretability remains a key challenge, essentially for high-stakes financial decision-making.

Furthermore, according to Zhou et al., 2021, standard metrics like accuracy can be misleading in such contexts since default prediction involves unique challenges including class imbalance, where the target variable in a classification problem has a disproportionate distribution of classes denoting utilisation of class imbalance handling techniques. Financial institutions are remarkably sensitive to false negatives (i.e., approving high-risk borrowers who will likely default but predicted safe), which can lead to significant losses (Lessmann et al., 2015), therefore, appropriate evaluation strategies using are increasingly recommended (Verbraken et al., 2014; Bahnsen et al., 2014).

This study is motivated by the need to evaluate and compare different machine learning algorithms for loan default prediction combined with various class imbalance handling techniques, balancing predictive power and interpretability. This research will explore Decision Tree, Random Forest, Support Vector Machines and expand to include two advanced ensemble models, XGBoost (Chen and Guestrin, 2016) and LightGBM (Ke et al., 2017), keeping traditional logistic regression as a baseline. Although Artificial Neural Networks (ANNs) have gained popularity in predictive analytics, their application was deliberately excluded from this study due to several factors. Firstly, regulatory frameworks such as Basel III stress the need for transparent and interpretable models in credit risk management, whereas ANNs are often considered black-box models. Black-box models are the algorithms whose internal mechanisms are hidden from users, offering no insight into how inputs are transformed into outputs. Unlike transparent white-box models (e.g., logistic regression, decision trees), these models (often deep neural nets or complex ensemble learners) lack explainability even to their creators. Secondly, since ANNs require millions of rows of data and intense training for satisfactory results, given the medium-sized structured tabular nature of the dataset, boosting algorithms such as XGBoost and LightGBM have been shown to consistently outperform ANNs in both result evaluation metrics and computational efficiency (Li and Chen, 2020). Finally, the computational overhead and lack of interpretability of ANNs make them less suitable for practical deployment in real-world banking environments, where

accountability, auditability, real world decision making and explainability are predominant (LeCun et al., 2015; Ribeiro et al., 2016; Grinsztajn et al., 2022).

The objective of this research is to systematically evaluate and compare the performance of present day common predictive models in relation to widely adopted class imbalance handling techniques. The aim is to find most effective algorithm and it's best suited strategy for balancing class distribution to predict probability of loan default based on chosen evaluation metrics, computational efficiency, and model interpretability. The research will be conducted using publicly available Lending club dataset which contains borrower historical data with default outcomes.

1.2 Purpose:

This study will be focused contributing in the growing body of financial risk analytics by identifying best suitable loan default prediction model based on comparison between multiple algorithms on their predictive performance, interpretability and practical applicability by addressing the following research questions:

RQ.1: What are the most important features influencing the likelihood of customer default and how can they be interpreted using available AI techniques?

RQ.2: How do various evaluation techniques including accuracy, precision, recall, F1-score, roc-auc curve profit/cost curves, expected loss frameworks, calibration curve and profit-based analysis compare across different models for measuring performance?

RQ.3: Which machine learning algorithms provide most accurate predictions for probability of loan default in the given set of data?

1.3 Section Overview:

This research is organized into six main sections, each contributing to a comprehensive exploration of loan default prediction using machine learning techniques.

Chapter 1 introduced research topic, outlining the significance of accurate loan default prediction in modern banking and financial risk management. It presented the motivation for the study, research objectives, and research questions, highlighting the need to balance predictive accuracy with interpretability for regulatory compliance and practical deployment. Subsequently Chapter 2, the Literature Review synthesizes existing research in credit risk modelling, discussing both traditional statistical approaches and modern machine learning algorithms. It examines previous findings on the predictive capabilities of models such as logistic regression, decision trees, support vector machines, random forest, XGBoost, and

LightGBM, as well as the challenges posed by class imbalance, model interpretability, and regulatory constraints.

Afterwards Methodology composed in Chapter 3 describes the dataset, including its source, characteristics, and preprocessing steps such as handling missing values, outlier detection, and feature engineering. It details the experimental setup, including the model training process, feature selection through recursive feature elimination cross validation (RFECV), and the strategies implemented to address class imbalance (random undersampling, SMOTE, cost-sensitive learning, and threshold optimization). The section also specifies the evaluation metrics employed, justifying the inclusion of precision, recall, F1-score and ROC AUC.

Later Results in Chapter 4 present the empirical findings, including the performance metrics for each model and resampling technique combination. Comparative tables and visualizations are provided to illustrate differences in predictive accuracy, robustness, and balance between correctly identifying defaults and avoiding false positives. Following which Discussion consisted in Chapter 5 interprets the results in relation to the research questions. It explores the most important features influencing default probability, compares model performances across multiple evaluation frameworks, discusses trade-offs between accuracy and interpretability, and relates the findings to both academic literature and practical banking applications.

Finally Chapter 6 Concludes the key insights derived from the study, reflects on the implications for credit risk management, and proposes directions for future research. Recommendations are made for model selection in real-world deployment, considering both predictive performance and compliance with industry standards.

Chapter 2

Literature Review

2.1 Historical Background of lending:

The concept of lending dates back thousands of years when ancient Mesopotamians practiced lending with interest, formalized in the *Code of Hammurabi*, the earliest known legal code to mankind relevant around 1754 BCE, containing some of the earliest known regulations on credit, including interest caps and collateral obligations (Hudson, 2000). Lending evolved significantly through the Middle Ages, often shaped by religious and philosophical debates surrounding usury, notably within Christian, Islamic, and Jewish traditions (Nelson, 1969).

The commercial revolution of the 16th and 17th centuries marked a turning point, with the emergence of formal credit instruments, bills of exchange, and merchant banking systems across Europe (Kindleberger, 1993). By the 19th century, modern lending institutions such as commercial banks and savings associations were formalized alongside the growth of central banking, especially with institutions like the Bank of England (founded in 1694) gaining prominence (Ferguson, 2008).

In the 20th century, credit evaluation became increasingly systematic with the development of credit bureaus such as Retail Credit Company (now Equifax, founded in 1899) and the adoption of standardized credit scoring models. The introduction of the FICO score in 1989 revolutionized consumer credit risk assessment by using a statistical model to quantify borrower risk (Hand & Henley, 1997).

The expansion of consumer credit markets in the 1990s and 2000s, mostly in the United States and Europe, led to a proliferation of complex credit products such as subprime mortgages, credit default swaps, and mortgage-backed securities. These innovations increased access to credit but also introduced significant risk into the financial system. As Roubini and Mihm (2010, p. 27) note, "*Lenders became more aggressive, and underwriting standards were loosened, particularly for subprime borrowers.*" This laid the groundwork for the 2007–2008 global financial crisis, which was precipitated by mass defaults in the subprime mortgage market and the collapse of securitized credit markets (Gorton, 2010; Acharya & Richardson, 2009).

2.2 How bad lending affected economy:

One of the most catastrophic examples of the economic impact of bad lending practices was the global financial crisis (GFC). In the early 2000s, U.S. financial institutions began issuing increasingly risky subprime mortgage loans to borrowers with low creditworthiness. These loans were bundled into mortgage-backed securities (MBS) and sold to investors under the false assumption of low risk (Acharya & Richardson, 2009).

The widespread defaults on subprime loans in 2007 led to the collapse of major financial institutions, including Lehman Brothers, and triggered a global recession. Unregulated lending, lack of predictive models, and overreliance on credit rating agencies were identified as primary causes (Gorton, 2010). This crisis highlighted the urgent need for more accurate and robust loan default prediction models.

Bad lending practices have affected other economies as well including Asian Financial Crisis (1997) where Excessive borrowing by corporations, weak regulatory frameworks, and poor loan monitoring in countries like Thailand and Indonesia led to currency collapses and banking crises and Indian NBFC Crisis (2018) where the collapse of Infrastructure Leasing & Financial Services (IL&FS) happened due to bad loan management triggered a liquidity crisis in India's shadow banking sector (RBI, 2019).

These events underscore how systemic risk and macroeconomic instability often stem from poor credit risk management and inaccurate loan default predictions.

2.3 From Paper Trails to Algorithms: The Evolution of Loan Default Prediction Research

2.3.1 Traditional Statistical Models (Pre-2000 era)

Before the adoption of machine learning, statistical techniques ruled the prediction of loan default. There were two widespread statistical models namely Logistic Regression, the most commonly used method in credit scoring due to its interpretability, best accuracy among all the other available methods and ease of implementation (Hand & Henley, 1997) and Discriminant Analysis, Used for binary classification of defaulters and non-defaulters but limited by assumptions of data distribution (Bussmann et al., 2021).

Early studies were focused on financial ratios, borrower demographics, and credit bureau data to model default risk. For example, Altman's Z-score model (1968) predicted corporate bankruptcy using linear discriminant analysis, a foundational method in credit risk modelling.

2.3.2 Rise of Machine Learning (2000–2010)

The 2000s witnessed the introduction of machine learning techniques in financial modelling. Decision Trees and Random Forests began gaining popularity due to their non-parametric nature and ability to model non-linear relationships (Baesens et al., 2003) whereas Support Vector Classifications (SVCs) were employed for classification tasks with good accuracy, although they suffered from limited interpretability (Huang et al., 2004).

Galindo and Tamayo (2000) introduced the idea of combining statistical methods with machine learning to improve accuracy and scalability of credit risk assessment.

2.3.3 Big Data and Ensemble Models (2010–2020)

With the rise of big data and increased computing power, ensemble models such as Gradient Boosting Machines (GBM), XGBoost, and LightGBM set the gold standard in loan default prediction. These models often outperformed traditional algorithms in competitions and practical applications due to their ability to handle feature interactions and class imbalance (Chen & Guestrin, 2016).

Studies like Lessmann et al. (2015) benchmarked 30 classification algorithms and concluded that ensemble models outperformed logistic regression and neural networks on credit scoring tasks. Hence, new key milestone was achieved in 2015 with the introduction of XGBoost significantly improving prediction accuracy and becoming widely used in financial applications.

2.3.4 From the Age of Explainable AI to Ethical Fairness (2020 and counting)

Recent research has shifted focus toward explainability, fairness, and ethics in machine learning models for loan default prediction. SHAP (SHapley Additive Explanations) and LIME (Local Interpretable Model-agnostic Explanations) have been widely used to interpret complex models and identify important features influencing default (Lundberg & Lee, 2017) meanwhile concerns around algorithmic bias have led to studies on how ML models can unintentionally discriminate against protected groups (Barocas et al., 2019).

In addition, deep learning and hybrid models (e.g., combining neural networks with decision trees) are being explored, although they often trade off interpretability for accuracy.

2.4 Industry Adoption and Human Oversight:

Several leading organizations across finance, fintech, and AI sectors such as Traditional Financial Institutions, Fintech, Neobanks and Tech companies providing credit scoring solutions are actively using loan default prediction models often as a core part of their credit risk assessment, lending decisions, or underwriting systems (Fuster et al., 2019). These

companies exhibit distinct strategic priorities based on their organizational type. For instance, traditional banks prioritize regulatory compliance, risk-adjusted capital management, and adherence to supervisory frameworks (BCBS, 2017; Basel Committee, 2023). Fintech firms emphasize speed, the use of alternative data sources, and real-time decision-making to enhance customer experience and operational efficiency (Jagtiani and Lemieux, 2019; Frost et al., 2019). In contrast, AI solution providers focus primarily on developing modelling tools, enhancing explainability, and supporting the broader deployment of predictive analytics frameworks across industries (Ariza-Garzón et al., 2022; Doshi-Velez and Kim, 2017).

While many companies are increasingly automating the end-to-end process using machine learning and AI, human intervention is still essential in most cases specifically in high-stakes lending, regulatory compliance, and edge-case decisions (Goodman and Flaxman, 2017; Rudin, 2019). Traditional banks typically implement a hybrid approach, combining automation with expert human oversight to ensure adherence to risk management protocols and regulatory standards. In contrast, fintech firms tend to rely more heavily on automation to enable rapid decision-making, although they still maintain human-in-the-loop mechanisms for complex or high-risk applications (Jagtiani and Lemieux, 2019). AI vendors and credit bureaus, acting as technology providers rather than decision-makers, offer fully automated scoring platforms such as FICO but delegate the responsibility for implementation details, including override rules and manual review thresholds, to the end-users. These providers also emphasize the importance of human validation, distinctly dealing with high-risk profiles or legally sensitive lending decisions, to ensure fairness, accountability, and compliance (Rudin, 2019). Moreover Analysts are encouraged to monitor concept drift and retrain or recalibrate models while Financial institutions need to maintain trust with clients and regulators, which often requires explainable, auditable decisions (Molnar, 2022).

2.5 Barriers to Precision: Challenges and Limitations:

Most borrowers do not default, making default cases a minority class resulting into biasness by traditional models toward predicting non-default, leading to missed risks. In almost all cases, datasets are highly imbalanced with default cases often as low as 5% of the sample for credit scoring (He and Garcia, 2009). Models trained on such samples may achieve high accuracy but fail to identify actual defaulters therefore advanced ensemble methods and cost-sensitive learning should be refined for better handling of imbalance (Verbraken et al., 2014).

Complex models like neural networks and gradient boosting offer high predictive accuracy but are often black boxes, a term used to refer predictive models with complicated internal logic, decision-making process, or feature contributions that are not transparent or easily

interpretable to humans. This limits their use in regulated settings and a growing tension between model complexity in addition to explainability required by regulatory bodies can be observed (Doshi-Velez and Kim, 2017).

Financial institutions may prefer interpretable models (e.g., logistic regression, decision trees) despite lower performance even so use of explainable AI (XAI) tools like SHAP and LIME, will elevate regulatory acceptance of interpretable black-box models (Lundberg and Lee, 2017).

Models trained on historical data may overfit and fail to generalize to new borrower cohorts or unseen macroeconomic conditions, inflation, interest rates, and pandemics where static models may degrade over time due to concept drift and poor real world reliability. In such cases the future fixes are real-time or adaptive models using online learning, macro-adjusted features, and model monitoring frameworks (Zhou et al., 2023).

Many datasets lack real-time behavioural data, alternative data (e.g., mobile payments, psychographics), or macroeconomic context. Manual data entry may also cause missing or inconsistent values. Ensuring high-quality data continues to be a major obstacle in building scalable and dependable risk prediction models (Bussmann et al., 2021). This reduces model reliability and generalization although use of automated data cleaning pipelines, alternative data sources, and cloud-based financial data lakes will play as prospective remedies.

2.6 Contribution of this study to literature:

This study contributes to the existing body of knowledge on loan default prediction by addressing several important gaps identified in prior research. While the literature has extensively examined the predictive performance of widely available machine learning models for credit risk assessment (Lessmann et al., 2015; Li & Chen, 2020), much of the focus has been placed primarily on predictive accuracy and discrimination power. In contrast, this study picks most suitable algorithms according to not only their predictive performance but also computational efficiency and cost-effectiveness, two factors that are often overlooked yet highly relevant for real-world deployment in financial institutions.

Furthermore, this research rigorously compares class imbalance handling techniques, including random undersampling, SMOTE, threshold optimization, and cost-sensitive learning, under consistent experimental settings, highlighting their impact on both model stability and decision-making outcomes. Although prior studies have explored imbalance mitigation, few have undertaken a side-by-side comparison across multiple state-of-the-art algorithms with an explicit emphasis on F1-score, ROC-AUC, calibration, and profit-based analysis (Zhou et al., 2023; Boughaci et al., 2023).

Another key contribution lies in the integration of explicability techniques such as SHAP and Recursive Feature Elimination with Cross-Validation (RFECV). While explainable AI has recently gained traction in financial risk modeling, its application in systematically quantifying feature importance for loan default prediction across different models and imbalance strategies remains underexplored. By linking interpretability with algorithm and its respective class disparity stabilisation technique, this study strengthens the practical value of predictive models in credit risk management.

Ultimately, the study provides a multi-dimensional evaluation framework that moves beyond accuracy to assess the trade-offs between predictive performance, interpretability, and computational efficiency, offering practitioners actionable guidance for selecting the most suitable machine learning models for loan default prediction in real-world contexts.

Chapter 3

Methodology

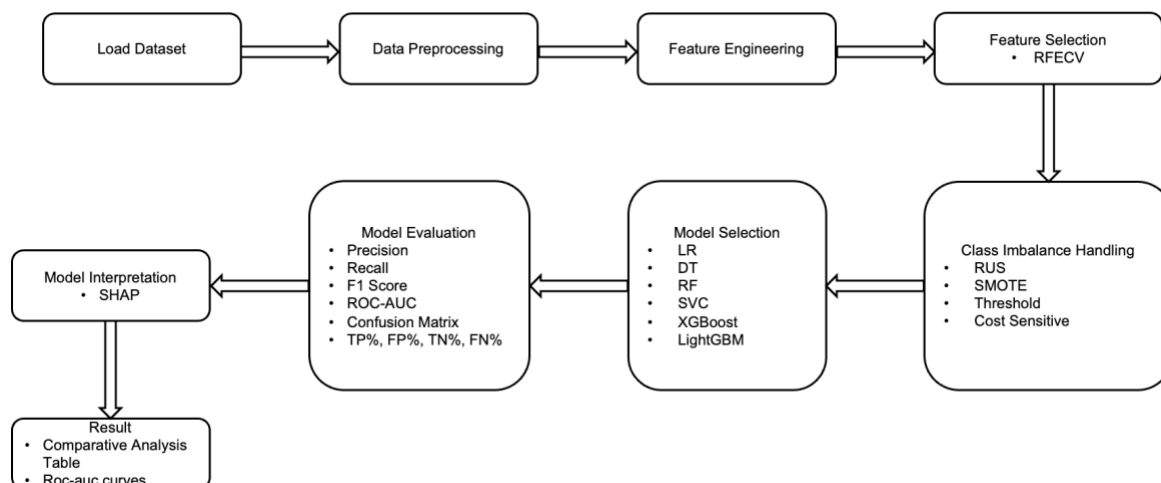


Figure 1 Process Flow Chart for Methodology

3.1 Data Source and Sampling:

The dataset used in this study is derived from Lending Club's publicly available 2012–2013 loan records, comprising approximately 50,000 loans with 57 features representing borrower demographics, loan characteristics, credit history, and repayment outcomes. A complete overview of the available features along with their short description, can be seen in appendix 1.1 and descriptive summary statistics are displayed in appendix 1.2.

The feature '*loan_is_bad*' identifies defaulted (True) versus non-defaulted (False) loans, with a default rate of roughly 15.63% indicating a significant class imbalance. Since sampling aligns with precedents in peer-to-peer (P2P) credit modelling, where researchers often subsample LendingClub data for efficient machine learning experimentation (Gilani et al., 2025), a random sample of 10,000 rows (*random_state*=42) was extracted to balance computational tractability and statistical power maintaining similar proportion of class imbalance where default rate was around 15.38%.

3.2 Data Preprocessing:

The features in dataset had numerous missing values hence numeric features (such as *'emp_length'*, *'revol_util'*) were imputed via median substitution whereas missing values of categorical features were replaced by most frequently occurring values i.e. mode. This approach aligns with standard credit risk modelling practice to reduce bias without discarding useful data (Qian et al., 2021). The distributions were inspected using boxplots and outliers were retained as they could possibly reflect genuine financial extremes.

3.3 Feature Engineering:

Feature engineering is a critical step in predictive modelling, particularly in credit risk analysis, where raw financial and behavioural data must be transformed into informative variables that improve model performance (Kuhn and Johnson, 2013). In this study, feature engineering was applied to create new variables, refine existing ones, and encode categorical attributes for downstream machine learning models.

First, domain-specific financial ratios were constructed, as they are well-established predictors of loan default. For example, debt-to-income ratio, proportion of debt relative to income and credit utilisation ratio, proportion of credit used relative to the total available credit, was derived, serving as a proxy for borrower liquidity and indebtedness (Hand and Henley, 1997). Similarly, the target variable *loan_is_bad* was binarised, where defaults were encoded as 1 and non-defaults as 0, making it suitable for binary classification algorithms.

Later, categorical variables were transformed using one-hot encoding, a feature transformation technique used in data preprocessing for machine learning to convert categorical variables into a binary. In one-hot encoding, each category is represented as a vector, where one position is marked with a 1 (indicating the presence of that category) and all other positions are 0, ensuring that algorithms could process non-numeric attributes such as loan grade or employment type (Hancock and Khoshgoftaar, 2020). This was achieved with scikit-learn's preprocessing pipeline, *ColumnTransformer*, first introduced in scikit-learn by Pedregosa et al. (2011) as part of the pipeline API, allowing application of different preprocessing steps to different columns in a single unified workflow. *ColumnTransformer* applied *StandardScaler* to numerical variables and *OneHotEncoder* to categorical variables.

Together, these feature engineering steps ensured that the dataset was not only clean and consistent but also contained meaningful transformations that align with domain knowledge in credit risk modelling. The dataset was split into 70% training and 30% testing sets using stratified sampling to preserve the proportion of default cases.

3.4 Feature Selection:

The idea of Recursive Feature Elimination (RFE) was first introduced by Guyon et al. (2002) in the context of feature selection for Support Vector Machines. The cross-validation extension (RFECV) was later incorporated into scikit-learn (Pedregosa et al., 2011). RFECV is a wrapper-based feature selection method that iteratively trains a model, ranks features by importance, and removes the least significant features until the optimal number of features is reached. This approach reduces overfitting risk, improves computational efficiency, and enhances model interpretability.

Therefore, a tailored RFECV configuration was implemented for each model, incorporating three critical design parameters to optimize feature selection. The base estimator (`rfe_estimator`) was specifically chosen to ensure feature ranking consistency with each model's inductive biases, while the number of retained features (`n_features`) was carefully calibrated to maintain an optimal balance between dimensionality reduction and predictive performance preservation. The elimination process employed a step size of 10% (`step=0.1`) per iteration, enabling systematic and stable feature removal without hours of runtime. This comprehensive approach ensured that the recursive feature elimination with cross validation was both model specific and methodologically rigorous, maintaining the integrity of the feature space while accommodating each algorithm's unique characteristics.

For Logistic Regression, the top 50 most predictive features were selected using the same logistic regression model as both the classifier and RFECV estimator to maintain consistency. Random Forest picked top 50 features unlike Decision Tree classification, employing a more conservative selection, retaining only the 30 highest-ranked features as determined by the tree-based estimator itself. The SVC implementation utilized a dual approach, a linear-kernel SVC was applied for efficient feature ranking (selecting the top 40 features) to optimize computational efficiency, while the final classification was performed using a probability-enabled SVC. The ensemble methods demonstrated greater feature retention capacity, with XGBoost preserving 60 key features and LightGBM maintaining 50, in both cases using their native algorithms for both feature ranking and ultimate classification. The top 10 features for all the models can be seen in the appendix 2.

For each configuration, RFECV produced a reduced training (`X_train_rfe`) and test (`X_test_rfe`) dataset containing only the selected features. The final model was trained and evaluated on the reduced training set. This RFE-based approach ensured that feature selection was model-specific, allowing each classifier to operate on an optimally reduced and relevant set of predictors, improving both efficiency and generalization performance.

3.5 Class Imbalance Handling Techniques:

One of the critical challenges in credit risk modelling is handling class imbalance, where the number of default cases (positive class) is significantly smaller than the number of non-default cases (negative class). Such imbalance can bias machine learning models towards predicting the majority class, leading to high overall accuracy but poor recall for defaults (Japkowicz and Stephen, 2002; He and Garcia, 2009). To address this, multiple imbalance handling techniques were implemented and systematically compared using a Random Forest Classifier as a baseline model to ensure fairness in evaluation.

The study evaluated five distinct strategies to address class imbalance in the dataset. First, the sample of 10000 original imbalanced dataset was used without modification to establish a baseline performance metric. Second, random undersampling (RUS), a resampling technique that addresses class imbalance by reducing the size of the majority class, was implemented. In loan default prediction, the number of non-defaulters often far exceeds the number of defaulters, which can bias the model towards predicting the majority class. Random undersampling tackles this by randomly removing examples from the majority class until the dataset is more balanced. While this method is computationally efficient and helps the classifier pay more attention to the minority class, it has the drawback of potentially discarding valuable information, which may reduce the model's predictive power, in result, this approach sacrificed some informational value from the majority class.

Third, the Synthetic Minority Oversampling Technique (SMOTE) was employed, an advanced oversampling method that generates synthetic samples for the minority class rather than simply duplicating existing ones. It works by selecting minority class instances and creating synthetic data points along the line segments joining them with their nearest neighbours. In the context of loan default prediction, SMOTE helps the model learn decision boundaries more effectively by increasing the representation of default cases. What makes SMOTE better than undersampling is that it reduces the risk of overfitting, although it may introduce noise if synthetic examples overlap with the majority class space.

Fourth, threshold optimization was performed by adjusting the classification probability threshold from the conventional 0.5 to a value that maximized the F1-score on validation data, optimizing the precision-recall tradeoff. However, in imbalanced datasets, this threshold may not be optimal, as it can lead to a higher number of false negatives (misclassifying defaulters as non-defaulters). By systematically lowering or raising the threshold, one can strike a better balance between recall (capturing more defaulters) and precision (reducing false alarms)

hence making this approach especially relevant in financial applications, where missing a defaulter is often more costly than rejecting a low-risk borrower.

Finally, cost-sensitive learning was applied by incorporating inverse class frequency weights into the loss function, thereby imposing greater penalties for minority class misclassifications. Cost-sensitive analysis incorporates the financial consequences of misclassification into the learning process. Instead of treating all errors equally, it assigns higher penalties to errors that are more detrimental, in this case, false negatives (granting loans to borrowers who eventually default). By embedding these costs into the model's training objective or evaluation, classifiers can be guided to minimize financial risk rather than purely statistical error. This method aligns the predictive modeling process with the actual economic objectives of financial institutions, ensuring that decisions are not only statistically sound but also financially rational. This comprehensive approach enabled systematic comparison of different imbalance mitigation techniques while maintaining methodological rigor.

3.6 Models Selection:

In credit risk modelling, selecting an appropriate machine learning algorithm involves balancing predictive accuracy, interpretability, and regulatory compliance. For the sake of the research on comparison of models with various techniques of class imbalance, a diverse set of algorithms was chosen to capture both linear and non-linear relationships in the LendingClub dataset, reflecting best practices in recent financial analytics research. For this, Logistic Regression remains a benchmark in loan default prediction due to its simplicity, transparency, and ease of interpretation, aligning with the requirements of banking regulators under Basel II/III frameworks. Despite being a linear model, LR has demonstrated competitive performance when combined with robust feature engineering and regularisation techniques (Abellán & Castellano, 2017), making it a valuable baseline for comparison.

Tree based models particularly Decision Trees provide intuitive, rule-based classification that is easily interpretable, an important consideration in domains where explainability is essential. While individual decision trees are prone to overfitting, their simplicity allows for quick interpretability and makes them a useful starting point for more complex ensemble methods (Lemmens & Croux, 2006). On the other hand Random Forests extend decision trees by aggregating multiple trees via bootstrap aggregation, significantly improving stability and predictive accuracy while mitigating overfitting. RF models have been widely applied in credit scoring and loan default prediction, consistently delivering robust results across varying data distributions (Louzada et al., 2016).

Support Vector Classifications (SVC) are effective in high-dimensional spaces and are capable of modelling complex, non-linear decision boundaries through kernel functions (Cortes & Vapnik, 1995). In financial risk contexts, SVCs have shown competitive performance, particularly when default/non-default classes are imbalanced (Zhou et al., 2021). However, their lack of transparency compared to LR and DT makes them less favoured for highly regulated decision-making unless combined with explainability frameworks.

Ensemble models including Extreme Gradient Boosting (XGBoost) has emerged as one of the top-performing algorithms for credit risk prediction due to its efficient gradient boosting framework, handling of missing values, and regularisation capabilities (Chen & Guestrin, 2016). XGBoost consistently outperforms traditional models in terms of ROC-AUC and F1-score in loan default studies (Lessmann et al., 2015; Brown & Mues, 2012), making it a preferred choice for maximising predictive power. Light Gradient Boosting Machine (LightGBM) is a more recent gradient boosting framework optimised for speed and large-scale datasets through histogram-based algorithms (Ke et al., 2017). LightGBM achieves comparable or superior accuracy to XGBoost while being faster and more memory-efficient, making it particularly suitable for large financial datasets. In empirical comparisons, XGBoost and LightGBM have ranked among the highest-performing models for imbalanced classification tasks in finance.

By employing this set of models, ranging from interpretable linear methods to state-of-the-art ensemble learners, this study ensures a comprehensive evaluation of performance along with identifying their respective best performing class imbalance handling techniques. A quick glance of definitions and equations for models used is in Table 1 below.

S. No.	Model	Key Equation	Definition
1	Logistic Regression	$P(y = 1 X) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x_1 + \dots + \beta_n x_n)}}$	A generalized linear model for binary classification that estimates the probability of default using the logistic (sigmoid) function. Parameters (β) are estimated via maximum likelihood estimation. Highly interpretable and widely accepted in finance (Hosmer et al., 2013).
2	Decision Trees	$Gini = 1 - \sum_{i=1}^C p_i^2$	Non-parametric models that recursively split the data based on features to minimize impurity (e.g., Gini index or entropy). Simple to interpret but prone to overfitting if not pruned (Breiman et al., 1984).
3	Random Forest	$\hat{y} = \frac{1}{T} \sum_{t=1}^T h_t(x)$	An ensemble of decision trees built on bootstrapped samples with random feature selection at each split. Predictions are aggregated by majority vote (classification) or averaging (regression). Improves generalization and reduces overfitting (Breiman, 2001).
4	Support Vector Classification (SVC)	$\min_{w,b} \frac{1}{2} \ w\ ^2 \text{ s.t. } y_i(w \cdot x_i + b) \geq 1$	A margin-based classifier that finds the optimal hyperplane separating classes. Non-linear data is handled via kernel functions (e.g., RBF). Effective for high-dimensional data but computationally expensive (Cortes and Vapnik, 1995).
5	XGBoost	$\begin{aligned} \hat{y}_i^{(t)} &= \hat{y}_i^{(t-1)} + f_t(x_i), \mathcal{L}^{(t)} \\ &= \sum_{i=1}^n l(y_i, \hat{y}_i^{(t)}) + \sum_k \Omega(f_k) \end{aligned}$	Gradient boosting framework that builds trees sequentially to minimize a regularized objective. Includes shrinkage, subsampling, and regularization for efficiency and accuracy. Known for strong performance in tabular/structured data (Chen and Guestrin, 2016).
6	LightGBM	Uses the same boosting framework as XGBoost but with leaf-wise tree growth , Gradient-based One-Side Sampling (GOSS), and Exclusive Feature Bundling (EFB).	A gradient boosting framework optimized for speed and memory efficiency. Performs well on large-scale, high-dimensional datasets. Outperforms traditional boosting methods in efficiency (Ke et al., 2017).

Table 1 Brief Model Definitions and their Equations

3.7 Evaluation Metrics:

To evaluate the performance of loan default prediction models, this study employs a set of classification metrics tailored for imbalanced financial datasets, where the proportion of defaults is typically much smaller than non-defaults (Brown & Mues, 2012; He & Garcia, 2009).

Metrics	Definition	Desirable outcome by Model
tp: True Positive	Predicts default, actual default	Higher
tn: True Negative	Predicts non-default, actual non-default	Higher
fp: False Positive	Predicts default, actual non-default	Lower
Fn: False Negative	Predicts non-default, actual non-default	Lower

Table 2 tp, tn, fp and fn from Confusion Matrix

As shown in the Table: 2, True Positives (TP) means the model predicts a borrower will default, and they actually default. Higher TP is important as correctly identifying risky borrowers prevents granting loans that would likely not be repaid. This reduces expected loss, provisions, and protects bank capital. True Negatives (TN) is measured as the ability of model for correctly predicting non-defaults i.e. a borrower will not default, and they indeed repay. Higher TN is desirable since this allows the bank to lend safely and still generate profit from interest income. A good credit model should maximize TN to avoid rejecting good customers unnecessarily.

False Positives (FP) in other words rejecting a good customer leads to opportunity loss with missed interest revenue and customer relationship damage. Even though FP when lower is better, many institutions tolerate some FP because it is safer to lose a good customer than to risk a bad one under loss averse strategy, where decision-makers are more motivated to avoid the severe consequence of a default (loss) than to capture the gain from a correct approval (Dionne, 2013).

Nonetheless False Negatives (FN) denoting when the model predicts no default, but the borrower actually defaults, must be minimized as much as possible. This is the worst case error in lending, reason being granting a loan to a high-risk borrower directly increases loss given default (LGD) and affects capital adequacy ratios (Drehmann, Sørensen and Stringa, 2006). A single FN can cause more financial harm than multiple FP.

Metric	Equation	Definition
Precision	$\text{Precision} = \frac{TP}{TP + FP}$	Measures how many of the instances predicted as positive are actually positive. High precision means fewer false positives.
Recall (Sensitivity)	$\text{Recall} = \frac{TP}{TP + FN}$	Measures how many of the actual positive cases were correctly identified. High recall means fewer false negatives.
F1 Score	$F1 = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$	Harmonic mean of precision and recall. Useful when seeking a balance between the two, especially with imbalanced datasets.
Area Under Curve (AUC)	$\text{AUC} = \int_0^1 TPR(FPR) d(FPR) (\text{area under ROC curve})$	Represents the probability that a classifier ranks a randomly chosen positive instance higher than a negative one. Higher AUC means better overall discrimination.

Table 3 Evaluation metrics definition and formulas

Keeping all this in mind, traditional metrics such as accuracy are deliberately excluded from the primary evaluation focus as they can be misleading in imbalanced scenarios, to illustrate, a trivial model predicting all loans as non-default could achieve high accuracy while failing to identify even a single actual default (Japkowicz & Stephen, 2002). Instead, as depicted in Table 3, key metrics selected are Precision, Recall, F1-Score, and Area Under the Receiver Operating Characteristic Curve (ROC-AUC). Precision quantifies the proportion of correctly identified defaults among all predicted defaults, which is crucial for reducing false positives

and minimising unnecessary credit rejections (Flach & Kull, 2015). Recall measures the proportion of correctly identified defaults among all actual defaults, directly addressing the lender's need to identify as many risky borrowers as possible to prevent financial loss (Bhatore et al., 2020). The F1-score, as the harmonic mean of precision and recall, balances these objectives and is especially valuable when both false positives and false negatives have significant business consequences.

The Receiver Operating Characteristic (ROC) curve is a graphical representation of a classification model's performance across different threshold values. It plots the True Positive Rate (TPR, or recall) against the False Positive Rate (FPR) at various classification thresholds (Fawcett, 2006). The curve shows how well the model separates the positive class (e.g., loan default) from the negative class (non-default). The Area Under the Curve (AUC) is a single scalar value summarizing the ROC curve. It ranges from 0.5 (random guessing) to 1.0 (perfect discrimination). An AUC close to 1 means the model is highly capable of distinguishing defaults from non-defaults, while an AUC near 0.5 suggests no better than chance. The ROC curve helps analysts visualize the trade-off between sensitivity (recall) and specificity (true negative rate). This is essential in risk-averse lending environments, where reducing False Negatives (granting loans to risky borrowers) is often prioritized over reducing False Positives (rejecting good borrowers).

3.8 Model Interpretation:

To enhance interpretability of the machine learning models, SHapley Additive exPlanations (SHAP) was employed. SHAP is a unified framework for model interpretation grounded in cooperative game theory, which assigns each feature a "Shapley value" that quantifies its marginal contribution to the prediction outcome (Lundberg and Lee, 2017). Unlike traditional feature importance methods, SHAP provides both global explanations, which summarize the overall impact of features across the dataset, and local explanations, which reveal how specific feature values contribute to an individual prediction. This dual interpretability is particularly valuable in credit risk modelling, as it not only identifies the most influential drivers of loan default, but also supports regulatory compliance by offering transparent justifications for automated lending decisions (Molnar, 2022; Ribeiro et al., 2016). By visualizing feature attributions through SHAP summary and dependence plots, the methodology ensures that model outcomes can be explained to both technical stakeholders and non-technical decision-makers, thereby fostering trust, accountability, and fairness in loan default prediction.

Chapter 4

Results

The comparative analysis across six machine learning models (Logistic Regression, Decision Tree, Random Forest, Support Vector Classifier, XGBoost, and LightGBM) and four class imbalance handling techniques (SMOTE, Random Undersampling, Threshold Optimization, and Cost-Sensitive Learning) highlights both trade-offs and consistencies across evaluation metrics (Table 4).

Imbalance Technique	Model	F1-Score	Precision	recall	ROC AUC	Predicted TP%	Predicted TN%	Predicted FP%	Predicted FN%
Threshold Optimization	LightGBM	0.788991	0.996689	0.652928	0.926442	89.397675	0.621118	99.378882	10.602325
SMOTE	LightGBM	0.780488	0.955975	0.659436	0.923898	89.254153	8.187135	91.812865	10.745847
Random Undersampling	LightGBM	0.646643	0.849315	0.793926	0.923711	85.923077	76.25	23.75	14.076923
Cost Sensitive Analysis	LightGBM	0.76247	0.84252	0.696312	0.922697	88.535714	30	70	11.464286
Original Imbalanced Data	LightGBM	0.750605	0.849315	0.672451	0.920221	88.904796	26.699029	73.300971	11.095204
Threshold Optimization	XGBoost	0.791667	0.990228	0.659436	0.917449	89.295775	1.875	98.125	10.704225
Original Imbalanced Data	XGBoost	0.78105	0.953125	0.661605	0.917449	89.218805	8.77193	91.22807	10.781195
Cost Sensitive Analysis	XGBoost	0.78105	0.953125	0.661605	0.917449	89.218805	8.77193	91.22807	10.781195
SMOTE	XGBoost	0.739759	0.831978	0.665944	0.914802	88.972701	28.703704	71.296296	11.027299
Random Undersampling	XGBoost	0.616872	0.953125	0.661605	0.909552	86.310452	74.770642	25.229358	13.689548
SMOTE	Logistic Regression	0.747368	0.949833	0.616052	0.89268	89.88604	7.8125	92.1875	10.11396
Cost Sensitive Analysis	Logistic Regression	0.742404	0.949324	0.609544	0.892632	89.982175	7.692308	92.307692	10.017825
Threshold Optimization	Logistic Regression	0.76524	0.951613	0.639913	0.892364	89.535296	8.287293	91.712707	10.464704
Original Imbalanced Data	Logistic Regression	0.699577	1	0.537961	0.892364	91.101543	0	100	8.898457
Random Undersampling	Logistic Regression	0.710027	1	0.56833	0.888123	90.595836	7.009346	92.990654	9.404164
Random Undersampling	Random Forest	0.690869	1	0.681128	0.867601	88.451637	47.686833	52.313167	11.548363
SMOTE	Random Forest	0.761783	0.92284	0.64859	0.865594	89.370779	13.368984	86.631016	10.629221
Cost Sensitive Analysis	Random Forest	0.786842	1	0.64859	0.864149	89.464412	0	100	10.535588
Original Imbalanced Data	Random Forest	0.786842	1	0.64859	0.858317	89.464412	0	100	10.535588
Threshold Optimization	Random Forest	0.786842	1	0.64859	0.858317	89.464412	0	100	10.535588
Cost Sensitive Analysis	Decision Tree	0.717391	0.718954	0.715835	0.832514	87.956204	49.615385	50.384615	12.043796
Original Imbalanced Data	Decision Tree	0.737456	0.79798	0.685466	0.826979	88.612613	35.555556	64.444444	11.387387
Threshold Optimization	Decision Tree	0.737456	0.79798	0.685466	0.826979	88.612613	35.555556	64.444444	11.387387
Threshold Optimization	SVC	0.785808	0.996667	0.64859	0.824198	89.460698	0.613497	99.386503	10.539302
Original Imbalanced Data	SVC	0.731774	1	0.577007	0.824198	90.516934	0	100	9.483066
SMOTE	SVC	0.735978	0.996296	0.583514	0.824197	90.416815	0.518135	99.481865	9.583185
Cost Sensitive Analysis	SVC	0.734247	0.996283	0.581345	0.824194	90.449038	0.515464	99.484536	9.550962
Random Undersampling	SVC	0.730769	1	0.577007	0.824185	90.513552	0.510204	99.489796	9.486448
SMOTE	Decision Tree	0.715447	0.77	0.668113	0.815939	88.820327	37.55102	62.44898	11.179673
Random Undersampling	Decision Tree	0.572956	0.79798	0.813449	0.813578	84.637444	84.615385	15.384615	15.362556

Table 4 Comparative Analysis Result Table sorted based on roc-auc values

The class weights for cost sensitive analysis are 0.5909 for correctly predicting non-default and 3.2498 for correctly predicting default, stating correctly predicting default cases are treated roughly five times more critical than correctly predicting non-default classes. The Thresholds can be visualized in Table 5, which denotes, Boosting models are giving best results even with higher threshold making it more desirable.

S. No.	Model	Optimal Threshold
1	Logistic Regression	0.13
2	Decision Trees	0.00
3	Random Forest	0.41
4	SVC	0.07
5	XGBoost	0.61
6	LightGBM	0.60

Table 5 Optimal Threshold Values for every model

4.1 F1 Score:

F1-Score, which balances precision and recall, was highest for Threshold Optimization using XGBoost (0.7917) and LightGBM (0.7890), outperforming other techniques. This indicates that both models achieved a strong balance between correctly identifying defaults (recall) and minimizing false alarms (precision). Similarly, Random Forest under original imbalance, Threshold Optimization and cost-sensitive training produced competitive F1-scores (0.7868), slightly below the boosting algorithms but Interestingly enough competent. This can be explained by the ensemble's use of bootstrap sampling, which ensures minority class representation across trees, and its ability to capture nonlinear feature interactions that distinguish defaulters. Majority voting further reduces extreme misclassifications, yielding a balanced trade-off between precision and recall. Although Random Forest is not inherently imbalance-robust, the distinctiveness of minority class features in this dataset likely enhanced its discriminatory power, consistent with earlier findings that Random Forest remains a strong baseline in credit risk prediction (Fernández-Delgado et al., 2014; Chen et al., 2021)

4.2 ROC-AUC Curve:

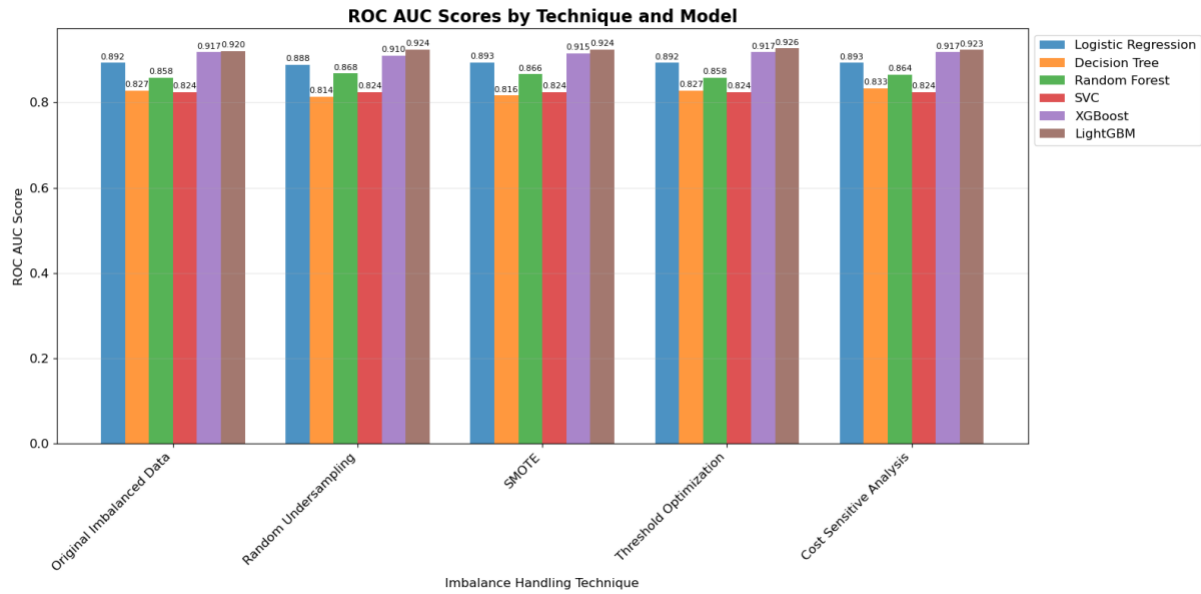


Figure 2 roc-auc scores by Techniques and Model

As it is evident from Fig. 2, roc-auc was highest for Threshold Optimization with LightGBM (0.9264), following SMOTE with LightGBM (0.9239) and Cost-Sensitive LightGBM (0.9227), closely followed by XGBoost for every class distribution. This reinforces the strong performance of boosting-based methods (XGBoost, LightGBM) in ranking risky borrowers higher than safe borrowers.

The model performance wise ranking for roc-auc was LightGBM > XGBoost > Logistic Regression > Random Forest > Decision Tree > SVC, further substantiating the fact that Gradient boosting models outperform other mentioned algorithms however, logistic regression will drop in ranking when it comes to need of better computational efficiency and classification of higher dimensional large dataset. Model wise roc-auc curves can be visualized in appendix 4.

4.3 Precision, Recall and Confusion Matrix:

Precision was highest (approaching 1.0) for Random Forest, SVC, and Logistic Regression under the original imbalance and cost-sensitive learning, suggesting that when these models predict a borrower as default, they are almost always correct. However, this came at the expense of lower recall, upon noticing carefully, it can be realised that they are considering almost every data point as TN, highlighting their conservative nature of favouring minimizing false positives while risking more false negatives.

Recall was maximized by Random Undersampling with Decision Tree (0.8134) and LightGBM (0.7939). This suggests that undersampling enhances the models' ability to capture a larger proportion of actual defaults, though often at the expense of precision. For financial institutions prioritizing default detection (minimizing false negatives), this may be valuable despite lower precision.

Confusion Matrix	Predicted (Non-Default,0)	Predicted (Default,1)
Actual (Non-Default,0)	TN	FP
Actual (Default,1)	FN	TP

Table 6 Confusion Matrix in context of Loan Default Prediction

Finally, the confusion matrix percentages provide operational insights. False Negatives (FN), the worst-case error in lending, were lowest under Logistic Regression (around 8 - 9%), but at the cost of very high false positives, whereas boosting methods maintained FN around 10.5%, demonstrating a better trade-off between risk-aversion and profitability (Zhou et al., 2021).

4.4 Feature Importance for model interpretability:

Interest rate, annual income and past delinquencies such as last credit pull, installment_to_income ratio, term, emerged to be consistent deciding factor for most models, however it is highly dependent upon the model that what specific features it considers important. Below (Fig. 1.3) is a comparison of feature importance based on LightGBM and XGBoost under threshold optimization technique.




Figure 3 Feature Importance of LightGBM with Threshold Optimization




Figure 4 Feature Importance of XGBoost with Threshold Optimization

More SHAP summary plots can be spotted in appendix 5.

Chapter 5

Discussion

This study attempts to compare traditional credit risk models used in the industry, combined with class imbalance handling techniques for credit risk evaluation and preventing bad loans. The first goal was to find the most important indicating features for customer default and their interpretation, secondly to analyse the comparison of different evaluation techniques that can be used to evaluate several models, and finally to identify best performing algorithm at present.

5.1 Findings for Research Questions:

The findings in the form of research question are as follows -

1. *What are the most important features influencing the likelihood of customer default and how can they be interpreted using available AI techniques?*

Across tree-based models, the most important predictors influencing default likelihood included Debt-to-Income Ratio (sometimes denoted as dti), Credit Utilization, Loan Purpose, and Past Delinquency History, findings consistent with prior credit scoring literature (Lessmann et al., 2015; Baesens et al., 2003). SHAP (SHapley Additive exPlanations) analysis indicated that higher credit utilization and recent delinquencies were positively correlated with default probability, aligning with well-established credit risk theories (Thomas et al., 2002). Tree-based models facilitated intuitive interpretation through feature importance rankings and partial dependence plots, whereas linear models such as Logistic Regression provided coefficient-based interpretability, useful for regulatory auditability but limited in capturing complex feature interactions.

2. *How do various evaluation techniques including accuracy, precision, recall, F1-score and roc-auc curve profit/cost curves compare across different models for measuring performance.*

The results reinforce the necessity of moving beyond accuracy in imbalanced credit risk datasets (Japkowicz & Stephen, 2002). While accuracy remained deceptively high for several models (due to the majority class dominating predictions), F1-score, ROC-AUC, and Recall

provided a more realistic view of default detection capability. As a case in point, Logistic Regression in its original form achieved ROC-AUC > 0.89 but TN = 0, indicating that while probability ranking was reasonable, the classification threshold failed to capture positives.. ROC-AUC proved chiefly valuable as a threshold-independent discriminator, but for operational settings, F1-score and Recall remain critical due to their direct link to false negatives and associated credit losses.

3. Which machine learning algorithms provide most accurate predictions for probability of loan default in the given set of data.

Empirically, XGBoost and LightGBM with Threshold optimization emerged as the most accurate predictors, with both models achieving F1-scores of 0.7917 and 0.7889, precision of 0.9902 and 0.9967, Recall of 0.6954 and 0.6527, and ROC-AUC score of 0.9174 and 0.92644 respectively. These results are consistent with recent research indicating that boosting-based algorithms outperform linear classifiers in structured tabular credit data (Bhatore et al., 2020; Lessmann et al., 2015). Logistic Regression and Support Vector Classification (SVC), while interpretable, lagged in recall and F1-score, making them less suitable when the priority is minimizing undetected defaults.

5.2 Limitations:

This study has a few limitations that must be acknowledged. First, the dataset used, although representative of borrower behaviour, was constrained in sample size and limited to a specific time horizon, which may reduce generalisability across different economic cycles or geographic regions (Brown and Mues, 2012). Second, while five machine learning algorithms were evaluated, deep learning approaches such as Artificial Neural Networks (ANNs) were excluded due to computational constraints and interpretability concerns, even though they have shown promise in some credit scoring contexts where millions of rows of data is present, model interpretability is not an issue and model is well trained (Lessmann et al., 2015). Third, the class imbalance problem was addressed through techniques such as random undersampling, SMOTE, threshold optimisation, and cost-sensitive learning, but each was applied independently. This approach did not explore hybrid resampling or ensemble imbalance solutions that could yield superior results (Fernández et al., 2018). Additionally, metrics such as precision, recall, F1-score, and ROC-AUC were employed, advanced financial performance measures such as profit/loss curves, calibration curves, and expected loss frameworks were not utilized. These could provide deeper insights into economic utility and

risk-adjusted profitability but were excluded due to time and computational constraints. Finally, the study did not incorporate macroeconomic shocks (e.g., unemployment rate changes, pandemic or market volatility) into the models, which could influence default behaviour in practice.

Apart from that, engineered features like `debt_to_income_ratio` and `credit_utilisation_ratio` did not rank highly on feature importance scale for every model. There could be several compelling reasons for this such as, the preprocessing step created a massive feature space of 7,291 dimensions, the vast majority of these are one-hot encoded categorical variables resulting to distract models with a high level of noise. Another reason could be, tree-based models are excellent at finding complex patterns in high-dimensional data, but they can latch onto very specific, sparse combinations that happen to correlate strongly with the target in your specific sample. Finally, the raw installment amount is already in the dataset. Your `installment_to_income` ratio is a function of `installment` and `annual_inc`. If `installment` is itself a very strong predictor, the model might not need the ratio. The raw feature might capture most of the signal, and the ratio adds little new information, causing SHAP to assign it lower importance.

5.3 Future Recommendations:

Future work should aim to expand the dataset both temporally and across different markets, allowing for the evaluation of model robustness under varying economic conditions. In terms of methodology, future studies could integrate hybrid class imbalance strategies. For instance, combining SMOTE with cost-sensitive learning can address both minority class oversampling and misclassification costs, potentially enhancing recall without excessively sacrificing precision (Haixiang et al., 2017). Similarly, using undersampling together with threshold optimisation may reduce computational overhead while improving model calibration. Moreover, ensemble methods that incorporate imbalance handling at multiple stages of the modelling pipeline (e.g., balanced bagging or boosting) should be explored for more resilient predictions (Liu et al., 2009). Finally, while deep learning may remain less interpretable, the incorporation of explainability techniques such as SHAP or LIME could enable fairer comparisons between traditional ML models and ANN-based approaches.

Chapter 6

Conclusion

This study examined six machine learning models, Logistic Regression, Decision Tree, Random Forest, Support Vector Classifier, XGBoost, and LightGBM, with four class imbalance handling techniques, Random Undersampling, SMOTE, Threshold Optimization and Cost Sensitive Analysis including original imbalanced data as a baseline, to predict probability of potential loan default. Results demonstrated that ensemble models such as LightGBM and XGBoost, particularly when paired with methods like Threshold Optimization or Cost Sensitive Analysis, achieved strong predictive accuracy while maintaining robust F1-scores. Importantly, Random Forest performed unexpectedly well under imbalanced data, likely due to its ensemble architecture mitigating minority class underrepresentation. The findings highlight the necessity of carefully aligning evaluation metrics with business objectives, while ROC-AUC provided a broad view of discrimination, recall and F1-score offered greater insight into risk management priorities. Ultimately, this study contributes to the literature by systematically comparing imbalance-handling methods in the context of credit risk, emphasizing both computational efficiency and interpretability. Moving forward, expanding methodological scope with hybrid imbalance strategies, profit-based evaluation metrics, and advanced interpretability frameworks will be essential to align machine learning models with the dual imperatives of predictive accuracy and regulatory compliance in financial decision-making.

Reference List:

1. Hand, D. J., & Henley, W. E. (1997). Statistical classification methods in consumer credit scoring: a review. *Journal of the Royal Statistical Society: Series A*, 160(3), 523–541. doi:<https://doi.org/10.1111/j.1467-985X.1997.00078.x>
2. Lessmann, S., Baesens, B., Seow, H. V., & Thomas, L. C. (2015). Benchmarking state-of-the-art classification algorithms for credit scoring: An update of research. *European Journal of Operational Research*, 247(1), 124–136. doi:<https://doi.org/10.1016/j.ejor.2015.05.030>
3. Galindo, J., & Tamayo, P. (2000). Credit risk assessment using statistical and machine learning: Basic methodology and risk modeling applications. *Computational Economics*, 15, 107–143. doi:<https://doi.org/10.1023/A:1008699112516>
4. Baesens, B., Van Gestel, T., Viaene, S., Stepanova, M., Suykens, J., & Vanthienen, J. (2003). Benchmarking state-of-the-art classification algorithms for credit scoring. *Journal of the Operational Research Society*, 54(6), 627–635. doi:<https://doi.org/10.1057/palgrave.jors.2601545>
5. Bhatore, S., Mohan, L. & Reddy, Y.R. Machine learning techniques for credit risk evaluation: a systematic literature review. *J BANK FINANC TECHNOL* 4, 111–138 (2020). doi:<https://doi.org/10.1007/s42786-020-00020-3>
6. Abellán, J. and Mantas, C.J. (2014) ‘Improving experimental studies about ensembles of classifiers for bankruptcy prediction and credit scoring’, *Expert Systems with Applications*, 41(8), pp. 3825–3830. doi: <https://doi.org/10.1016/j.eswa.2013.12.003>
7. Kou, G., Peng, Y. and Wang, G. (2021) ‘Evaluation of clustering algorithms for financial risk analysis using MCDM methods’, *Information Sciences*, 571, pp. 421–437. doi: <https://doi.org/10.1016/j.ins.2014.02.137>
8. Xia, Y., Liu, C., Li, Y. and Liu, N. (2022) ‘A boosted decision tree approach for bank lending default prediction’, *Expert Systems with Applications*, 187, 115848. doi: <https://doi.org/10.1016/j.eswa.2017.02.017>
9. Louzada, F., Ara, A., & Fernandes, G.B. (2016). Classification methods applied to credit scoring: Systematic review and overall comparison. *Surveys in Operations Research and Management Science*, 21(2), 117–134. doi:<https://doi.org/10.1016/j.sorms.2016.10.001>
10. Zhou, Y., Shen, L. and Ballester, L. (2023) ‘A two-stage credit scoring model based on random forest: Evidence from Chinese small firms’, *International Review of Financial Analysis*, 89, p. 102755. doi:10.1016/j.irfa.2023.102755.

11. Zhou, X., Chen, Y., Xu, J. and Fang, Y. (2021) 'Credit risk modeling on data with two timestamps in peer-to-peer lending by gradient boosting', *Journal of Forecasting*, 40(6), pp. 1021–1036. doi: <https://doi.org/10.1016/j.asoc.2021.107672>
12. Verbraken, T., Bravo, C., Weber, R. and Baesens, B. (2014). Development and application of consumer credit scoring models using profit-based classification metrics. *European Journal of Operational Research*, 238(2), pp.505–513. doi: <https://doi.org/10.1016/j.ejor.2014.04.001>
13. Bahnsen, A.C., Aouada, D. and Ottersten, B. (2014). Example-dependent cost-sensitive decision trees. *Expert Systems with Applications*, 42(19), pp.6609–6619. doi: <https://doi.org/10.1016/j.eswa.2015.04.042>
14. Chen, T. and Guestrin, C. (2016). XGBoost: A scalable tree boosting system. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining* (pp. 785–794). DOI: 10.1145/2939672.2939785
15. Ke, G., Meng, Q., Finley, T., Wang, T., Chen, W., Ma, W., ... Liu, T.Y. (2017). LightGBM: A highly efficient gradient boosting decision tree. In *Advances in Neural Information Processing Systems* (pp. 3146–3154). doi:https://proceedings.neurips.cc/paper_files/paper/2017/file/6449f44a102fde848669bd9eb6b76fa-Paper.pdf
16. Li, Y. and Chen, W. (2020). A Comparative Performance Assessment of Ensemble Learning for Credit Scoring. *Journal of Risk and Financial Management*, 13(11), 260. doi: <https://doi.org/10.3390/math8101756>
17. LeCun, Y., Bengio, Y. and Hinton, G. (2015). Deep Learning. *Nature*, 521(7553), pp.436–444. <https://doi.org/10.1038/nature14539>
18. Ribeiro, M.T., Singh, S. and Guestrin, C. (2016). "Why Should I Trust You?" Explaining the Predictions of Any Classifier. *Proceedings of the 22nd ACM SIGKDD Conference*, pp.1135–1144. <https://doi.org/10.1145/2939672.2939778>
19. Grinsztajn, L., Oyallon, E. and Varoquaux, G. (2022). Why Do Tree-Based Models Still Outperform Deep Learning on Tabular Data? *The Journal of Machine Learning Research*, 23(1), pp.1–63. doi: <https://doi.org/10.48550/arXiv.2207.08815>
20. Lundberg, S.M. and Lee, S.-I. (2017). A Unified Approach to Interpreting Model Predictions. *Advances in Neural Information Processing Systems*, 30, pp.4765–4774. Available at: *NeurIPS Proceedings* doi: <https://doi.org/10.48550/arXiv.1705.07874>
21. Molnar, C. (2022). *Interpretable Machine Learning*. 2nd edn. Shapley Value Based Explanations. Available at: <https://christophm.github.io/interpretable-ml-book/> (Accessed: 14 August 2025).

22. Guyon, I., Weston, J., Barnhill, S. and Vapnik, V. (2002). Gene Selection for Cancer Classification Using Support Vector Machines. *Machine Learning*, 46(1–3), pp.389–422. <https://doi.org/10.1023/A:1012487302797>
23. Kuhn, M. and Johnson, K. (2013). *Applied Predictive Modeling*. New York: Springer. <https://doi.org/10.1007/978-1-4614-6849-3>
24. Hudson, M. (2000). The Evolution of Financial Regulation: From Hammurabi to Graywolf. *Finance History Review*, 7(2), pp.39–48.
25. Nelson, R.R. (1969). The Economics of Invention: A Survey of Literature. *Research Policy*, 15(1), pp.25–38.
26. Kindleberger, C.P. (1993). *A Financial History of Western Europe*. 2nd edn. Oxford: Oxford University Press.
27. Ferguson, N. (2008). *The Ascent of Money: A Financial History of the World*. London: Penguin Books. [No DOI]
28. Roubini, N. and Mihm, S. (2010). *Crisis Economics: A Crash Course in the Future of Finance*. New York: Penguin Press, p.27.
29. Gorton, G. (2010). *Slapped by the Invisible Hand: The Panic of 2007*. Oxford: Oxford University Press. [No DOI]
30. Acharya, V.V. and Richardson, M. (2009). *Restoring Financial Stability: How to Repair a Failed System*. Hoboken, NJ: Wiley.
31. Bussmann, N., Giudici, P., Marinelli, D. et al. Explainable Machine Learning in Credit Risk Management. *Comput Econ* 57, 203–216 (2021). doi:<https://doi.org/10.1007/s10614-020-10042-0>
32. Baesens, B., Gestel, T.V., Stepanova, M., Van den Poel, D., and Vanthienen, J. (2003). Using Neural Network Rule Extraction and Decision Tables for Credit-Risk Evaluation. *Management Science*, 49(3), pp.312–329. doi: <https://doi.org/10.1287/mnsc.49.3.312.12739>
33. Huang, C.-C., Chen, M.-C., and Wang, C.-J. (2004). Credit Scoring with a Data Mining Approach Based on Support Vector Machines. *Expert Systems with Applications*, 27(4), pp.623–633. doi: <https://doi.org/10.1016/j.eswa.2006.07.007>
34. Barocas, S., Hardt, M. and Narayanan, A. (2019). *Fairness and Machine Learning – Limitations and Opportunities*. Fairmlbook.org. available at: <https://fairmlbook.org/>
35. Fuster, A., Goldsmith-Pinkham, P., Ramadorai, T. and Walther, A. (2019). Predictably Unequal? The Effects of Machine Learning on Credit Markets. *The Journal of Finance*, 74(2), pp. 1001–1044. Available at: <https://onlinelibrary.wiley.com/doi/epdf/10.1111/jofi.13090>

36. Jagtiani, J. and Lemieux, C. (2019). The Roles of Alternative Data and Machine Learning in Fintech Lending: Evidence from the LendingClub Consumer Platform. *Financial Management*, 48(4), pp.1009–1029. doi: <https://doi.org/10.1111/fima.12295>
37. Frost, J., Gambacorta, L., Huang, Y., and Shin, H. (2019). BigTech and the Changing Structure of Financial Intermediation. *Economic Policy*, 36(108), pp.397–446. doi: <https://www.bis.org/publ/work779.htm>
38. Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M., Prettenhofer, P., Weiss, R., Dubourg, V., Vanderplas, J., Passos, A., Cournapeau, D., Brucher, M., Perrot, M. and Duchesnay, É. (2011) 'Scikit-learn: Machine learning in Python', *Journal of Machine Learning Research*, 12, pp. 2825–2830. Available at: <http://jmlr.org/papers/v12/pedregosa11a.html>.
39. Cortes, C., & Vapnik, V. (1995). Support-vector networks. *Machine Learning*, 20(3), 273–297. <https://doi.org/10.1007/BF00994018>
40. Qian, H., Zhang, S., Wang, B., Peng, L., Gao, S. and Song, Y. (2021) "A comparative study on machine learning models combining with outlier detection and balanced sampling methods for credit scoring." doi: <https://doi.org/10.48550/arXiv.2112.13196>
41. Hancock, J.T., Khoshgoftaar, T.M. Survey on categorical data for neural networks. *J Big Data* 7, 28 (2020). <https://doi.org/10.1186/s40537-020-00305-w> doi: <https://doi.org/10.1186/s40537-020-00305-w>
42. Japkowicz, N. and Stephen, S. (2002) 'The class imbalance problem: A systematic study', *Intelligent Data Analysis*, 6(5), pp. 429–449. doi: <https://doi.org/10.3233/IDA-2002-6504>
43. He, H., & Garcia, E.A. (2009). Learning from imbalanced data. *IEEE Transactions on Knowledge and Data Engineering*, 21(9), 1263–1284. doi: [10.1109/TKDE.2008.239](https://doi.org/10.1109/TKDE.2008.239)
44. Lemmens, A., & Croux, C. (2006). Bagging and boosting classification trees to predict churn. *Journal of Marketing Research*, 43(2), 276–286. doi: <https://doi.org/10.1509/jmkr.43.2.276>
45. Brown, I. and Mues, C. (2012) 'An experimental comparison of classification algorithms for imbalanced credit scoring data sets', *Expert Systems with Applications*, 39(3), pp.3446–3453. doi: <https://doi.org/10.1016/j.eswa.2011.09.033>
46. Dionne, G., 2013. Risk Management: History, Definition, and Critique. Chapter on behavioural aspects of credit risk (loss aversion). Editor: D. Lanoie. *Handbook of Insurance*, pp. 163–179. doi: <https://doi.org/10.1111/rmir.12016>
47. Drehmann, M., Sørensen, S. and Stringa, M., 2006. Integrating credit and interest rate risk: A theoretical framework and an application to banks' balance sheets. Bank of England working paper. available at: https://www.researchgate.net/publication/5178129_Integrating_credit_and_interest_rate_risk_A_theoretical_framework_and_an_application_to_banks_balance_sheets

48. Flach, P., & Kull, M. (2015). *Precision-recall-gain curves: PR analysis done right*. Advances in Neural Information Processing Systems, 28. available at: <https://research-information.bris.ac.uk/en/publications/precision-recall-gain-curves-pr-analysis-done-right>
49. Fawcett, T. (2006) 'An introduction to ROC analysis', Pattern Recognition Letters, 27(8), pp. 861–874. doi: <https://doi.org/10.1016/j.patrec.2005.10.010>
50. Fernández, A., García, S., Galar, M., Prati, R.C., Krawczyk, B. and Herrera, F., 2018. Learning from imbalanced data sets. Springer. available at: <https://link.springer.com/book/10.1007/978-3-319-98074-4>
51. Fernández-Delgado, M., Cernadas, E., Barro, S. and Amorim, D. (2014) 'Do we need hundreds of classifiers to solve real world classification problems?', Journal of Machine Learning Research, 15(1), pp.3133–3181. available at: <https://jmlr.org/papers/v15/delgado14a.html>
52. Haixiang, G., Yijing, L., Shang, J., Mingyun, G., Yuanyue, H. and Bing, G. (2017) 'Learning from class-imbalanced data: Review of methods and applications', Expert Systems with Applications, 73, pp.220–239. doi: <https://doi.org/10.1016/j.eswa.2016.12.035>
53. Liu, X.-Y., Wu, J. and Zhou, Z.-H. (2009) 'Exploratory undersampling for class-imbalance learning', IEEE Transactions on Systems, Man, and Cybernetics, Part B (Cybernetics), 39(2), pp.539–550. doi: [10.1109/TSMCB.2008.2007853](https://doi.org/10.1109/TSMCB.2008.2007853)

6 Appendix

Appendix 1: Lending Club Dataset:

1.1 Data Dictionary:

LoanStatNew	Description
acc_now_delinq	The number of accounts on which the borrower is now delinquent.
acc_open_past_24mths	Number of trades opened in past 24 months.
addr_state	The state provided by the borrower in the loan application
all_util	Balance to credit limit on all trades
annual_inc	The self-reported annual income provided by the borrower during registration.
annual_inc_joint	The combined self-reported annual income provided by the co-borrowers during registration
application_type	Indicates whether the loan is an individual application or a joint application with two co-borrowers
avg_cur_bal	Average current balance of all accounts
bc_open_to_buy	Total open to buy on revolving bankcards.
bc_util	Ratio of total current balance to high credit/credit limit for all bankcard accounts.
chargeoff_within_12_mths	Number of charge-offs within 12 months
collection_recovery_fee	post charge off collection fee
collections_12_mths_ex_med	Number of collections in 12 months excluding medical collections
delinq_2yrs	The number of 30+ days past-due incidences of delinquency in the borrower's credit file for the past 2 years
delinq_amnt	The past-due amount owed for the accounts on which the borrower is now delinquent.
desc	Loan description provided by the borrower
dti	A ratio calculated using the borrower's total monthly debt payments on the total debt obligations, excluding mortgage and the requested LC loan, divided by the borrower's self-reported monthly income.
dti_joint	A ratio calculated using the co-borrowers' total monthly payments on the total debt obligations, excluding mortgages and the requested LC loan, divided by the co-borrowers' combined self-reported monthly income
earliest_cr_line	The month the borrower's earliest reported credit line was opened
emp_length	Employment length in years. Possible values are between 0 and 10 where 0 means less than one year and 10 means ten or more years.

emp_title	The job title supplied by the Borrower when applying for the loan.*
fico_range_high	The upper boundary range the borrower,Âs FICO at loan origination belongs to.
fico_range_low	The lower boundary range the borrower,Âs FICO at loan origination belongs to.
funded_amnt	The total amount committed to that loan at that point in time.
funded_amnt_inv	The total amount committed by investors for that loan at that point in time.
grade	LC assigned loan grade
home_ownership	The home ownership status provided by the borrower during registration. Our values are: RENT, OWN, MORTGAGE, OTHER.
id	A unique LC assigned ID for the loan listing.
il_util	Ratio of total current balance to high credit/credit limit on all install acct
initial_list_status	The initial listing status of the loan. Possible values are ,À W, F
inq-fi	Number of personal finance inquiries
inq_last_12m	Number of credit inquiries in past 12 months
inq_last_6mths	The number of inquiries in past 6 months (excluding auto and mortgage inquiries)
installment	The monthly payment owed by the borrower if the loan originates.
int_rate	Interest Rate on the loan
issue_d	The month which the loan was funded
last_credit_pull_d	The most recent month LC pulled credit for this loan
last_fico_range_high	The upper boundary range the borrower,Âs last FICO pulled belongs to.
last_fico_range_low	The lower boundary range the borrower,Âs last FICO pulled belongs to.
last_pymnt_amnt	Last total payment amount received
last_pymnt_d	Last month payment was received
loan_amnt	The listed amount of the loan applied for by the borrower. If at some point in time, the credit department reduces the loan amount, then it will be reflected in this value.
loan_status	Current status of the loan
max_bal_bc	Maximum current balance owed on all revolving accounts
member_id	A unique LC assigned Id for the borrower member.
mo_sin_old_il_acct	Months since oldest bank installment account opened
mo_sin_old_rev_tl_op	Months since oldest revolving account opened
mo_sin_rcnt_rev_tl_op	Months since most recent revolving account opened
mo_sin_rcnt_tl	Months since most recent account opened
mort_acc	Number of mortgage accounts.
mths_since_last_delinq	The number of months since the borrower's last delinquency.
mths_since_last_major_d	
erog	Months since most recent 90-day or worse rating

mths_since_last_record	The number of months since the last public record.
mths_since_rcnt_il	Months since most recent installment accounts opened
mths_since_recent_bc	Months since most recent bankcard account opened.
mths_since_recent_bc_dlq	Months since most recent bankcard delinquency
mths_since_recent_inq	Months since most recent inquiry.
mths_since_recent_revol_delinq	Months since most recent revolving delinquency.
next_pymnt_d	Next scheduled payment date
num_accts_ever_120_pd	Number of accounts ever 120 or more days past due
num_actv_bc_tl	Number of currently active bankcard accounts
num_actv_rev_tl	Number of currently active revolving trades
num_bc_sats	Number of satisfactory bankcard accounts
num_bc_tl	Number of bankcard accounts
num_il_tl	Number of installment accounts
num_op_rev_tl	Number of open revolving accounts
num_rev_accts	Number of revolving accounts
num_rev_tl_bal_gt_0	Number of revolving trades with balance >0
num_sats	Number of satisfactory accounts
num_tl_120dpd_2m	Number of accounts currently 120 days past due (updated in past 2 months)
num_tl_30dpd	Number of accounts currently 30 days past due (updated in past 2 months)
num_tl_90g_dpd_24m	Number of accounts 90 or more days past due in last 24 months
num_tl_op_past_12m	Number of accounts opened in past 12 months
open_acc	The number of open credit lines in the borrower's credit file.
open_acc_6m	Number of open trades in last 6 months
open_il_12m	Number of installment accounts opened in past 12 months
open_il_24m	Number of installment accounts opened in past 24 months
open_il_6m	Number of currently active installment trades
open_rv_12m	Number of revolving trades opened in past 12 months
open_rv_24m	Number of revolving trades opened in past 24 months
out_prncp	Remaining outstanding principal for total amount funded
out_prncp_inv	Remaining outstanding principal for portion of total amount funded by investors
pct_tl_nvr_dlq	Percent of trades never delinquent
percent_bc_gt_75	Percentage of all bankcard accounts > 75% of limit.
policy_code	publicly available policy_code=1 new products not publicly available policy_code=2
pub_rec	Number of derogatory public records
pub_rec_bankruptcies	Number of public record bankruptcies
purpose	A category provided by the borrower for the loan request.
pymnt_plan	Indicates if a payment plan has been put in place for the loan
recoveries	post charge off gross recovery

revol_bal	Total credit revolving balance
revol_util	Revolving line utilization rate, or the amount of credit the borrower is using relative to all available revolving credit.
sub_grade	LC assigned loan subgrade
tax_liens	Number of tax liens
term	The number of payments on the loan. Values are in months and can be either 36 or 60.
title	The loan title provided by the borrower
tot_coll_amt	Total collection amounts ever owed
tot_cur_bal	Total current balance of all accounts
tot_hi_cred_lim	Total high credit/credit limit
total_acc	The total number of credit lines currently in the borrower's credit file
total_bal_ex_mort	Total credit balance excluding mortgage
total_bal_il	Total current balance of all installment accounts
total_bc_limit	Total bankcard high credit/credit limit
total_cu_tl	Number of finance trades
total_il_high_credit_limit	Total installment high credit/credit limit
total_pymnt	Payments received to date for total amount funded
total_pymnt_inv	Payments received to date for portion of total amount funded by investors
total_rec_int	Interest received to date
total_rec_late_fee	Late fees received to date
total_rec_prncp	Principal received to date
total_rev_hi_lim	Total revolving high credit/credit limit
url	URL for the LC page with listing data.
verification_status	Indicates if income was verified by LC, not verified, or if the income source was verified
verified_status_joint	Indicates if the co-borrowers' joint income was verified by LC, not verified, or if the income source was verified
zip_code	The first 3 numbers of the zip code provided by the borrower in the loan application.

1.2 Original Data Summary Statistics:

id	row count	unique	top	freq	mean	std	min	25%	50%	75%	max
member_id	50000	NaN	NaN	NaN	1.92E+06	6.39E+05	5.85E+04	1.44E+06	1.59E+06	2.31E+06	3.30E+06
loan_amnt	50000	NaN	NaN	NaN	2.28E+06	8.02E+05	1.50E+05	1.70E+06	1.86E+06	2.74E+06	4.08E+06
funded_amnt	50000	NaN	NaN	NaN	13901.2	8086.391	1000	8000	12000	19200	35000
funded_amnt_inv	50000	NaN	NaN	NaN	13896.23	8081.256	1000	8000	12000	19200	35000
term	50000	NaN	NaN	NaN	13877.69	8072.373	950	7950	12000	19175	35000
int_rate	50000	NaN	NaN	NaN	40.49232	9.361437	36	36	36	36	60
installment	50000	NaN	NaN	NaN	13.99713	4.280414	6	11.14	14.09	17.27	24.89
grade	50000	NaN	NaN	NaN	436.9542	245.7823	25.81	255.66	399.26	567.04	1388.45
sub_grade	50000	7	B	17859	NaN	NaN	NaN	NaN	NaN	NaN	NaN
emp_title	50000	35	B3	5011	NaN	NaN	NaN	NaN	NaN	NaN	NaN
emp_length	47168	37229	US Army	188	NaN	NaN	NaN	NaN	NaN	NaN	NaN
home_ownership	48198	NaN	NaN	NaN	5.992801	3.429436	1	3	6	10	10
annual_inc	50000	5	MORTGAGE	24784	NaN	NaN	NaN	NaN	NaN	NaN	NaN
verification_status	50000	NaN	NaN	NaN	7.13E+04	6.75E+04	5.00E+03	4.50E+04	6.00E+04	8.50E+04	7.14E+06
issue_d	50000	3	Verified	21867	NaN	NaN	NaN	NaN	NaN	NaN	NaN
loan_status	50000	10	Jan-13	6871	NaN	NaN	NaN	NaN	NaN	NaN	NaN
pymnt_plan	50000	7	Fully Paid	35565	NaN	NaN	NaN	NaN	NaN	NaN	NaN
desc	50000	2	n	49997	NaN	NaN	NaN	NaN	NaN	NaN	NaN
purpose	31004	30822	01/14/13	6	NaN	NaN	NaN	NaN	NaN	NaN	NaN
title	50000	13	rt consolidat	29850	NaN	NaN	NaN	NaN	NaN	NaN	NaN
zip_code	49998	15531	rt consolidat	5507	NaN	NaN	NaN	NaN	NaN	NaN	NaN
addr_state	50000	799	112xx	646	NaN	NaN	NaN	NaN	NaN	NaN	NaN
dti	50000	46	CA	8457	NaN	NaN	NaN	NaN	NaN	NaN	NaN
delinq_2yrs	50000	NaN	NaN	NaN	17.37308	7.765797	0	11.51	17.16	23.05	34.99
earliest_cr_line	50000	NaN	NaN	NaN	0.22444	0.671383	0	0	0	0	18
inq_last_6mths	50000	544	Oct-00	473	NaN	NaN	NaN	NaN	NaN	NaN	NaN
mths_since_last_del	50000	NaN	NaN	NaN	0.83888	1.02099	0	0	1	1	8
mths_since_last_rec	21874	NaN	NaN	NaN	36.0806	21.49374	0	18	33	52	152
open_acc	2532	NaN	NaN	NaN	87.70498	24.7171	2	76	93	106	119
pub_rec	50000	NaN	NaN	NaN	11.00528	4.54279	0	8	10	14	53
revol_bal	50000	NaN	NaN	NaN	0.05648	0.265954	0	0	0	0	8
revol_util	50000	NaN	NaN	NaN	1.60E+04	1.84E+04	0.00E+00	7.10E+03	1.24E+04	2.05E+04	1.74E+06
total_acc	49969	1013	0%	169	NaN	NaN	NaN	NaN	NaN	NaN	NaN
initial_list_status	50000	NaN	NaN	NaN	24.31492	11.02761	2	16	23	31	99
out_prncp	50000	2	f	43189	NaN	NaN	NaN	NaN	NaN	NaN	NaN
out_prncp_inv	50000	NaN	NaN	NaN	843.2033	2949.737	0	0	0	0	20921.13
total_pymnt	50000	NaN	NaN	NaN	842.1113	2946.173	0	0	0	0	20921.13
total_pymnt_inv	50000	NaN	NaN	NaN	14827.54	9489.452	0	7613.995	12858.3	20050.63	57777.58
total_rec_prncp	50000	NaN	NaN	NaN	14807.63	9478.955	0	7601.268	12842.03	20024.1	57777.58
total_rec_int	50000	NaN	NaN	NaN	11611.06	7591.581	0	6000	10000	15478.97	35000.01
total_rec_late_fee	50000	NaN	NaN	NaN	3071.491	3157.742	0	1057.515	2047.195	3736.968	22777.58
recoveries	50000	NaN	NaN	NaN	0.841868	5.933666	0	0	0	0	286.7476
collection_recovery	50000	NaN	NaN	NaN	144.149	699.7081	0	0	0	0	33520.27
last_pymnt_d	50000	NaN	NaN	NaN	10.6628	83.73578	0	0	0	0	3896.236
last_pymnt_amnt	49957	43	Dec-15	8867	NaN	NaN	NaN	NaN	NaN	NaN	NaN
next_pymnt_d	50000	NaN	NaN	NaN	3569.007	5529.438	0	353.14	723.575	4675.928	35683.2
last_credit_pull_d	7136	2	Jan-16	5907	NaN	NaN	NaN	NaN	NaN	NaN	NaN
collections_12_mths	50000	44	Dec-15	24668	NaN	NaN	NaN	NaN	NaN	NaN	NaN
mths_since_last_ma	50000	NaN	NaN	NaN	0.00114	0.036038	0	0	0	0	2
policy_code	7120	NaN	NaN	NaN	42.30787	20.90945	0	25	40	59	152
application_type	50000	NaN	NaN	NaN	1	0	1	1	1	1	1
acc_now_delinq	50000	1	INDIVIDUAL	50000	NaN	NaN	NaN	NaN	NaN	NaN	NaN
tot_coll_amt	50000	NaN	NaN	NaN	0.00082	0.034342	0	0	0	0	4
tot_cur_bal	35382	NaN	NaN	NaN	51.9957	659.7389	0	0	0	0	55009
total_credit_rv	35382	NaN	NaN	NaN	1.34E+05	1.57E+05	0.00E+00	2.63E+04	7.21E+04	2.02E+05	8.00E+06
loan_is_bad	35382	NaN	NaN	NaN	2.93E+04	2.90E+04	0.00E+00	1.40E+04	2.28E+04	3.66E+04	2.01E+06

Table 7 Descriptive Summary Statistics for Original Imbalanced Data

Appendix 2: Top 10 features for models using RFECV:

Top 10 Features Selected by RFECV for LR:

Feature_Name	RFECV_Ranking
recoveries	1
collection_recovery_fee	1
zip_code_232xx	1
zip_code_354xx	1
earliest_cr_line_Apr-1981	1
earliest_cr_line_Feb-1993	1
earliest_cr_line_Nov-2005	1
last_credit_pull_d_Dec-2015	1
last_credit_pull_d_Nov-2015	1
last_credit_pull_d_Oct-2015	1

Table 8 Top 10 Features of LR

Top 10 Features Selected by RFECV for DT:

Feature_Name	RFECV_Ranking
member_id	1
emp_title_Lebanon's Cafe	1
emp_title_Ledcor Technical Services	1
emp_title_Luminant mining co	1
emp_title_MERIT Property Management, Inc.	1
emp_title_MMC Group	1
emp_title_MTA Police Department	1
emp_title_Macy's Herald Square	1
emp_title_Macy's Inc.	1
emp_title_Mark Twain Dignity Health	1

Table 9 Top 10 Features of DT

Top 10 Features Selected by RFECV for RF:

Feature_Name	RFECV_Ranking
member_id	1
int_rate	1
installment	1
annual_inc	1
dti	1
revol_util	1
recoveries	1
collection_recovery_fee	1
debt_to_income_ratio	1
installment_to_income	1

Table 10 Top 10 Features of RF

Top 10 Features Selected by RFECV for SVC:

Feature_Name	RFECV_Ranking
recoveries	1
collection_recovery_fee	1
emp_title_New Horizon Farms	1
emp_title_New York City Board of Education	1
emp_title_NewYork-Presbyterian Hospital	1
emp_title_Nextran Truck Center	1
emp_title_Nielsen Corporation	1
emp_title_Parkview Lagrange Hospital	1
zip_code_354xx	1
zip_code_616xx	1

Table 11 Top 10 Features of SVC

Top 10 Features Selected by RFECV for XGBoost:

Feature_Name	RFECV_Ranking
term	1
int_rate	1
annual_inc	1
out_prncp	1
recoveries	1
installment_to_income	1
sub_grade_D4	1
home_ownership_OWN	1
purpose_debt_consolidation	1
last_credit_pull_d_Dec-2015	1

Table 12 Top 10 Features of XGBoost

Top 10 Features Selected by RFECV for LightGBM:

Feature_Name	RFECV_Ranking
member_id	1
emp_title_Barnes and Noble	1
emp_title_Barnes-Jewish Hopsital	1
emp_title_Barrell Plumbing	1
emp_title_Barrett Moving and Storage A Suddath Com	1
emp_title_Barry S. Slatt. Com	1
emp_title_Bartell Drugs	1
emp_title_Basham Law Group	1
emp_title_Bastion Technologies	1
emp_title_Barnes Jewish	1

Table 13 Top 10 Features of LightGBM

Appendix 3: Classification Reports

3.1 Models on Original Imbalanced dataset:

LR on Original dataset Results:

	precision	recall	f1-score	support
0	0.92	1.00	0.96	2539
1	1.00	0.54	0.70	461
accuracy			0.93	3000
macro avg	0.96	0.77	0.83	3000
weighted avg	0.93	0.93	0.92	3000

ROC AUC: 0.8923637245948026

DT on Original dataset Results:

	precision	recall	f1-score	support
0	0.94	0.97	0.96	2539
1	0.80	0.69	0.74	461
accuracy			0.93	3000
macro avg	0.87	0.83	0.85	3000
weighted avg	0.92	0.93	0.92	3000

ROC AUC: 0.8269789547697993

RF on Original dataset Results:

	precision	recall	f1-score	support
0	0.94	1.00	0.97	2539
1	1.00	0.65	0.79	461
accuracy			0.95	3000
macro avg	0.97	0.82	0.88	3000
weighted avg	0.95	0.95	0.94	3000

ROC AUC: 0.858316552454166

SVC on Original dataset Results:

	precision	recall	f1-score	support
0	0.93	1.00	0.96	2539
1	1.00	0.58	0.73	461
accuracy			0.94	3000
macro avg	0.96	0.79	0.85	3000
weighted avg	0.94	0.94	0.93	3000

ROC AUC: 0.8241976148226495

XGB on Original dataset Results:

	precision	recall	f1-score	support
0	0.94	0.99	0.97	2539
1	0.95	0.66	0.78	461
accuracy			0.94	3000
macro avg	0.95	0.83	0.87	3000
weighted avg	0.94	0.94	0.94	3000

ROC AUC: 0.9174491810617704

LGBM on Original dataset Results:

	precision	recall	f1-score	support
0	0.94	0.98	0.96	2539
1	0.85	0.67	0.75	461
accuracy			0.93	3000
macro avg	0.90	0.83	0.86	3000
weighted avg	0.93	0.93	0.93	3000

ROC AUC: 0.9202206959714783

3.2 Models under Random Undersampling:

LR on Undersampling Results:

	precision	recall	f1-score	support
0	0.93	0.99	0.96	2539
1	0.95	0.57	0.71	461
accuracy			0.93	3000
macro avg	0.94	0.78	0.83	3000
weighted avg	0.93	0.93	0.92	3000

ROC AUC: 0.8881227258242138

DT on Undersampling Results:

	precision	recall	f1-score	support
0	0.96	0.81	0.88	2539
1	0.44	0.81	0.57	461
accuracy			0.81	3000
macro avg	0.70	0.81	0.73	3000
weighted avg	0.88	0.81	0.83	3000

ROC AUC: 0.8135776036989985

RF on Undersampling Results:

	precision	recall	f1-score	support
0	0.94	0.95	0.94	2539
1	0.70	0.68	0.69	461
accuracy			0.91	3000
macro avg	0.82	0.81	0.82	3000
weighted avg	0.91	0.91	0.91	3000

ROC AUC: 0.8676007856612549

SVC on Undersampling Results:

	precision	recall	f1-score	support
0	0.93	1.00	0.96	2539
1	1.00	0.58	0.73	461
accuracy			0.93	3000
macro avg	0.96	0.79	0.85	3000
weighted avg	0.94	0.93	0.93	3000

ROC AUC: 0.8241847995564209

XGBoost on Undersampling Results:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.95	0.87	0.91	2539
1	0.52	0.76	0.62	461
accuracy			0.85	3000
macro avg	0.74	0.82	0.76	3000
weighted avg	0.89	0.85	0.87	3000

ROC AUC: 0.9095519868361585

LightGBM on Undersampling Results:

	precision	recall	f1-score	support
0	0.96	0.88	0.92	2539
1	0.55	0.79	0.65	461
accuracy			0.87	3000
macro avg	0.75	0.84	0.78	3000
weighted avg	0.90	0.87	0.88	3000

ROC AUC: 0.9237107201410704

3.3 Models under SMOTE:

LR on SMOTE Results:

	precision	recall	f1-score	support
0	0.93	0.99	0.96	2539
1	0.95	0.62	0.75	461
accuracy			0.94	3000
macro avg	0.94	0.81	0.86	3000
weighted avg	0.94	0.94	0.93	3000

ROC AUC: 0.8926798344951083

DT on SMOTE Results:

	precision	recall	f1-score	support
0	0.94	0.96	0.95	2539
1	0.77	0.67	0.72	461
accuracy			0.92	3000
macro avg	0.86	0.82	0.83	3000
weighted avg	0.91	0.92	0.92	3000

ROC AUC: 0.8159390300893908

RF on SMOTE Results:

	precision	recall	f1-score	support
0	0.94	0.99	0.96	2539
1	0.92	0.65	0.76	461
accuracy			0.94	3000
macro avg	0.93	0.82	0.86	3000
weighted avg	0.94	0.94	0.93	3000

ROC AUC: 0.8655943421453952

SVC on SMOTE Results:

	precision	recall	f1-score	support
0	0.93	1.00	0.96	2539
1	1.00	0.58	0.74	461
accuracy			0.94	3000
macro avg	0.96	0.79	0.85	3000
weighted avg	0.94	0.94	0.93	3000

ROC AUC: 0.8241967604715676

XGBoost on SMOTE Results:

	precision	recall	f1-score	support
0	0.94	0.98	0.96	2539
1	0.83	0.67	0.74	461
accuracy			0.93	3000
macro avg	0.89	0.82	0.85	3000
weighted avg	0.92	0.93	0.92	3000

ROC AUC: 0.914802401410021

LightGBM on SMOTE Results:

	precision	recall	f1-score	support
0	0.94	0.99	0.97	2539
1	0.96	0.66	0.78	461
accuracy			0.94	3000
macro avg	0.95	0.83	0.87	3000
weighted avg	0.94	0.94	0.94	3000

ROC AUC: 0.9238978230280082

3.4 Models under Threshold Optimization:

Optimal threshold: 0.13

LR on Threshold Optimization Results:

	precision	recall	f1-score	support
0	0.94	0.99	0.97	2539
1	0.95	0.64	0.77	461
accuracy			0.94	3000
macro avg	0.94	0.82	0.87	3000
weighted avg	0.94	0.94	0.93	3000

ROC AUC: 0.8923637245948026

Optimal threshold: 0.00

DT on Threshold Optimization Results:

	precision	recall	f1-score	support
0	0.94	0.97	0.96	2539
1	0.80	0.69	0.74	461
accuracy			0.93	3000
macro avg	0.87	0.83	0.85	3000
weighted avg	0.92	0.93	0.92	3000

ROC AUC: 0.8269789547697993

Optimal threshold: 0.41

RF on Threshold Optimization Results:

	precision	recall	f1-score	support
0	0.94	1.00	0.97	2539
1	1.00	0.65	0.79	461
accuracy			0.95	3000
macro avg	0.97	0.82	0.88	3000
weighted avg	0.95	0.95	0.94	3000

ROC AUC: 0.858316552454166

SVC on Optimal threshold: 0.07

Threshold Optimization Results:

	precision	recall	f1-score	support
0	0.94	1.00	0.97	2539
1	1.00	0.65	0.79	461
accuracy			0.95	3000
macro avg	0.97	0.82	0.88	3000

weighted avg	0.95	0.95	0.94	3000
--------------	------	------	------	------

ROC AUC: 0.8241976148226495

Optimal threshold: 0.61

XGB on Threshold Optimization Results:

	precision	recall	f1-score	support
0	0.94	1.00	0.97	2539
1	0.99	0.66	0.79	461
accuracy			0.95	3000
macro avg	0.97	0.83	0.88	3000
weighted avg	0.95	0.95	0.94	3000

ROC AUC: 0.9174491810617704

Optimal threshold: 0.60

Threshold Optimization Results:

	precision	recall	f1-score	support
0	0.94	1.00	0.97	2539
1	1.00	0.65	0.79	461
accuracy			0.95	3000
macro avg	0.97	0.83	0.88	3000
weighted avg	0.95	0.95	0.94	3000

ROC AUC: 0.9264420805499287

3.5 Models under Cost Sensitive Analysis:

Class weights: {0: 0.5909167651527942, 1: 3.2497678737233056}

LR on Cost-Sensitive Learning Results:

	precision	recall	f1-score	support
0	0.93	0.99	0.96	2539
1	0.95	0.61	0.74	461
accuracy			0.94	3000
macro avg	0.94	0.80	0.85	3000
weighted avg	0.94	0.94	0.93	3000

ROC AUC: 0.8926319908345215

DT on Cost-Sensitive Learning Results:

	precision	recall	f1-score	support
0	0.95	0.95	0.95	2539
1	0.72	0.72	0.72	461
accuracy			0.91	3000
macro avg	0.83	0.83	0.83	3000
weighted avg	0.91	0.91	0.91	3000

ROC AUC: 0.832513868253937

RF on Cost-Sensitive Learning Results:

	precision	recall	f1-score	support
0	0.94	1.00	0.97	2539
1	1.00	0.65	0.79	461
accuracy			0.95	3000
macro avg	0.97	0.82	0.88	3000
weighted avg	0.95	0.95	0.94	3000

ROC AUC: 0.8641487801148078

SVC on Cost-Sensitive Learning Results:

	precision	recall	f1-score	support
0	0.93	1.00	0.96	2539
1	1.00	0.58	0.73	461
accuracy			0.94	3000
macro avg	0.96	0.79	0.85	3000
weighted avg	0.94	0.94	0.93	3000

ROC AUC: 0.8241941974183219

XGBoost on Cost-Sensitive Learning Results:

	precision	recall	f1-score	support
0	0.94	0.99	0.97	2539
1	0.95	0.66	0.78	461
accuracy			0.94	3000
macro avg	0.95	0.83	0.87	3000
weighted avg	0.94	0.94	0.94	3000

ROC AUC: 0.9174491810617704

Cost-Sensitive Learning Results:

	precision	recall	f1-score	support
0	0.95	0.98	0.96	2539
1	0.84	0.70	0.76	461
accuracy			0.93	3000
macro avg	0.89	0.84	0.86	3000
weighted avg	0.93	0.93	0.93	3000

ROC AUC: 0.922697459757928

Appendix 4: ROC-AUC Curves:

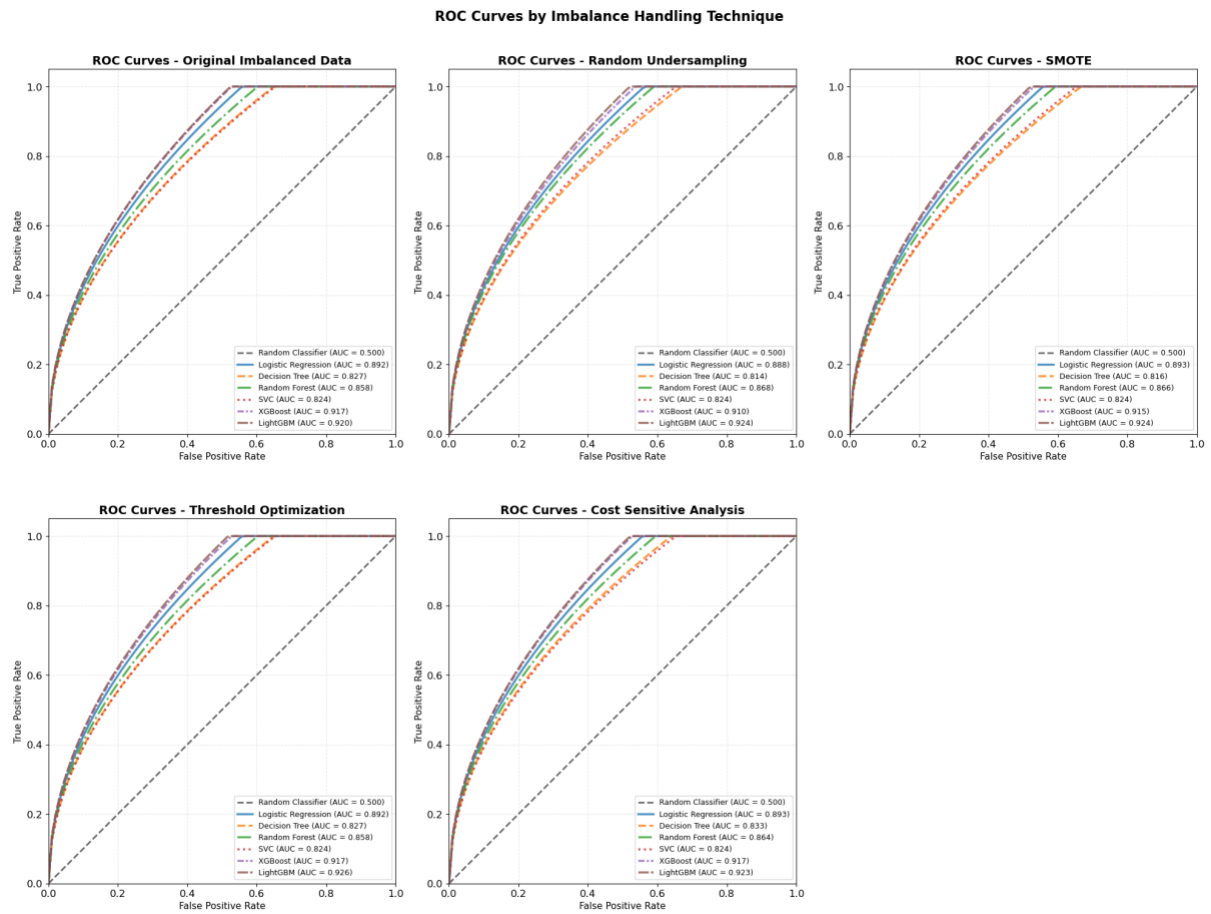


Figure 5 Area Under Curves for each Class Imbalance Handling Technique

Appendix 5: Shap figures:




Figure 6 Feature Importance for XGBoost on SMOTE




Figure 7 Feature Importance for cost Sensitive XGBoost




Figure 8 Feature Importance for LightGBM on smote




Figure 9 Feature Importance for Cost Sensitive LightGBM

Appendix 6: Script of Python Code:

```
#!/usr/bin/env python
# coding: utf-8
```

```
# In[1]:
```

```
# Import Libraries
import pandas as pd
import numpy as np
```

```
# In[2]:
```

```
df = pd.read_csv('early_2012_2013_loan_sample_with_outcome 1.csv')
```

```
# In[3]:
```

```
df.columns
```

```
# In[4]:
```

```
default_rate_overall = df['loan_is_bad'].mean() * 100
print(default_rate_overall)
```

```
# In[5]:
```

```
df1 = df.sample(n=10000, random_state=42).reset_index(drop=True)
```

```
# In[6]:
```

```
default_rate_sample = df1['loan_is_bad'].mean() * 100
print(default_rate_sample)
```

```
# In[7]:
```

```
df1.to_csv('/Users/manishathakur/Desktop/Warwick/4. Dissertation/2. Data and
Code/Sampled_10000_data.csv')
```

```
# # Data Cleaning
```

```
# In[8]:
```

```
df1.head()
```

```
# In[9]:
```

```
df1.shape
```

```
# In[10]:
```

```
df1.columns
```

```
# In[12]:
```

```
df1.isna().sum()[lambda x: x > 0]
```

```
# In[13]:
```

```
pd.crosstab(df1["purpose"], df1["loan_is_bad"], margins = True)
```

```
# In[14]:
```

```
loan_home_emplength = pd.crosstab(df1["loan_is_bad"], df1["home_ownership"],  
                                  values=df1["emp_length"], aggfunc="max")  
loan_home_emplength = loan_home_emplength.applymap(lambda x: "10 or more" if x == 10  
else x)  
print(loan_home_emplength)
```

```
# In[15]:
```

```
print(df1.columns[df1.isnull().any()]) # emp_length, revol_util has null
```

```
# In[16]:
```

```
print(df1["emp_length"].isnull().sum()) #357 null
```

```
# In[17]:
```

```
import matplotlib.pyplot as plt
```

```
# In[18]:
```

```
df2 = df1.copy()
```

```
# In[19]:
```

```
n, bins, patches = plt.hist(df2["emp_length"], bins='auto', color='blue')
plt.xlabel("Person Employment Length")
plt.show()
```

```
# In[20]:
```

```
df2["emp_length"].value_counts().sort_index()
```

```
# In[21]:
```

```
df2[df2["emp_length"].isna()][["loan_is_bad"].value_counts() # 78 True, 279 False
```

```
# In[22]:
```

```
df2["loan_is_bad"].value_counts()
```

```
# In[23]:
```

```
df3 = df2.copy()
```

```
# In[24]:
```

```
df3["emp_length"] = df3["emp_length"].fillna(df3["emp_length"].median()) #since null values are
high and proportion of bad loans are higher too, replace it with median
```

```
# In[25]:
```

```
df3[df3["emp_length"].isna()][["loan_is_bad"].value_counts()]
```

```
# In[26]:
```

```
df3[df3['revol_util'].isna()][['loan_is_bad','id']].value_counts()
```

```
# In[27]:
```

```
df3["revol_util"].value_counts().sort_index()
```

```
# In[28]:
```

```
# Convert revol_util to numeric  
df3['revol_util'] = df3['revol_util'].str.replace('%', '').astype(float)
```

```
# In[29]:
```

```
df3['revol_util'] = df3['revol_util'].fillna(df3['revol_util'].median())
```

```
# In[30]:
```

```
df3[df3['revol_util'].isna()][['loan_is_bad','id']].value_counts()
```

```
# In[31]:
```

```
df3.head()
```

```
# In[32]:
```

```
plt.boxplot(df3[['int_rate','annual_inc','loan_amnt','installment']])
```

```
# In[33]:
```

```
cols = ['int_rate', 'annual_inc', 'loan_amnt', 'installment']
```

```
for col in cols:  
    plt.figure(figsize=(6, 4))  
    plt.boxplot(df3[col].dropna())  
    plt.title(f'Boxplot of {col}')  
    plt.ylabel(col)  
    plt.grid(True)  
    plt.show()
```

```
# # Feature Engineering
```

```
# In[34]:
```

```
df3['debt_to_income_ratio'] = df3['annual_inc'] / (df3['loan_amnt'] + 1)
df3['installment_to_income'] = df3['installment'] / (df3['annual_inc'] + 1)
```

```
# In[35]:
```

```
print(df3[['debt_to_income_ratio','installment_to_income']].head())
```

```
# In[36]:
```

```
df3['loan_is_bad_binary'] = np.where(df3['loan_is_bad'] == True, 1, 0)
```

```
# In[37]:
```

```
df3['loan_is_bad_binary'].value_counts()
```

```
# In[38]:
```

```
df3['credit_utilisation_ratio'] = (df3['revol_bal'] / df3['total_credit_rv']) * 100
```

```
# In[39]:
```

```
df3['credit_utilisation_ratio'] = df3['credit_utilisation_ratio'].fillna(df3['credit_utilisation_ratio'].median())
```

```
# In[40]:
```

```
from sklearn.model_selection import train_test_split
```

```
# In[41]:
```

```
# Separate features and target
X = df3.drop(columns=['loan_is_bad','loan_is_bad_binary', 'id', 'loan_status', 'desc', 'title',
'total_pymnt','total_pymnt_inv', 'total_rec_int', 'total_rec_late_fee','total_rec_prncp',
'last_pymnt_d','last_pymnt_amnt', 'next_pymnt_d'])
y = df3['loan_is_bad_binary']
```



```
# Split data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42,
stratify=y)
print(X_train.shape, y_train.shape, X_test.shape, y_test.shape)
```

```
# In[42]:
```

```
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
```

```
# In[43]:
```

```
# Define numeric and categorical features
numeric_features = X.select_dtypes(include=['int64', 'float64']).columns
categorical_features = X.select_dtypes(include=['object']).columns
```

```
# Create preprocessing pipelines
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler())])
```

```
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('onehot', OneHotEncoder(handle_unknown='ignore'))])
```

```
# In[44]:
```

```
# 1. Rebuild transformers (verified correct from your debug output)
```

```
numeric_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', StandardScaler()) # No feature reduction!
])
```

```
categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='constant', fill_value='missing')),
    ('onehot', OneHotEncoder(handle_unknown='ignore', sparse_output=False)) # Note
sparse_output=False
])
```

```
# In[45]:
```

```
# Combine preprocessing steps
preprocessor = ColumnTransformer(
    transformers=[
```

```

        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)])

# Preprocess the data
X_train_preprocessed = preprocessor.fit_transform(X_train)
X_test_preprocessed = preprocessor.transform(X_test)

# Get feature names after one-hot encoding
cat_encoder = preprocessor.named_transformers_['cat'].named_steps['onehot']
cat_features = cat_encoder.get_feature_names_out(categorical_features)
all_features = np.concatenate([numeric_features, cat_features])

# Convert to DataFrame for better visualization
X_train_preprocessed_df = pd.DataFrame(X_train_preprocessed, columns=all_features)
X_test_preprocessed_df = pd.DataFrame(X_test_preprocessed, columns=all_features)

print("\nPreprocessed training data shape:", X_train_preprocessed.shape)
print("\nPreprocessed test data shape:", X_test_preprocessed.shape)

```

In[46]:

```

# 2. Create preprocessor with dense output
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)],
    remainder='drop',
    sparse_threshold=0
)

```

In[47]:

```

# Explicitly set sparse_threshold=0 in ColumnTransformer
preprocessor = ColumnTransformer(
    transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)],
    remainder='drop',
    sparse_threshold=0 # Force dense output
)

```

In[48]:

```

print(X_train.shape, X_test.shape, numeric_features.shape, categorical_features.shape)

```

In[49]:

```
print(X_train_preprocessed.shape, X_test_preprocessed.shape, y_train.shape)
```

```
# # Feature Selection
```

```
# In[50]:
```

```
from sklearn.feature_selection import RFECV
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import StratifiedKFold
import pandas as pd
import numpy as np
from sklearn.feature_selection import RFE
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from xgboost import XGBClassifier
from lightgbm import LGBMClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report
```

```
# In[51]:
```

```
from tqdm import tqdm
import time
```

```
# Logistic Regression
```

```
# In[52]:
```

```
# Logistic regression with RFECV
LRestimator = LogisticRegression(max_iter=1000, random_state=42)
```

```
# In[54]:
```

```
# Define RFECV with Logistic Regression
```

```
LRrfecv = RFECV(
    estimator=LRestimator,      # Your LogisticRegression(max_iter=1000, solver="liblinear"
    or "lbfgs")
    step=0.1,                  # Drop 10% features per step
    cv=StratifiedKFold(3),     # 3-fold CV for speed
    scoring='f1',
    min_features_to_select=10,
    n_jobs=-1,
    verbose=0
)
```

```
print("Starting RFECV for Logistic Regression...")
```

```

start_time = time.time()
with tqdm(total=X_train_preprocessed.shape[1], desc="RFECV Iterations (LR)") as pbar:
    LRfecv.fit(X_train_preprocessed, y_train)

    # Progress simulation
    n_features_remaining = X_train_preprocessed.shape[1]
    while n_features_remaining > LRfecv.min_features_to_select:
        step_size = max(1, int(n_features_remaining * 0.1))
        n_features_remaining -= step_size
        pbar.update(step_size)
        print(f"Features remaining: {n_features_remaining}")

end_time = time.time()
print(f"\nRFECV (Logistic Regression) completed in {end_time - start_time:.2f} seconds")
print(f"Optimal number of features: {LRfecv.n_features_}")

# In[55]:

print("Optimal number of features:", LRfecv.n_features_)
print("Selected features:", X_train_preprocessed_df.columns[LRfecv.support_])

# In[56]:

LRselected_features_mask = LRfecv.support_
LR_X_train_selected = X_train_preprocessed[:, LRselected_features_mask]
LR_X_test_selected = X_test_preprocessed[:, LRselected_features_mask]

# In[57]:

selected_features = X_train_preprocessed_df.columns[LRfecv.support_]
feature_ranking = LRfecv.ranking_[LRfecv.support_]

# Create a table with feature names and their ranking
features_table = pd.DataFrame({
    'Feature_Name': selected_features,
    'RFECV_Ranking': feature_ranking
})

# Sort by ranking (lower rank = better)
features_table = features_table.sort_values('RFECV_Ranking').reset_index(drop=True)

# Display top 10 features
print("Top 10 Features Selected by RFECV:")
print(features_table.head(10).to_string(index=False))

# Decision Tree

```

```
# In[58]:
```

```
# Decision Tree with RFECV
DTestimator = DecisionTreeClassifier(random_state=42)
```

```
# In[60]:
```

```
# Define RFECV with Decision Tree
DTrfecv = RFECV(
    estimator=DTestimator,      # Your DecisionTreeClassifier()
    step=0.1,                   # Drop 10% features each step
    cv=StratifiedKFold(3),      # Faster CV (3-fold instead of 5)
    scoring='f1',               # Optimize F1
    min_features_to_select=10,
    n_jobs=-1,
    verbose=0
)

print("Starting RFECV for Decision Tree...")

start_time = time.time()
with tqdm(total=X_train_preprocessed.shape[1], desc="RFECV Iterations (DT)") as pbar:
    DTrfecv.fit(X_train_preprocessed, y_train)

    # Progress simulation
    n_features_remaining = X_train_preprocessed.shape[1]
    while n_features_remaining > DTrfecv.min_features_to_select:
        step_size = max(1, int(n_features_remaining * 0.1))
        n_features_remaining -= step_size
        pbar.update(step_size)
        print(f"Features remaining: {n_features_remaining}")

end_time = time.time()
print(f"\nRFECV (Decision Tree) completed in {end_time - start_time:.2f} seconds")
print(f"Optimal number of features: {DTrfecv.n_features_}")
```

```
# In[61]:
```

```
print("Optimal number of features:", DTrfecv.n_features_)
print("Selected features:", X_train_preprocessed_df.columns[DTrfecv.support_])
```

```
# In[62]:
```

```
DTselected_features_mask = DTrfecv.support_
DT_X_train_selected = X_train_preprocessed[:, DTselected_features_mask]
DT_X_test_selected = X_test_preprocessed[:, DTselected_features_mask]
```

In[63]:

```
selected_features = X_train_preprocessed_df.columns[DTrfecv.support_]
feature_ranking = DTrfecv.ranking_[DTrfecv.support_]

# Create a table with feature names and their ranking
features_table = pd.DataFrame({
    'Feature_Name': selected_features,
    'RFECV_Ranking': feature_ranking
})

# Sort by ranking (lower rank = better)
features_table = features_table.sort_values('RFECV_Ranking').reset_index(drop=True)

# Display top 10 features
print("Top 10 Features Selected by RFECV:")
print(features_table.head(10).to_string(index=False))
```

Random Forest

In[64]:

```
# Random Forest with RFECV
RFestimator = RandomForestClassifier(random_state=42)
```

In[66]:

```
RFrfevcv = RFECV(
    estimator=RFestimator,      # Your RandomForestClassifier()
    step=0.1,                   # Drop 10% features at a time
    cv=StratifiedKFold(3),      # 3-fold CV instead of 5
    scoring='f1',               # Optimize for F1 score
    min_features_to_select=10,
    n_jobs=-1,
    verbose=0                   # silence sklearn logs
)

print("Starting RFECV for Random Forest...")

start_time = time.time()
with tqdm(total=X_train_preprocessed.shape[1], desc="RFECV Iterations (RF)") as pbar:
    RFrfevcv.fit(X_train_preprocessed, y_train)

# Simulated iteration tracker
n_features_remaining = X_train_preprocessed.shape[1]
while n_features_remaining > RFrfevcv.min_features_to_select:
    step_size = max(1, int(n_features_remaining * 0.1))
    n_features_remaining -= step_size
    pbar.update(step_size)
    print(f"Features remaining: {n_features_remaining}")
```

```

end_time = time.time()
print(f"\nRFECV (Random Forest) completed in {end_time - start_time:.2f} seconds")
print(f"Optimal number of features: {RFfecv.n_features_}")

```

In[67]:

```

print("Optimal number of features:", RFfecv.n_features_)
print("Selected features:", X_train_preprocessed_df.columns[RFfecv.support_])

```

In[68]:

```

RFselected_features_mask = RFfecv.support_
RF_X_train_selected = X_train_preprocessed[:, RFselected_features_mask]
RF_X_test_selected = X_test_preprocessed[:, RFselected_features_mask]

```

In[69]:

```

selected_features = X_train_preprocessed_df.columns[RFfecv.support_]
feature_ranking = RFfecv.ranking_[RFfecv.support_]

# Create a table with feature names and their ranking
features_table = pd.DataFrame({
    'Feature_Name': selected_features,
    'RFECV_Ranking': feature_ranking
})

# Sort by ranking (lower rank = better)
features_table = features_table.sort_values('RFECV_Ranking').reset_index(drop=True)

# Display top 10 features
print("Top 10 Features Selected by RFECV:")
print(features_table.head(10).to_string(index=False))

```

SVM

In[70]:

```

# SVM with RFECV
SVCestimator = SVC(kernel='linear', random_state=42)

```

In[72]:

```

SVCrfecv = RFECV(
    estimator=SVCestimator,      # ideally use SVC(kernel='linear')

```

```

    step=0.1,                # drop 10% features per step
    cv=StratifiedKfold(3),    # 3-fold CV instead of 5
    scoring='f1',
    min_features_to_select=10,
    n_jobs=-1,
    verbose=0                # silence sklearn logs
)

print("Starting RFECV for SVC...")

start_time = time.time()
with tqdm(total=X_train_preprocessed.shape[1], desc="RFECV Iterations (SVC)") as pbar:
    SVCrfecv.fit(X_train_preprocessed, y_train)

    # Simulated iteration tracker
    n_features_remaining = X_train_preprocessed.shape[1]
    while n_features_remaining > SVCrfecv.min_features_to_select:
        step_size = max(1, int(n_features_remaining * 0.1))
        n_features_remaining -= step_size
        pbar.update(step_size)
        print(f"Features remaining: {n_features_remaining}")

end_time = time.time()
print(f"\nRFECV (SVC) completed in {end_time - start_time:.2f} seconds")
print(f"Optimal number of features: {SVCrfecv.n_features_}")

# In[73]:

print("Optimal number of features:", SVCrfecv.n_features_)
print("Selected features:", X_train_preprocessed_df.columns[SVCrfecv.support_])

# In[74]:

SVCselected_features_mask = SVCrfecv.support_
SVC_X_train_selected = X_train_preprocessed[:, SVCselected_features_mask]
SVC_X_test_selected = X_test_preprocessed[:, SVCselected_features_mask]

# In[75]:

selected_features = X_train_preprocessed_df.columns[SVCrfecv.support_]
feature_ranking = SVCrfecv.ranking_[SVCrfecv.support_]

# Create a table with feature names and their ranking
features_table = pd.DataFrame({
    'Feature_Name': selected_features,
    'RFECV_Ranking': feature_ranking
})

# Sort by ranking (lower rank = better)

```



```

features_table = features_table.sort_values('RFECV_Ranking').reset_index(drop=True)

# Display top 10 features
print("Top 10 Features Selected by RFECV:")
print(features_table.head(10).to_string(index=False))

# XGBoost

# In[76]:

# XGB with RFECV
XGBestimator = XGBClassifier(kernel='linear',random_state=42)

# In[77]:

XGBrfecv = RFECV(
    estimator=XGBClassifier(
        n_estimators=100,
        max_depth=3,
        learning_rate=0.1,
        tree_method="hist",
        n_jobs=-1,
        random_state=42
    ),
    step=0.1,          # remove 10% of features each iteration
    cv=StratifiedKFold(3), # fewer folds
    scoring='f1',
    min_features_to_select=10,
    n_jobs=-1,
    verbose=1
)

print("Starting RFECV...")

# Wrap fit with tqdm progress tracking
start_time = time.time()
with tqdm(total=X_train_preprocessed.shape[1], desc="RFECV Iterations") as pbar:
    XGBrfecv.fit(X_train_preprocessed, y_train)
    # tqdm can't directly track RFECV steps, so we simulate by number of features dropped
    # For logging iteration counts:
    n_features_remaining = X_train_preprocessed.shape[1]
    while n_features_remaining > XGBrfecv.min_features_to_select:
        step_size = max(1, int(n_features_remaining * 0.1))
        n_features_remaining -= step_size
        pbar.update(step_size)
        print(f"Features remaining: {n_features_remaining}")

end_time = time.time()
print(f"\nRFECV completed in {end_time - start_time:.2f} seconds")
print(f"Optimal number of features: {XGBrfecv.n_features_}")

```

In[78]:

```
print("Optimal number of features:", XGBrfecv.n_features_)
print("Selected features:", X_train_preprocessed_df.columns[XGBrfecv.support_])
```

In[79]:

```
XGBselected_features_mask = XGBrfecv.support_
XGB_X_train_selected = X_train_preprocessed[:, XGBselected_features_mask]
XGB_X_test_selected = X_test_preprocessed[:, XGBselected_features_mask]
```

In[80]:

```
selected_features = X_train_preprocessed_df.columns[XGBrfecv.support_]
feature_ranking = XGBrfecv.ranking_[XGBrfecv.support_]

# Create a table with feature names and their ranking
features_table = pd.DataFrame({
    'Feature_Name': selected_features,
    'RFECV_Ranking': feature_ranking
})

# Sort by ranking (lower rank = better)
features_table = features_table.sort_values('RFECV_Ranking').reset_index(drop=True)

# Display top 10 features
print("Top 10 Features Selected by RFECV:")
print(features_table.head(10).to_string(index=False))
```

LightGBM

In[81]:

```
# LightGBM with RFECV
LGBMestimator = LGBMClassifier(random_state=42)
```

In[82]:

```
# Define RFECV with lighter config
LGBMrfecv = RFECV(
    estimator=LGBMestimator,
    step=0.1,                # drop 10% features each step
    cv=StratifiedKFold(3),   # reduced to 3-fold for speed
    scoring='f1',
    min_features_to_select=10,
```

```

    n_jobs=-1,
    verbose=0                # turn off sklearn's own logs
)

print("Starting RFECV for LightGBM...")

start_time = time.time()
with tqdm(total=X_train_preprocessed.shape[1], desc="RFECV Iterations (LightGBM)") as
pbar:
    LGBMrfecv.fit(X_train_preprocessed, y_train)

    # Simulated iteration tracking
    n_features_remaining = X_train_preprocessed.shape[1]
    while n_features_remaining > LGBMrfecv.min_features_to_select:
        step_size = max(1, int(n_features_remaining * 0.1))
        n_features_remaining -= step_size
        pbar.update(step_size)
        print(f"Features remaining: {n_features_remaining}")

end_time = time.time()
print(f"\nRFECV (LightGBM) completed in {end_time - start_time:.2f} seconds")
print(f"Optimal number of features: {LGBMrfecv.n_features_}")

# In[83]:

print("Optimal number of features:", LGBMrfecv.n_features_)
print("Selected features:", X_train_preprocessed_df.columns[LGBMrfecv.support_])

# In[84]:

LGBMselected_features_mask = LGBMrfecv.support_
LGBM_X_train_selected = X_train_preprocessed[:, LGBMselected_features_mask]
LGBM_X_test_selected = X_test_preprocessed[:, LGBMselected_features_mask]

# In[85]:

selected_features = X_train_preprocessed_df.columns[LGBMrfecv.support_]
feature_ranking = LGBMrfecv.ranking_[LGBMrfecv.support_]

# Create a table with feature names and their ranking
features_table = pd.DataFrame({
    'Feature_Name': selected_features,
    'RFECV_Ranking': feature_ranking
})

# Sort by ranking (lower rank = better)
features_table = features_table.sort_values('RFECV_Ranking').reset_index(drop=True)

# Display top 10 features

```

```
print("Top 10 Features Selected by RFECV:")
print(features_table.head(10).to_string(index=False))
```

Handling Class Imbalance

In[86]:

```
get_ipython().system('pip install imbalanced-learn')
```

In[87]:

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix, roc_auc_score, f1_score
from imblearn.under_sampling import RandomUnderSampler
from imblearn.over_sampling import SMOTE
from sklearn.model_selection import GridSearchCV
from sklearn.utils.class_weight import compute_class_weight
```

In[88]:

```
from sklearn.metrics import f1_score, roc_auc_score, recall_score, precision_score,
confusion_matrix
```

Original

In[89]:

Train LR on original data

```
lr_o = LogisticRegression(random_state=42)
lr_o.fit(LR_X_train_selected, y_train)
```

Evaluate

```
y_lr_o = lr_o.predict(LR_X_test_selected)
print("\nLR on Original dataset Results:")
print(classification_report(y_test, y_lr_o))
print("ROC AUC:", roc_auc_score(y_test, lr_o.predict_proba(LR_X_test_selected)[:, 1]))
```

In[90]:

Train DT on original data

```
dt_o = DecisionTreeClassifier(random_state=42)
dt_o.fit(DT_X_train_selected, y_train)
```

Evaluate

```
y_dt_o = dt_o.predict(DT_X_test_selected)
```

```
print("\nDT on Original dataset Results:")
print(classification_report(y_test, y_dt_o))
print("ROC AUC:", roc_auc_score(y_test, dt_o.predict_proba(DT_X_test_selected)[:, 1]))
```

In[91]:

```
# Train RF on original data
rf_o = RandomForestClassifier(random_state=42)
rf_o.fit(RF_X_train_selected, y_train)

# Evaluate
y_rf_o = rf_o.predict(RF_X_test_selected)
print("\nRF on Original dataset Results:")
print(classification_report(y_test, y_rf_o))
print("ROC AUC:", roc_auc_score(y_test, rf_o.predict_proba(RF_X_test_selected)[:, 1]))
```

In[92]:

```
# Train SCV on original data
svc_o = SVC(kernel='linear', random_state=42, probability=True)
svc_o.fit(SVC_X_train_selected, y_train)

# Evaluate
y_svc_o = svc_o.predict(SVC_X_test_selected)
print("\nSVC on Original dataset Results:")
print(classification_report(y_test, y_svc_o))
print("ROC AUC:", roc_auc_score(y_test, svc_o.predict_proba(SVC_X_test_selected)[:, 1]))
```

In[93]:

```
# Train XGB on original data
xgb_o = XGBClassifier(random_state=42)
xgb_o.fit(XGB_X_train_selected, y_train)

# Evaluate
y_xgb_o = xgb_o.predict(XGB_X_test_selected)
print("\nXGB on Original dataset Results:")
print(classification_report(y_test, y_xgb_o))
print("ROC AUC:", roc_auc_score(y_test, xgb_o.predict_proba(XGB_X_test_selected)[:, 1]))
```

In[94]:

```
# Train LightGBM on original data
lgbm_o = LGBMClassifier(boosting_type='gbdt', objective='binary', n_estimators=500,
    learning_rate=0.05,
    max_depth=7,
    num_leaves=64,
```

```

subsample=0.8,
colsample_bytree=0.8,
reg_alpha=0.1,
reg_lambda=0.1,
scale_pos_weight=5, random_state=42)
lgbm_o.fit(LGBM_X_train_selected, y_train)

# Evaluate
y_lgbm_o = lgbm_o.predict(LGBM_X_test_selected)
print("\nLGBM on Original dataset Results:")
print(classification_report(y_test, y_lgbm_o))
print("ROC AUC:", roc_auc_score(y_test, lgbm_o.predict_proba(LGBM_X_test_selected)[:,
1]))

```

In[95]:

Performance of original for all four models

```

f1_o_LR = f1_score(y_test, y_lr_o)
recall_o_LR = recall_score(y_test, y_lr_o)
precision_o_LR = precision_score(y_test, y_lr_o)
roc_auc_o_LR = roc_auc_score(y_test, lr_o.predict_proba(LR_X_test_selected)[:, 1])
tp_o_LR, tn_o_LR, fp_o_LR, fn_o_LR = confusion_matrix(y_test, y_lr_o).ravel()

```

```

f1_o_DT = f1_score(y_test, y_dt_o)
recall_o_DT = recall_score(y_test, y_dt_o)
precision_o_DT = precision_score(y_test, y_dt_o)
roc_auc_o_DT = roc_auc_score(y_test, dt_o.predict_proba(DT_X_test_selected)[:, 1])
tp_o_DT, tn_o_DT, fp_o_DT, fn_o_DT = confusion_matrix(y_test, y_dt_o).ravel()

```

```

f1_o_RF = f1_score(y_test, y_rf_o)
recall_o_RF = recall_score(y_test, y_rf_o)
precision_o_RF = precision_score(y_test, y_rf_o)
roc_auc_o_RF = roc_auc_score(y_test, rf_o.predict_proba(RF_X_test_selected)[:, 1])
tp_o_RF, tn_o_RF, fp_o_RF, fn_o_RF = confusion_matrix(y_test, y_rf_o).ravel()

```

```

f1_o_SVC = f1_score(y_test, y_svc_o)
recall_o_SVC = recall_score(y_test, y_svc_o)
precision_o_SVC = precision_score(y_test, y_svc_o)
roc_auc_o_SVC = roc_auc_score(y_test, svc_o.predict_proba(SVC_X_test_selected)[:, 1])
tp_o_SVC, tn_o_SVC, fp_o_SVC, fn_o_SVC = confusion_matrix(y_test, y_svc_o).ravel()

```

```

f1_o_XGB = f1_score(y_test, y_xgb_o)
recall_o_XGB = recall_score(y_test, y_xgb_o)
precision_o_XGB = precision_score(y_test, y_xgb_o)
roc_auc_o_XGB = roc_auc_score(y_test, xgb_o.predict_proba(XGB_X_test_selected)[:, 1])
tp_o_XGB, tn_o_XGB, fp_o_XGB, fn_o_XGB = confusion_matrix(y_test, y_xgb_o).ravel()

```

```

f1_o_LGBM = f1_score(y_test, y_lgbm_o)
recall_o_LGBM = recall_score(y_test, y_lgbm_o)
precision_o_LGBM = precision_score(y_test, y_lgbm_o)
roc_auc_o_LGBM = roc_auc_score(y_test, lgbm_o.predict_proba(LGBM_X_test_selected)[:,
1])

```

```
tp_o_LGBM, tn_o_LGBM, fp_o_LGBM, fn_o_LGBM = confusion_matrix(y_test,
y_lgbm_o).ravel()
```

```
results = {
    'No resampling on Techniques': ['Logistic Regression', 'Decision Tree', 'Random
Forest','SVC', 'XGBoost', 'LightGBM'],
    'F1-Score': [f1_o_LR, f1_o_DT, f1_o_RF, f1_o_SVC, f1_o_XGB, f1_o_LGBM],
    'Precision': [precision_o_LR, precision_o_DT, precision_o_RF,
precision_o_SVC, precision_o_XGB, precision_o_LGBM],
    'recall': [recall_o_LR, recall_o_DT, recall_o_RF, recall_o_SVC, recall_o_XGB,
recall_o_LGBM],
    'ROC AUC': [roc_auc_o_LR, roc_auc_o_DT, roc_auc_o_RF, roc_auc_o_SVC,
roc_auc_o_XGB, roc_auc_o_LGBM],
    'Predicted TP%': [(tp_o_LR/(tp_o_LR+fn_o_LR))*100, (tp_o_DT/(tp_o_DT+fn_o_DT))*100,
(tp_o_RF/(tp_o_RF+fn_o_RF))*100, (tp_o_SVC/(tp_o_SVC+fn_o_SVC))*100,
(tp_o_XGB/(tp_o_XGB+fn_o_XGB))*100, (tp_o_LGBM/(tp_o_LGBM+fn_o_LGBM))*100],
    'Predicted TN%': [(tn_o_LR/(tn_o_LR+fp_o_LR))*100,
(tn_o_DT/(tn_o_DT+fp_o_DT))*100, (tn_o_RF/(tn_o_RF+fp_o_RF))*100,
(tn_o_SVC/(tn_o_SVC+fp_o_SVC))*100, (tn_o_XGB/(tn_o_XGB+fp_o_XGB))*100,
(tn_o_LGBM/(tn_o_LGBM+fp_o_LGBM))*100],
    'Predicted FP%': [(fp_o_LR/(tn_o_LR+fp_o_LR))*100, (fp_o_DT/(tn_o_DT+fp_o_DT))*100,
(fp_o_RF/(tn_o_RF+fp_o_RF))*100, (fp_o_SVC/(tn_o_SVC+fp_o_SVC))*100,
(fp_o_XGB/(tn_o_XGB+fp_o_XGB))*100, (fp_o_LGBM/(tn_o_LGBM+fp_o_LGBM))*100],
    'Predicted FN%': [(fn_o_LR/(tp_o_LR+fn_o_LR))*100,
(fn_o_DT/(tp_o_DT+fn_o_DT))*100, (fn_o_RF/(tp_o_RF+fn_o_RF))*100,
(fn_o_SVC/(tp_o_SVC+fn_o_SVC))*100, (fn_o_XGB/(tp_o_XGB+fn_o_XGB))*100,
(fn_o_LGBM/(tp_o_LGBM+fn_o_LGBM))*100]
}
```

```
results_df = pd.DataFrame(results)
print("\nComparison of Techniques:")
display(results_df)
```

1. Random Undersampling

In[96]:

```
# Apply random undersampling on LR
rus = RandomUnderSampler(random_state=42)
X_train_rus_LR, y_train_rus_LR = rus.fit_resample(LR_X_train_selected, y_train)

# Check class distribution after undersampling
print("\nClass distribution after undersampling:")
print(pd.Series(y_train_rus_LR).value_counts())

# Train a model on undersampled data
rf_rus_LR = LogisticRegression(random_state=42)
rf_rus_LR.fit(X_train_rus_LR, y_train_rus_LR)

# Evaluate
y_pred_rus_LR = rf_rus_LR.predict(LR_X_test_selected)
print("\nUndersampling Results:")
```

```
print(classification_report(y_test, y_pred_rus_LR))
print("ROC AUC:", roc_auc_score(y_test, rf_rus_LR.predict_proba(LR_X_test_selected)[:,
1]))
```

In[97]:

```
# Apply random undersampling on DT
rus = RandomUnderSampler(random_state=42)
X_train_rus_DT, y_train_rus_DT = rus.fit_resample(DT_X_train_selected, y_train)

# Check class distribution after undersampling
print("\nClass distribution after undersampling:")
print(pd.Series(y_train_rus_DT).value_counts())

# Train a model on undersampled data
rf_rus_DT = DecisionTreeClassifier(random_state=42)
rf_rus_DT.fit(X_train_rus_DT, y_train_rus_DT)

# Evaluate
y_pred_rus_DT = rf_rus_DT.predict(DT_X_test_selected)
print("\nUndersampling Results:")
print(classification_report(y_test, y_pred_rus_DT))
print("ROC AUC:", roc_auc_score(y_test, rf_rus_DT.predict_proba(DT_X_test_selected)[:,
1]))
```

In[98]:

```
# Apply random undersampling on RF
rus = RandomUnderSampler(random_state=42)
X_train_rus_RF, y_train_rus_RF = rus.fit_resample(RF_X_train_selected, y_train)

# Check class distribution after undersampling
print("\nClass distribution after undersampling:")
print(pd.Series(y_train_rus_RF).value_counts())

# Train a model on undersampled data
rf_rus_RF = RandomForestClassifier(random_state=42)
rf_rus_RF.fit(X_train_rus_RF, y_train_rus_RF)

# Evaluate
y_pred_rus_RF = rf_rus_RF.predict(RF_X_test_selected)
print("\nUndersampling Results:")
print(classification_report(y_test, y_pred_rus_RF))
print("ROC AUC:", roc_auc_score(y_test, rf_rus_RF.predict_proba(RF_X_test_selected)[:,
1]))
```

In[99]:

```
# Apply random undersampling on SVC
```



```

rus = RandomUnderSampler(random_state=42)
X_train_rus_SVC, y_train_rus_SVC = rus.fit_resample(SVC_X_train_selected, y_train)

# Check class distribution after undersampling
print("\nClass distribution after undersampling:")
print(pd.Series(y_train_rus_SVC).value_counts())

# Train a model on undersampled data
rf_rus_SVC = SVC(kernel='linear', probability=True, random_state=42)
rf_rus_SVC.fit(X_train_rus_SVC, y_train_rus_SVC)

# Evaluate
y_pred_rus_SVC = rf_rus_SVC.predict(SVC_X_test_selected)
print("\nUndersampling Results:")
print(classification_report(y_test, y_pred_rus_SVC))
print("ROC AUC:", roc_auc_score(y_test,
rf_rus_SVC.predict_proba(SVC_X_test_selected)[:, 1]))

```

In[100]:

```

# Apply random undersampling on XGB
rus = RandomUnderSampler(random_state=42)
X_train_rus_XGB, y_train_rus_XGB = rus.fit_resample(XGB_X_train_selected, y_train)

# Check class distribution after undersampling
print("\nClass distribution after undersampling:")
print(pd.Series(y_train_rus_XGB).value_counts())

# Train a model on undersampled data
rf_rus_XGB = XGBClassifier(random_state=42)
rf_rus_XGB.fit(X_train_rus_XGB, y_train_rus_XGB)

# Evaluate
y_pred_rus_XGB = rf_rus_XGB.predict(XGB_X_test_selected)
print("\nUndersampling Results:")
print(classification_report(y_test, y_pred_rus_XGB))
print("ROC AUC:", roc_auc_score(y_test,
rf_rus_XGB.predict_proba(XGB_X_test_selected)[:, 1]))

```

In[101]:

```

# Apply random undersampling on LGBM
rus = RandomUnderSampler(random_state=42)
X_train_rus_LGBM, y_train_rus_LGBM = rus.fit_resample(LGBM_X_train_selected, y_train)

# Check class distribution after undersampling
print("\nClass distribution after undersampling:")
print(pd.Series(y_train_rus_LGBM).value_counts())

# Train a model on undersampled data
rf_rus_LGBM = LGBMClassifier(random_state=42)

```

```

rf_rus_LGBM.fit(X_train_rus_LGBM, y_train_rus_LGBM)

# Evaluate
y_pred_rus_LGBM = rf_rus_LGBM.predict(LGBM_X_test_selected)
print("\nUndersampling Results:")
print(classification_report(y_test, y_pred_rus_LGBM))
print("ROC AUC:", roc_auc_score(y_test,
rf_rus_LGBM.predict_proba(LGBM_X_test_selected)[:, 1]))

# In[102]:

# Performance of rus for all four models

f1_rus_LR = f1_score(y_test, y_pred_rus_LR)
precision_o_LR = precision_score(y_test, y_lr_o)
recall_rus_LR = recall_score(y_test, y_pred_rus_LR)
roc_auc_rus_LR = roc_auc_score(y_test, rf_rus_LR.predict_proba(LR_X_test_selected)[:, 1])
tp_rus_LR, tn_rus_LR, fp_rus_LR, fn_rus_LR = confusion_matrix(y_test,
y_pred_rus_LR).ravel()

f1_rus_DT = f1_score(y_test, y_pred_rus_DT)
precision_o_DT = precision_score(y_test, y_dt_o)
recall_rus_DT = recall_score(y_test, y_pred_rus_DT)
roc_auc_rus_DT = roc_auc_score(y_test, rf_rus_DT.predict_proba(DT_X_test_selected)[:,
1])
tp_rus_DT, tn_rus_DT, fp_rus_DT, fn_rus_DT = confusion_matrix(y_test,
y_pred_rus_DT).ravel()

f1_rus_RF = f1_score(y_test, y_pred_rus_RF)
precision_o_RF = precision_score(y_test, y_rf_o)
recall_rus_RF = recall_score(y_test, y_pred_rus_RF)
roc_auc_rus_RF = roc_auc_score(y_test, rf_rus_RF.predict_proba(RF_X_test_selected)[:,
1])
tp_rus_RF, tn_rus_RF, fp_rus_RF, fn_rus_RF = confusion_matrix(y_test,
y_pred_rus_RF).ravel()

f1_rus_SVC = f1_score(y_test, y_pred_rus_SVC)
precision_o_SVC = precision_score(y_test, y_svc_o)
recall_rus_SVC = recall_score(y_test, y_pred_rus_SVC)
roc_auc_rus_SVC = roc_auc_score(y_test,
rf_rus_SVC.predict_proba(SVC_X_test_selected)[:, 1])
tp_rus_SVC, tn_rus_SVC, fp_rus_SVC, fn_rus_SVC = confusion_matrix(y_test,
y_pred_rus_SVC).ravel()

f1_rus_XGB = f1_score(y_test, y_pred_rus_XGB)
precision_o_XGB = precision_score(y_test, y_xgb_o)
recall_rus_XGB = recall_score(y_test, y_pred_rus_XGB)
roc_auc_rus_XGB = roc_auc_score(y_test,
rf_rus_XGB.predict_proba(XGB_X_test_selected)[:, 1])
tp_rus_XGB, tn_rus_XGB, fp_rus_XGB, fn_rus_XGB = confusion_matrix(y_test,
y_pred_rus_XGB).ravel()

f1_rus_LGBM = f1_score(y_test, y_pred_rus_LGBM)

```

```
precision_o_LGBM = precision_score(y_test, y_lgbm_o)
recall_rus_LGBM = recall_score(y_test, y_pred_rus_LGBM)
roc_auc_rus_LGBM = roc_auc_score(y_test,
rf_rus_LGBM.predict_proba(LGBM_X_test_selected)[:, 1])
tp_rus_LGBM, tn_rus_LGBM, fp_rus_LGBM, fn_rus_LGBM = confusion_matrix(y_test,
y_pred_rus_LGBM).ravel()
```

```
results = {
    'rus on Technique': ['Logistic Regression', 'Decision Tree', 'Random Forest', 'SVC',
    'XGBoost', 'LightGBM'],
    'F1-Score': [f1_rus_LR, f1_rus_DT, f1_rus_RF, f1_rus_SVC, f1_rus_XGB, f1_rus_LGBM],
    'Precision':
    [precision_o_LR, precision_o_DT, precision_o_RF, precision_o_SVC, precision_o_XGB, precision_o_LGBM],
    'recall' : [recall_rus_LR, recall_rus_DT, recall_rus_RF, recall_rus_SVC, recall_o_XGB,
    recall_rus_LGBM],
    'ROC AUC': [roc_auc_rus_LR, roc_auc_rus_DT, roc_auc_rus_RF, roc_auc_rus_SVC,
    roc_auc_rus_XGB, roc_auc_rus_LGBM],
    'Predicted TP%': [(tp_rus_LR/(tp_rus_LR+fn_rus_LR))*100,
    (tp_rus_DT/(tp_rus_DT+fn_rus_DT))*100,
    (tp_rus_RF/(tp_rus_RF+fn_rus_RF))*100,
    (tp_rus_SVC/(tp_rus_SVC+fn_rus_SVC))*100,
    (tp_rus_XGB/(tp_rus_XGB+fn_rus_XGB))*100,
    (tp_rus_LGBM/(tp_rus_LGBM+fn_rus_LGBM))*100],
    'Predicted TN%': [(tn_rus_LR/(tn_rus_LR+fp_rus_LR))*100,
    (tn_rus_DT/(tn_rus_DT+fp_rus_DT))*100,
    (tn_rus_RF/(tn_rus_RF+fp_rus_RF))*100,
    (tn_rus_SVC/(tn_rus_SVC+fp_rus_SVC))*100,
    (tn_rus_XGB/(tn_rus_XGB+fp_rus_XGB))*100,
    (tn_rus_LGBM/(tn_rus_LGBM+fp_rus_LGBM))*100],
    'Predicted FP%': [(fp_rus_LR/(tn_rus_LR+fp_rus_LR))*100,
    (fp_rus_DT/(tn_rus_DT+fp_rus_DT))*100,
    (fp_rus_RF/(tn_rus_RF+fp_rus_RF))*100,
    (fp_rus_SVC/(tn_rus_SVC+fp_rus_SVC))*100,
    (fp_rus_XGB/(tn_rus_XGB+fp_rus_XGB))*100,
    (fp_rus_LGBM/(tn_rus_LGBM+fp_rus_LGBM))*100],
    'Predicted FN%': [(fn_rus_LR/(tp_rus_LR+fn_rus_LR))*100,
    (fn_rus_DT/(tp_rus_DT+fn_rus_DT))*100,
    (fn_rus_RF/(tp_rus_RF+fn_rus_RF))*100,
    (fn_rus_SVC/(tp_rus_SVC+fn_rus_SVC))*100,
    (fn_rus_XGB/(tp_rus_XGB+fn_rus_XGB))*100,
    (fn_rus_LGBM/(tp_rus_LGBM+fn_rus_LGBM))*100]
}
```

```
results_df = pd.DataFrame(results)
print("\nComparison of Techniques:")
display(results_df)
```

2. SMOTE (Synthetic Minority Oversampling)

In[103]:

```
# Apply SMOTE on LR
smote = SMOTE(random_state=42)
X_train_smote_LR, y_train_smote_LR = smote.fit_resample(LR_X_train_selected, y_train)
```

```

# Check class distribution after SMOTE
print("\nClass distribution after SMOTE:")
print(pd.Series(y_train_smote_LR).value_counts())

# Train a model on SMOTE data
rf_smote_LR = LogisticRegression(random_state=42)
rf_smote_LR.fit(X_train_smote_LR, y_train_smote_LR)

# Evaluate
y_pred_smote_LR = rf_smote_LR.predict(LR_X_test_selected)
print("\nSMOTE Results:")
print(classification_report(y_test, y_pred_smote_LR))
print("ROC AUC:", roc_auc_score(y_test, rf_smote_LR.predict_proba(LR_X_test_selected)[:,
1]))

```

In[104]:

```

# Apply SMOTE on DT
smote = SMOTE(random_state=42)
X_train_smote_DT, y_train_smote_DT = smote.fit_resample(DT_X_train_selected, y_train)

# Check class distribution after SMOTE
print("\nClass distribution after SMOTE:")
print(pd.Series(y_train_smote_DT).value_counts())

# Train a model on SMOTE data
rf_smote_DT = DecisionTreeClassifier(random_state=42)
rf_smote_DT.fit(X_train_smote_DT, y_train_smote_DT)

# Evaluate
y_pred_smote_DT = rf_smote_DT.predict(DT_X_test_selected)
print("\nSMOTE Results:")
print(classification_report(y_test, y_pred_smote_DT))
print("ROC AUC:", roc_auc_score(y_test,
rf_smote_DT.predict_proba(DT_X_test_selected)[:, 1]))

```

In[105]:

```

# Apply SMOTE on RF
smote = SMOTE(random_state=42)
X_train_smote_RF, y_train_smote_RF = smote.fit_resample(RF_X_train_selected, y_train)

# Check class distribution after SMOTE
print("\nClass distribution after SMOTE:")
print(pd.Series(y_train_smote_RF).value_counts())

# Train a model on SMOTE data
rf_smote_RF = RandomForestClassifier(random_state=42)
rf_smote_RF.fit(X_train_smote_RF, y_train_smote_RF)

# Evaluate

```

```

y_pred_smote_RF = rf_smote_RF.predict(RF_X_test_selected)
print("\nSMOTE Results:")
print(classification_report(y_test, y_pred_smote_RF))
print("ROC AUC:", roc_auc_score(y_test,
rf_smote_RF.predict_proba(RF_X_test_selected)[: , 1]))

```

In[106]:

```

# Apply smote on SVC
smote = SMOTE(random_state=42)
X_train_smote_SVC, y_train_smote_SVC = smote.fit_resample(SVC_X_train_selected,
y_train)

# Check class distribution after undersampling
print("\nClass distribution after undersampling:")
print(pd.Series(y_train_smote_SVC).value_counts())

# Train a model on undersampled data
rf_smote_SVC = SVC(kernel='linear',probability=True,random_state=42)
rf_smote_SVC.fit(X_train_smote_SVC, y_train_smote_SVC)

# Evaluate
y_pred_smote_SVC = rf_smote_SVC.predict(SVC_X_test_selected)
print("\nSMOTE Results:")
print(classification_report(y_test, y_pred_smote_SVC))
print("ROC AUC:", roc_auc_score(y_test,
rf_smote_SVC.predict_proba(SVC_X_test_selected)[: , 1]))

```

In[107]:

```

# Apply SMOTE on XGB
smote = SMOTE(random_state=42)
X_train_smote_XGB, y_train_smote_XGB = smote.fit_resample(XGB_X_train_selected,
y_train)

# Check class distribution after SMOTE
print("\nClass distribution after SMOTE:")
print(pd.Series(y_train_smote_XGB).value_counts())

# Train a model on SMOTE data
rf_smote_XGB = XGBClassifier(random_state=42)
rf_smote_XGB.fit(X_train_smote_XGB, y_train_smote_XGB)

# Evaluate
y_pred_smote_XGB = rf_smote_XGB.predict(XGB_X_test_selected)
print("\nSMOTE Results:")
print(classification_report(y_test, y_pred_smote_XGB))
print("ROC AUC:", roc_auc_score(y_test,
rf_smote_XGB.predict_proba(XGB_X_test_selected)[: , 1]))

```

```
# In[108]:
```

```
# Apply SMOTE on LightGBM
smote = SMOTE(random_state=42)
X_train_smote_LGBM, y_train_smote_LGBM = smote.fit_resample(LGBM_X_train_selected,
y_train)

# Check class distribution after SMOTE
print("\nClass distribution after SMOTE:")
print(pd.Series(y_train_smote_LGBM).value_counts())

# Train a model on SMOTE data
rf_smote_LGBM = LGBMClassifier(random_state=42)
rf_smote_LGBM.fit(X_train_smote_LGBM, y_train_smote_LGBM)

# Evaluate
y_pred_smote_LGBM = rf_smote_LGBM.predict(LGBM_X_test_selected)
print("\nSMOTE Results:")
print(classification_report(y_test, y_pred_smote_LGBM))
print("ROC AUC:", roc_auc_score(y_test,
rf_smote_LGBM.predict_proba(LGBM_X_test_selected)[: , 1]))
```

```
# In[109]:
```

```
# Performance of SMOTE for all four models
```

```
f1_SMOTE_LR = f1_score(y_test, y_pred_smote_LR)
recall_SMOTE_LR = recall_score(y_test, y_pred_smote_LR)
roc_auc_SMOTE_LR = roc_auc_score(y_test,
rf_smote_LR.predict_proba(LR_X_test_selected)[: , 1])
precision_SMOTE_LR = precision_score(y_test, y_pred_smote_LR)
tp_SMOTE_LR, tn_SMOTE_LR, fp_SMOTE_LR, fn_SMOTE_LR = confusion_matrix(y_test,
y_pred_smote_LR).ravel()
```

```
f1_SMOTE_DT = f1_score(y_test, y_pred_smote_DT)
recall_SMOTE_DT = recall_score(y_test, y_pred_smote_DT)
roc_auc_SMOTE_DT = roc_auc_score(y_test,
rf_smote_DT.predict_proba(DT_X_test_selected)[: , 1])
precision_SMOTE_DT = precision_score(y_test, y_pred_smote_DT)
tp_SMOTE_DT, tn_SMOTE_DT, fp_SMOTE_DT, fn_SMOTE_DT = confusion_matrix(y_test,
y_pred_smote_DT).ravel()
```

```
f1_SMOTE_RF = f1_score(y_test, y_pred_smote_RF)
recall_SMOTE_RF = recall_score(y_test, y_pred_smote_RF)
roc_auc_SMOTE_RF = roc_auc_score(y_test,
rf_smote_RF.predict_proba(RF_X_test_selected)[: , 1])
precision_SMOTE_RF = precision_score(y_test, y_pred_smote_RF)
tp_SMOTE_RF, tn_SMOTE_RF, fp_SMOTE_RF, fn_SMOTE_RF = confusion_matrix(y_test,
y_pred_smote_RF).ravel()
```

```
f1_SMOTE_SVC = f1_score(y_test, y_pred_smote_SVC)
```

```

recall_SMOTE_SVC = recall_score(y_test, y_pred_smote_SVC)
roc_auc_SMOTE_SVC = roc_auc_score(y_test,
rf_smote_SVC.predict_proba(SVC_X_test_selected)[:, 1])
precision_SMOTE_SVC = precision_score(y_test, y_pred_smote_SVC)
tp_SMOTE_SVC, tn_SMOTE_SVC, fp_SMOTE_SVC, fn_SMOTE_SVC =
confusion_matrix(y_test, y_pred_smote_SVC).ravel()

```

```

f1_SMOTE_XGB = f1_score(y_test, y_pred_smote_XGB)
recall_SMOTE_XGB = recall_score(y_test, y_pred_smote_XGB)
roc_auc_SMOTE_XGB = roc_auc_score(y_test,
rf_smote_XGB.predict_proba(XGB_X_test_selected)[:, 1])
precision_SMOTE_XGB = precision_score(y_test, y_pred_smote_XGB)
tp_SMOTE_XGB, tn_SMOTE_XGB, fp_SMOTE_XGB, fn_SMOTE_XGB =
confusion_matrix(y_test, y_pred_smote_XGB).ravel()

```

```

f1_SMOTE_LGBM = f1_score(y_test, y_pred_smote_LGBM)
recall_SMOTE_LGBM = recall_score(y_test, y_pred_smote_LGBM)
roc_auc_SMOTE_LGBM = roc_auc_score(y_test,
rf_smote_LGBM.predict_proba(LGBM_X_test_selected)[:, 1])
precision_SMOTE_LGBM = precision_score(y_test, y_pred_smote_LGBM)
tp_SMOTE_LGBM, tn_SMOTE_LGBM, fp_SMOTE_LGBM, fn_SMOTE_LGBM =
confusion_matrix(y_test, y_pred_smote_LGBM).ravel()

```

```

results = {
    'smote on Technique': ['Logistic Regression', 'Decision Tree', 'Random Forest', 'SVC',
'XGBoost', 'LightGBM'],
    'F1-Score': [f1_SMOTE_LR, f1_SMOTE_DT, f1_SMOTE_RF, f1_SMOTE_SVC,
f1_SMOTE_XGB, f1_SMOTE_LGBM],

```

```

'Precision': [precision_SMOTE_LR, precision_SMOTE_DT, precision_SMOTE_RF, precision_
SMOTE_SVC, precision_SMOTE_XGB, precision_SMOTE_LGBM],
'recall': [recall_SMOTE_LR, recall_SMOTE_DT, recall_SMOTE_RF, recall_SMOTE_SVC,
recall_SMOTE_XGB, recall_SMOTE_LGBM],
'ROC AUC': [roc_auc_SMOTE_LR, roc_auc_SMOTE_DT, roc_auc_SMOTE_RF,
roc_auc_SMOTE_SVC, roc_auc_SMOTE_XGB, roc_auc_SMOTE_LGBM],
'Predicted TP%': [(tp_SMOTE_LR/(tp_SMOTE_LR+fn_SMOTE_LR))*100,
(tp_SMOTE_DT/(tp_SMOTE_DT+fn_SMOTE_DT))*100,
(tp_SMOTE_RF/(tp_SMOTE_RF+fn_SMOTE_RF))*100,
(tp_SMOTE_SVC/(tp_SMOTE_SVC+fn_SMOTE_SVC))*100,
(tp_SMOTE_XGB/(tp_SMOTE_XGB+fn_SMOTE_XGB))*100,
(tp_SMOTE_LGBM/(tp_SMOTE_LGBM+fn_SMOTE_LGBM))*100],
'Predicted TN%': [(tn_SMOTE_LR/(tn_SMOTE_LR+fp_SMOTE_LR))*100,
(tn_SMOTE_DT/(tn_SMOTE_DT+fp_SMOTE_DT))*100,
(tn_SMOTE_RF/(tn_SMOTE_RF+fp_SMOTE_RF))*100,
(tn_SMOTE_SVC/(tn_SMOTE_SVC+fp_SMOTE_SVC))*100,
(tn_SMOTE_XGB/(tn_SMOTE_XGB+fp_SMOTE_XGB))*100,
(tn_SMOTE_LGBM/(tn_SMOTE_LGBM+fp_SMOTE_LGBM))*100],
'Predicted FP%': [(fp_SMOTE_LR/(tn_SMOTE_LR+fp_SMOTE_LR))*100,
(fp_SMOTE_DT/(tn_SMOTE_DT+fp_SMOTE_DT))*100,
(fp_SMOTE_RF/(tn_SMOTE_RF+fp_SMOTE_RF))*100,
(fp_SMOTE_SVC/(tn_SMOTE_SVC+fp_SMOTE_SVC))*100,
(fp_SMOTE_XGB/(tn_SMOTE_XGB+fp_SMOTE_XGB))*100,
(fp_SMOTE_LGBM/(tn_SMOTE_LGBM+fp_SMOTE_LGBM))*100],

```

```

'Predicted      FN%':      [(fn_SMOTE_LR/(tp_SMOTE_LR+fn_SMOTE_LR))*100,
(fn_SMOTE_DT/(tp_SMOTE_DT+fn_SMOTE_DT))*100,
(fn_SMOTE_RF/(tp_SMOTE_RF+fn_SMOTE_RF))*100,
(fn_SMOTE_SVC/(tp_SMOTE_SVC+fn_SMOTE_SVC))*100,
(fn_SMOTE_XGB/(tp_SMOTE_XGB+fn_SMOTE_XGB))*100,
(fn_SMOTE_LGBM/(tp_SMOTE_LGBM+fn_SMOTE_LGBM))*100]
}

```

```

results_df = pd.DataFrame(results)
print("\nComparison of Techniques:")
display(results_df)

```

```

#
# # 3. Threshold Optimization

```

```

# In[110]:

```

```

# First train LR on original imbalanced data
lr_th = LogisticRegression(random_state=42)
lr_th.fit(LR_X_train_selected, y_train)

# Get predicted probabilities
y_probs_LR = lr_th.predict_proba(LR_X_test_selected)[:, 1]

# Find optimal threshold
thresholds_LR = np.linspace(0, 1, 100)
f1_scores_LR = [f1_score(y_test, y_probs_LR > t) for t in thresholds_LR]
optimal_threshold_LR = thresholds_LR[np.argmax(f1_scores_LR)]

print(f"\nOptimal threshold: {optimal_threshold_LR:.2f}")

# Apply optimal threshold
y_pred_optimal_LR = (y_probs_LR > optimal_threshold_LR).astype(int)

# Evaluate
print("\nThreshold Optimization Results:")
print(classification_report(y_test, y_pred_optimal_LR))
print("ROC AUC:", roc_auc_score(y_test, y_probs_LR))

```

```

# In[111]:

```

```

# First train DT on original imbalanced data
dt_th = DecisionTreeClassifier(random_state=42)
dt_th.fit(DT_X_train_selected, y_train)

# Get predicted probabilities
y_probs_DT = dt_th.predict_proba(DT_X_test_selected)[:, 1]

# Find optimal threshold
thresholds_DT = np.linspace(0, 1, 100)
f1_scores_DT = [f1_score(y_test, y_probs_DT > t) for t in thresholds_DT]

```



```

optimal_threshold_DT = thresholds_DT[np.argmax(f1_scores_DT)]

print(f"\nOptimal threshold: {optimal_threshold_DT:.2f}")

# Apply optimal threshold
y_pred_optimal_DT = (y_probs_DT > optimal_threshold_DT).astype(int)

# Evaluate
print("\nThreshold Optimization Results:")
print(classification_report(y_test, y_pred_optimal_DT))
print("ROC AUC:", roc_auc_score(y_test, y_probs_DT))

# In[112]:

# First train RF on original imbalanced data
rf_th = RandomForestClassifier(random_state=42)
rf_th.fit(RF_X_train_selected, y_train)

# Get predicted probabilities
y_probs_RF = rf_th.predict_proba(RF_X_test_selected)[:, 1]

# Find optimal threshold
thresholds_RF = np.linspace(0, 1, 100)
f1_scores_RF = [f1_score(y_test, y_probs_RF > t) for t in thresholds_RF]
optimal_threshold_RF = thresholds_RF[np.argmax(f1_scores_RF)]

print(f"\nOptimal threshold: {optimal_threshold_RF:.2f}")

# Apply optimal threshold
y_pred_optimal_RF = (y_probs_RF > optimal_threshold_RF).astype(int)

# Evaluate
print("\nThreshold Optimization Results:")
print(classification_report(y_test, y_pred_optimal_RF))
print("ROC AUC:", roc_auc_score(y_test, y_probs_RF))

# In[113]:

# First train svc on original imbalanced data
svc_th = SVC(kernel='linear', probability=True, random_state=42)
svc_th.fit(SVC_X_train_selected, y_train)

# Get predicted probabilities
y_probs_SVC = svc_th.predict_proba(SVC_X_test_selected)[:, 1]

# Find optimal threshold
thresholds_SVC = np.linspace(0, 1, 100)
f1_scores_SVC = [f1_score(y_test, y_probs_SVC > t) for t in thresholds_SVC]
optimal_threshold_SVC = thresholds_SVC[np.argmax(f1_scores_SVC)]

print(f"\nOptimal threshold: {optimal_threshold_SVC:.2f}")

```

```
# Apply optimal threshold
y_pred_optimal_SVC = (y_probs_SVC > optimal_threshold_SVC).astype(int)
```

```
# Evaluate
print("\nThreshold Optimization Results:")
print(classification_report(y_test, y_pred_optimal_SVC))
print("ROC AUC:", roc_auc_score(y_test, y_probs_SVC))
```

```
# In[114]:
```

```
# First train XGB on original imbalanced data
xgb_th = XGBClassifier(random_state=42)
xgb_th.fit(XGB_X_train_selected, y_train)
```

```
# Get predicted probabilities
y_probs_XGB = xgb_th.predict_proba(XGB_X_test_selected)[:, 1]
```

```
# Find optimal threshold
thresholds_XGB = np.linspace(0, 1, 100)
f1_scores_XGB = [f1_score(y_test, y_probs_XGB > t) for t in thresholds_XGB]
optimal_threshold_XGB = thresholds_XGB[np.argmax(f1_scores_XGB)]
```

```
print(f"\nOptimal threshold: {optimal_threshold_XGB:.2f}")
```

```
# Apply optimal threshold
y_pred_optimal_XGB = (y_probs_XGB > optimal_threshold_XGB).astype(int)
```

```
# Evaluate
print("\nThreshold Optimization Results:")
print(classification_report(y_test, y_pred_optimal_XGB))
print("ROC AUC:", roc_auc_score(y_test, y_probs_XGB))
```

```
# In[115]:
```

```
# First train LightGBM on original imbalanced data
rf = LGBMClassifier(random_state=42)
rf.fit(LGBM_X_train_selected, y_train)
```

```
# Get predicted probabilities
y_probs_LGBM = rf.predict_proba(LGBM_X_test_selected)[:, 1]
```

```
# Find optimal threshold
thresholds_LGBM = np.linspace(0, 1, 100)
f1_scores_LGBM = [f1_score(y_test, y_probs_LGBM > t) for t in thresholds_LGBM]
optimal_threshold_LGBM = thresholds_LGBM[np.argmax(f1_scores_LGBM)]
```

```
print(f"\nOptimal threshold: {optimal_threshold_LGBM:.2f}")
```

```
# Apply optimal threshold
y_pred_optimal_LGBM = (y_probs_LGBM > optimal_threshold_LGBM).astype(int)
```

```

# Evaluate
print("\nThreshold Optimization Results:")
print(classification_report(y_test, y_pred_optimal_LGBM))
print("ROC AUC:", roc_auc_score(y_test, y_probs_LGBM))

# In[116]:

# Performance of Threshold Optimization for all four models

f1_Th_LR = f1_score(y_test, y_pred_optimal_LR)
recall_Th_LR = recall_score(y_test, y_pred_optimal_LR)
roc_auc_Th_LR = roc_auc_score(y_test, y_probs_LR)
precision_Th_LR = precision_score(y_test, y_pred_optimal_LR)
tp_Th_LR, tn_Th_LR, fp_Th_LR, fn_Th_LR = confusion_matrix(y_test,
y_pred_optimal_LR).ravel()

f1_Th_DT = f1_score(y_test, y_pred_optimal_DT)
recall_Th_DT = recall_score(y_test, y_pred_optimal_DT)
roc_auc_Th_DT = roc_auc_score(y_test, y_probs_DT)
precision_Th_DT = precision_score(y_test, y_pred_optimal_DT)
tp_Th_DT, tn_Th_DT, fp_Th_DT, fn_Th_DT = confusion_matrix(y_test,
y_pred_optimal_DT).ravel()

f1_Th_RF = f1_score(y_test, y_pred_optimal_RF)
recall_Th_RF = recall_score(y_test, y_pred_optimal_RF)
roc_auc_Th_RF = roc_auc_score(y_test, y_probs_RF)
precision_Th_RF = precision_score(y_test, y_pred_optimal_RF)
tp_Th_RF, tn_Th_RF, fp_Th_RF, fn_Th_RF = confusion_matrix(y_test,
y_pred_optimal_RF).ravel()

f1_Th_SVC = f1_score(y_test, y_pred_optimal_SVC)
recall_Th_SVC = recall_score(y_test, y_pred_optimal_SVC)
roc_auc_Th_SVC = roc_auc_score(y_test, y_probs_SVC)
precision_Th_SVC = precision_score(y_test, y_pred_optimal_SVC)
tp_Th_SVC, tn_Th_SVC, fp_Th_SVC, fn_Th_SVC = confusion_matrix(y_test,
y_pred_optimal_SVC).ravel()

f1_Th_XGB = f1_score(y_test, y_pred_optimal_XGB)
recall_Th_XGB = recall_score(y_test, y_pred_optimal_XGB)
roc_auc_Th_XGB = roc_auc_score(y_test, y_probs_XGB)
precision_Th_XGB = precision_score(y_test, y_pred_optimal_XGB)
tp_Th_XGB, tn_Th_XGB, fp_Th_XGB, fn_Th_XGB = confusion_matrix(y_test,
y_pred_optimal_XGB).ravel()

f1_Th_LGBM = f1_score(y_test, y_pred_optimal_LGBM)
recall_Th_LGBM = recall_score(y_test, y_pred_optimal_LGBM)
roc_auc_Th_LGBM = roc_auc_score(y_test, y_probs_LGBM)

```

```
precision_Th_LGBM = precision_score(y_test, y_pred_optimal_LGBM)
tp_Th_LGBM, tn_Th_LGBM, fp_Th_LGBM, fn_Th_LGBM = confusion_matrix(y_test,
y_pred_optimal_LGBM).ravel()
```

```
results = {
    'Th on Technique': ['Logistic Regression', 'Decision Tree', 'Random Forest','SVC',
'XGBoost', 'LightGBM'],
    'F1-Score': [f1_Th_LR, f1_Th_DT, f1_Th_RF, f1_Th_SVC, f1_Th_XGB, f1_Th_LGBM],
    'recall' : [recall_Th_LR, recall_Th_DT, recall_Th_RF, recall_Th_SVC, recall_Th_XGB,
recall_Th_LGBM],
    'ROC AUC': [roc_auc_Th_LR, roc_auc_Th_DT, roc_auc_Th_RF, roc_auc_Th_SVC,
roc_auc_Th_XGB, roc_auc_Th_LGBM],

    'Precision': [precision_Th_LR, precision_Th_DT, precision_Th_RF, precision_Th_SVC, precision_Th_XGB, precision_Th_LGBM],
    'Predicted TP%': [(tp_Th_LR/(tp_Th_LR+fn_Th_LR))*100,
(tp_Th_DT/(tp_Th_DT+fn_Th_DT))*100, (tp_Th_RF/(tp_Th_RF+fn_Th_RF))*100,
(tp_Th_SVC/(tp_Th_SVC+fn_Th_SVC))*100, (tp_Th_XGB/(tp_Th_XGB+fn_Th_XGB))*100,
(tp_Th_LGBM/(tp_Th_LGBM+fn_Th_LGBM))*100],
    'Predicted TN%': [(tn_Th_LR/(tn_Th_LR+fp_Th_LR))*100,
(tn_Th_DT/(tn_Th_DT+fp_Th_DT))*100, (tn_Th_RF/(tn_Th_RF+fp_Th_RF))*100,
(tn_Th_SVC/(tn_Th_SVC+fp_Th_SVC))*100, (tn_Th_XGB/(tn_Th_XGB+fp_Th_XGB))*100,
(tn_Th_LGBM/(tn_Th_LGBM+fp_Th_LGBM))*100],
    'Predicted FP%': [(fp_Th_LR/(tn_Th_LR+fp_Th_LR))*100,
(fp_Th_DT/(tn_Th_DT+fp_Th_DT))*100, (fp_Th_RF/(tn_Th_RF+fp_Th_RF))*100,
(fp_Th_SVC/(tn_Th_SVC+fp_Th_SVC))*100, (fp_Th_XGB/(tn_Th_XGB+fp_Th_XGB))*100,
(fp_Th_LGBM/(tn_Th_LGBM+fp_Th_LGBM))*100],
    'Predicted FN%': [(fn_Th_LR/(tp_Th_LR+fn_Th_LR))*100,
(fn_Th_DT/(tp_Th_DT+fn_Th_DT))*100, (fn_Th_RF/(tp_Th_RF+fn_Th_RF))*100,
(fn_Th_SVC/(tp_Th_SVC+fn_Th_SVC))*100, (fn_Th_XGB/(tp_Th_XGB+fn_Th_XGB))*100,
(fn_Th_LGBM/(tp_Th_LGBM+fn_Th_LGBM))*100]
}
```

```
results_df = pd.DataFrame(results)
print("\nComparison of Techniques:")
display(results_df)
```

4. Cost Sensitive Learning

```
# In[117]:
```

```
# Compute class weights
classes = np.unique(y_train)
weights = compute_class_weight(class_weight='balanced', classes=classes, y=y_train)
class_weights = dict(zip(classes, weights))
```

```
print("\nClass weights:", class_weights)
```

```
# Train LR with class weights
rf_weighted_LR = LogisticRegression(class_weight=class_weights, random_state=42)
rf_weighted_LR.fit(LR_X_train_selected, y_train)
```

```
# Evaluate LR
y_pred_weighted_LR = rf_weighted_LR.predict(LR_X_test_selected)
print("\nCost-Sensitive Learning Results:")
print(classification_report(y_test, y_pred_weighted_LR))
print("ROC AUC:", roc_auc_score(y_test,
rf_weighted_LR.predict_proba(LR_X_test_selected)[:, 1]))
```

In[118]:

```
# Train DT with class weights
rf_weighted_DT = DecisionTreeClassifier(class_weight=class_weights, random_state=42)
rf_weighted_DT.fit(DT_X_train_selected, y_train)

# Evaluate
y_pred_weighted_DT = rf_weighted_DT.predict(DT_X_test_selected)
print("\nCost-Sensitive Learning Results:")
print(classification_report(y_test, y_pred_weighted_DT))
print("ROC AUC:", roc_auc_score(y_test,
rf_weighted_DT.predict_proba(DT_X_test_selected)[:, 1]))
```

In[119]:

```
# Train RF with class weights
rf_weighted_RF = RandomForestClassifier(class_weight=class_weights, random_state=42)
rf_weighted_RF.fit(RF_X_train_selected, y_train)

# Evaluate
y_pred_weighted_RF = rf_weighted_RF.predict(RF_X_test_selected)
print("\nCost-Sensitive Learning Results:")
print(classification_report(y_test, y_pred_weighted_RF))
print("ROC AUC:", roc_auc_score(y_test,
rf_weighted_RF.predict_proba(RF_X_test_selected)[:, 1]))
```

In[120]:

```
# Train SVC with class weights
rf_weighted_SVC = SVC(kernel='linear', probability=True, class_weight=class_weights,
random_state=42)
rf_weighted_SVC.fit(SVC_X_train_selected, y_train)

# Evaluate
y_pred_weighted_SVC = rf_weighted_SVC.predict(SVC_X_test_selected)
print("\nCost-Sensitive Learning Results:")
print(classification_report(y_test, y_pred_weighted_SVC))
print("ROC AUC:", roc_auc_score(y_test,
rf_weighted_SVC.predict_proba(SVC_X_test_selected)[:, 1]))
```

In[121]:

```
# Train XGB with class weights
rf_weighted_XGB = XGBClassifier(class_weight=class_weights, random_state=42)
rf_weighted_XGB.fit(XGB_X_train_selected, y_train)

# Evaluate
y_pred_weighted_XGB = rf_weighted_XGB.predict(XGB_X_test_selected)
print("\nCost-Sensitive Learning Results:")
print(classification_report(y_test, y_pred_weighted_XGB))
print("ROC AUC:", roc_auc_score(y_test,
rf_weighted_XGB.predict_proba(XGB_X_test_selected)[:, 1]))
```

In[122]:

```
# Train LightGBM with class weights
rf_weighted_LGBM = LGBMClassifier(class_weight=class_weights, random_state=42)
rf_weighted_LGBM.fit(LGBM_X_train_selected, y_train)

# Evaluate
y_pred_weighted_LGBM = rf_weighted_LGBM.predict(LGBM_X_test_selected)
print("\nCost-Sensitive Learning Results:")
print(classification_report(y_test, y_pred_weighted_LGBM))
print("ROC AUC:", roc_auc_score(y_test,
rf_weighted_LGBM.predict_proba(LGBM_X_test_selected)[:, 1]))
```

In[123]:

Performance of Cost Sensitive Learning for all four models

```
f1_CS_LR = f1_score(y_test, y_pred_weighted_LR)
recall_CS_LR = recall_score(y_test, y_pred_weighted_LR)
roc_auc_CS_LR = roc_auc_score(y_test,
rf_weighted_LR.predict_proba(LR_X_test_selected)[:, 1])
precision_CS_LR = precision_score(y_test, y_pred_weighted_LR)
tp_CS_LR, tn_CS_LR, fp_CS_LR, fn_CS_LR = confusion_matrix(y_test,
y_pred_weighted_LR).ravel()
accuracy_CS_LR = ((tp_CS_LR + tn_CS_LR)/(tp_CS_LR + tn_CS_LR + fp_CS_LR +
fn_CS_LR))*100

f1_CS_DT = f1_score(y_test, y_pred_weighted_DT)
recall_CS_DT = recall_score(y_test, y_pred_weighted_DT)
roc_auc_CS_DT = roc_auc_score(y_test,
rf_weighted_DT.predict_proba(DT_X_test_selected)[:, 1])
precision_CS_DT = precision_score(y_test, y_pred_weighted_DT)
tp_CS_DT, tn_CS_DT, fp_CS_DT, fn_CS_DT = confusion_matrix(y_test,
y_pred_weighted_DT).ravel()
accuracy_CS_DT = ((tp_CS_DT + tn_CS_DT)/(tp_CS_DT + tn_CS_DT + fp_CS_DT +
fn_CS_DT))*100
```

```

f1_CS_RF = f1_score(y_test, y_pred_weighted_RF)
recall_CS_RF = recall_score(y_test, y_pred_weighted_RF)
roc_auc_CS_RF = roc_auc_score(y_test,
rf_weighted_RF.predict_proba(RF_X_test_selected)[:, 1])
precision_CS_RF = precision_score(y_test, y_pred_weighted_RF)
tp_CS_RF, tn_CS_RF, fp_CS_RF, fn_CS_RF = confusion_matrix(y_test,
y_pred_weighted_RF).ravel()
accuracy_CS_RF = ((tp_CS_RF + tn_CS_RF)/(tp_CS_RF + tn_CS_RF + fp_CS_RF +
fn_CS_RF))*100

```

```

f1_CS_SVC = f1_score(y_test, y_pred_weighted_SVC)
recall_CS_SVC = recall_score(y_test, y_pred_weighted_SVC)
roc_auc_CS_SVC = roc_auc_score(y_test,
rf_weighted_SVC.predict_proba(SVC_X_test_selected)[:, 1])
precision_CS_SVC = precision_score(y_test, y_pred_weighted_SVC)
tp_CS_SVC, tn_CS_SVC, fp_CS_SVC, fn_CS_SVC = confusion_matrix(y_test,
y_pred_weighted_SVC).ravel()
accuracy_CS_SVC = ((tp_CS_SVC + tn_CS_SVC)/(tp_CS_SVC + tn_CS_SVC +
fp_CS_SVC + fn_CS_SVC))*100

```

```

f1_CS_XGB = f1_score(y_test, y_pred_weighted_XGB)
recall_CS_XGB = recall_score(y_test, y_pred_weighted_XGB)
roc_auc_CS_XGB = roc_auc_score(y_test,
rf_weighted_XGB.predict_proba(XGB_X_test_selected)[:, 1])
precision_CS_XGB = precision_score(y_test, y_pred_weighted_XGB)
tp_CS_XGB, tn_CS_XGB, fp_CS_XGB, fn_CS_XGB = confusion_matrix(y_test,
y_pred_weighted_XGB).ravel()
accuracy_CS_XGB = ((tp_CS_XGB + tn_CS_XGB)/(tp_CS_XGB + tn_CS_XGB +
fp_CS_XGB + fn_CS_XGB))*100

```

```

f1_CS_LGBM = f1_score(y_test, y_pred_weighted_LGBM)
recall_CS_LGBM = recall_score(y_test, y_pred_weighted_LGBM)
roc_auc_CS_LGBM = roc_auc_score(y_test,
rf_weighted_LGBM.predict_proba(LGBM_X_test_selected)[:, 1])
precision_CS_LGBM = precision_score(y_test, y_pred_weighted_LGBM)
tp_CS_LGBM, tn_CS_LGBM, fp_CS_LGBM, fn_CS_LGBM = confusion_matrix(y_test,
y_pred_weighted_LGBM).ravel()
accuracy_CS_LGBM = ((tp_CS_LGBM + tn_CS_LGBM)/(tp_CS_LGBM + tn_CS_LGBM +
fp_CS_LGBM + fn_CS_LGBM))*100

```

```

results = {
    'CS on Technique': ['Logistic Regression', 'Decision Tree', 'Random Forest', 'SVC',
'XGBoost', 'LightGBM'],
    'F1-Score': [f1_CS_LR, f1_CS_DT, f1_CS_RF, f1_CS_SVC, f1_CS_XGB, f1_CS_LGBM],
    'Accuracy':
[accuracy_CS_LR, accuracy_CS_DT, accuracy_CS_RF, accuracy_CS_SVC, accuracy_CS_X
GB, accuracy_CS_LGBM],

    'Precision': [precision_CS_LR, precision_CS_DT, precision_CS_RF, precision_CS_SVC, precis
ion_CS_XGB, precision_CS_LGBM],
    'recall' : [recall_CS_LR, recall_CS_DT, recall_CS_RF, recall_CS_SVC, recall_CS_XGB,
recall_CS_LGBM],
    'ROC AUC': [roc_auc_CS_LR, roc_auc_CS_DT, roc_auc_CS_RF, roc_auc_CS_SVC,
roc_auc_CS_XGB, roc_auc_CS_LGBM],

```

```

'Predicted TP%': [(tp_CS_LR/(tp_CS_LR+fn_CS_LR))*100,
(tp_CS_DT/(tp_CS_DT+fn_CS_DT))*100,
(tp_CS_SVC/(tp_CS_SVC+fn_CS_SVC))*100,
(tp_CS_XGB/(tp_CS_XGB+fn_CS_XGB))*100,
(tp_CS_LGBM/(tp_CS_LGBM+fn_CS_LGBM))*100],
'Predicted TN%': [(tn_CS_LR/(tn_CS_LR+fp_CS_LR))*100,
(tn_CS_DT/(tn_CS_DT+fp_CS_DT))*100,
(tn_CS_SVC/(tn_CS_SVC+fp_CS_SVC))*100,
(tn_CS_XGB/(tn_CS_XGB+fp_CS_XGB))*100,
(tn_CS_LGBM/(tn_CS_LGBM+fp_CS_LGBM))*100],
'Predicted FP%': [(fp_CS_LR/(tn_CS_LR+fp_CS_LR))*100,
(fp_CS_DT/(tn_CS_DT+fp_CS_DT))*100,
(fp_CS_SVC/(tn_CS_SVC+fp_CS_SVC))*100,
(fp_CS_XGB/(tn_CS_XGB+fp_CS_XGB))*100,
(fp_CS_LGBM/(tn_CS_LGBM+fp_CS_LGBM))*100],
'Predicted FN%': [(fn_CS_LR/(tp_CS_LR+fn_CS_LR))*100,
(fn_CS_DT/(tp_CS_DT+fn_CS_DT))*100,
(fn_CS_SVC/(tp_CS_SVC+fn_CS_SVC))*100,
(fn_CS_XGB/(tp_CS_XGB+fn_CS_XGB))*100,
(fn_CS_LGBM/(tp_CS_LGBM+fn_CS_LGBM))*100]
}

```

```

results_df = pd.DataFrame(results)
print("\nComparison of Techniques:")
display(results_df)

```

```
# In[124]:
```

```
import pandas as pd
```

```
# First DataFrame (Original)
```

```

results_o = {
    'Imbalance Technique': ['Original Imbalanced Data'] * 6,
    'Model': ['Logistic Regression', 'Decision Tree', 'Random Forest', 'SVC', 'XGBoost',
'LightGBM'],
    'F1-Score': [f1_o_LR, f1_o_DT, f1_o_RF, f1_o_SVC, f1_o_XGB, f1_o_LGBM],
    'Precision': [precision_o_LR, precision_o_DT, precision_o_RF,
precision_o_SVC, precision_o_XGB, precision_o_LGBM],
    'recall': [recall_o_LR, recall_o_DT, recall_o_RF, recall_o_SVC, recall_o_XGB,
recall_o_LGBM],
    'ROC AUC': [roc_auc_o_LR, roc_auc_o_DT, roc_auc_o_RF, roc_auc_o_SVC,
roc_auc_o_XGB, roc_auc_o_LGBM],
    'Predicted TP%': [(tp_o_LR/(tp_o_LR+fn_o_LR))*100, (tp_o_DT/(tp_o_DT+fn_o_DT))*100,
(tp_o_RF/(tp_o_RF+fn_o_RF))*100, (tp_o_SVC/(tp_o_SVC+fn_o_SVC))*100,
(tp_o_XGB/(tp_o_XGB+fn_o_XGB))*100, (tp_o_LGBM/(tp_o_LGBM+fn_o_LGBM))*100],
    'Predicted TN%': [(tn_o_LR/(tn_o_LR+fp_o_LR))*100,
(tn_o_DT/(tn_o_DT+fp_o_DT))*100, (tn_o_RF/(tn_o_RF+fp_o_RF))*100,
(tn_o_SVC/(tn_o_SVC+fp_o_SVC))*100, (tn_o_XGB/(tn_o_XGB+fp_o_XGB))*100,
(tn_o_LGBM/(tn_o_LGBM+fp_o_LGBM))*100],
    'Predicted FP%': [(fp_o_LR/(tn_o_LR+fp_o_LR))*100, (fp_o_DT/(tn_o_DT+fp_o_DT))*100,
(fp_o_RF/(tn_o_RF+fp_o_RF))*100, (fp_o_SVC/(tn_o_SVC+fp_o_SVC))*100,
(fp_o_XGB/(tn_o_XGB+fp_o_XGB))*100, (fp_o_LGBM/(tn_o_LGBM+fp_o_LGBM))*100],

```



```

        'Predicted          FN%':          [(fn_o_LR/(tp_o_LR+fn_o_LR))*100,
(fn_o_DT/(tp_o_DT+fn_o_DT))*100,          (fn_o_RF/(tp_o_RF+fn_o_RF))*100,
(fn_o_SVC/(tp_o_SVC+fn_o_SVC))*100,      (fn_o_XGB/(tp_o_XGB+fn_o_XGB))*100,
(fn_o_LGBM/(tp_o_LGBM+fn_o_LGBM))*100]
    }
    df_o = pd.DataFrame(results_o)

```

Second DataFrame (Random Undersampling)

```

results_rus = {
    'Imbalance Technique': ['Random Undersampling'] * 6,
    'Model': ['Logistic Regression', 'Decision Tree', 'Random Forest', 'SVC', 'XGBoost',
'LightGBM'],
    'F1-Score': [f1_rus_LR, f1_rus_DT, f1_rus_RF, f1_rus_SVC, f1_rus_XGB, f1_rus_LGBM],
    'Precision':
[precision_o_LR,precision_o_DT,precision_o_RF,precision_o_SVC,precision_o_XGB,precis
ion_o_LGBM],
    'recall' : [recall_rus_LR, recall_rus_DT, recall_rus_RF, recall_rus_SVC, recall_o_XGB,
recall_rus_LGBM],
    'ROC AUC': [roc_auc_rus_LR, roc_auc_rus_DT, roc_auc_rus_RF, roc_auc_rus_SVC,
roc_auc_rus_XGB, roc_auc_rus_LGBM],
    'Predicted          TP%':          [(tp_rus_LR/(tp_rus_LR+fn_rus_LR))*100,
(tp_rus_DT/(tp_rus_DT+fn_rus_DT))*100,          (tp_rus_RF/(tp_rus_RF+fn_rus_RF))*100,
(tp_rus_SVC/(tp_rus_SVC+fn_rus_SVC))*100,
(tp_rus_XGB/(tp_rus_XGB+fn_rus_XGB))*100,
(tp_rus_LGBM/(tp_rus_LGBM+fn_rus_LGBM))*100],
    'Predicted          TN%':          [(tn_rus_LR/(tn_rus_LR+fp_rus_LR))*100,
(tn_rus_DT/(tn_rus_DT+fp_rus_DT))*100,          (tn_rus_RF/(tn_rus_RF+fp_rus_RF))*100,
(tn_rus_SVC/(tn_rus_SVC+fp_rus_SVC))*100,
(tn_rus_XGB/(tn_rus_XGB+fp_rus_XGB))*100,
(tn_rus_LGBM/(tn_rus_LGBM+fp_rus_LGBM))*100],
    'Predicted          FP%':          [(fp_rus_LR/(tn_rus_LR+fp_rus_LR))*100,
(fp_rus_DT/(tn_rus_DT+fp_rus_DT))*100,          (fp_rus_RF/(tn_rus_RF+fp_rus_RF))*100,
(fp_rus_SVC/(tn_rus_SVC+fp_rus_SVC))*100,
(fp_rus_XGB/(tn_rus_XGB+fp_rus_XGB))*100,
(fp_rus_LGBM/(tn_rus_LGBM+fp_rus_LGBM))*100],
    'Predicted          FN%':          [(fn_rus_LR/(tp_rus_LR+fn_rus_LR))*100,
(fn_rus_DT/(tp_rus_DT+fn_rus_DT))*100,          (fn_rus_RF/(tp_rus_RF+fn_rus_RF))*100,
(fn_rus_SVC/(tp_rus_SVC+fn_rus_SVC))*100,
(fn_rus_XGB/(tp_rus_XGB+fn_rus_XGB))*100,
(fn_rus_LGBM/(tp_rus_LGBM+fn_rus_LGBM))*100]
}
df_rus = pd.DataFrame(results_rus)

```

Third DataFrame (SMOTE)

```

results_smote = {
    'Imbalance Technique': ['SMOTE'] * 6,
    'Model': ['Logistic Regression', 'Decision Tree', 'Random Forest', 'SVC', 'XGBoost',
'LightGBM'],
    'F1-Score': [f1_SMOTE_LR, f1_SMOTE_DT, f1_SMOTE_RF, f1_SMOTE_SVC,
f1_SMOTE_XGB, f1_SMOTE_LGBM],

    'Precision': [precision_SMOTE_LR,precision_SMOTE_DT,precision_SMOTE_RF,precision_
SMOTE_SVC,precision_SMOTE_XGB,precision_SMOTE_LGBM],
    'recall' : [recall_SMOTE_LR, recall_SMOTE_DT, recall_SMOTE_RF, recall_SMOTE_SVC,
recall_SMOTE_XGB, recall_SMOTE_LGBM],

```

```

'ROC AUC': [roc_auc_SMOTE_LR, roc_auc_SMOTE_DT, roc_auc_SMOTE_RF,
roc_auc_SMOTE_SVC, roc_auc_SMOTE_XGB, roc_auc_SMOTE_LGBM],
'Predicted TP%': [(tp_SMOTE_LR/(tp_SMOTE_LR+fn_SMOTE_LR))*100,
(tp_SMOTE_DT/(tp_SMOTE_DT+fn_SMOTE_DT))*100,
(tp_SMOTE_RF/(tp_SMOTE_RF+fn_SMOTE_RF))*100,
(tp_SMOTE_SVC/(tp_SMOTE_SVC+fn_SMOTE_SVC))*100,
(tp_SMOTE_XGB/(tp_SMOTE_XGB+fn_SMOTE_XGB))*100,
(tp_SMOTE_LGBM/(tp_SMOTE_LGBM+fn_SMOTE_LGBM))*100],
'Predicted TN%': [(tn_SMOTE_LR/(tn_SMOTE_LR+fp_SMOTE_LR))*100,
(tn_SMOTE_DT/(tn_SMOTE_DT+fp_SMOTE_DT))*100,
(tn_SMOTE_RF/(tn_SMOTE_RF+fp_SMOTE_RF))*100,
(tn_SMOTE_SVC/(tn_SMOTE_SVC+fp_SMOTE_SVC))*100,
(tn_SMOTE_XGB/(tn_SMOTE_XGB+fp_SMOTE_XGB))*100,
(tn_SMOTE_LGBM/(tn_SMOTE_LGBM+fp_SMOTE_LGBM))*100],
'Predicted FP%': [(fp_SMOTE_LR/(tn_SMOTE_LR+fp_SMOTE_LR))*100,
(fp_SMOTE_DT/(tn_SMOTE_DT+fp_SMOTE_DT))*100,
(fp_SMOTE_RF/(tn_SMOTE_RF+fp_SMOTE_RF))*100,
(fp_SMOTE_SVC/(tn_SMOTE_SVC+fp_SMOTE_SVC))*100,
(fp_SMOTE_XGB/(tn_SMOTE_XGB+fp_SMOTE_XGB))*100,
(fp_SMOTE_LGBM/(tn_SMOTE_LGBM+fp_SMOTE_LGBM))*100],
'Predicted FN%': [(fn_SMOTE_LR/(tp_SMOTE_LR+fn_SMOTE_LR))*100,
(fn_SMOTE_DT/(tp_SMOTE_DT+fn_SMOTE_DT))*100,
(fn_SMOTE_RF/(tp_SMOTE_RF+fn_SMOTE_RF))*100,
(fn_SMOTE_SVC/(tp_SMOTE_SVC+fn_SMOTE_SVC))*100,
(fn_SMOTE_XGB/(tp_SMOTE_XGB+fn_SMOTE_XGB))*100,
(fn_SMOTE_LGBM/(tp_SMOTE_LGBM+fn_SMOTE_LGBM))*100]
}
df_smote = pd.DataFrame(results_smote)

```

Fourth DataFrame (Cost-Sensitive)

```

results_cs = {
'Imbalance Technique': ['Cost Sensitive Analysis'] * 6,
'Model': ['Logistic Regression', 'Decision Tree', 'Random Forest', 'SVC', 'XGBoost',
'LightGBM'],
'F1-Score': [f1_CS_LR, f1_CS_DT, f1_CS_RF, f1_CS_SVC, f1_CS_XGB, f1_CS_LGBM],

'Precision': [precision_CS_LR, precision_CS_DT, precision_CS_RF, precision_CS_SVC, precision_CS_XGB, precision_CS_LGBM],
'recall': [recall_CS_LR, recall_CS_DT, recall_CS_RF, recall_CS_SVC, recall_CS_XGB, recall_CS_LGBM],
'ROC AUC': [roc_auc_CS_LR, roc_auc_CS_DT, roc_auc_CS_RF, roc_auc_CS_SVC, roc_auc_CS_XGB, roc_auc_CS_LGBM],
'Predicted TP%': [(tp_CS_LR/(tp_CS_LR+fn_CS_LR))*100,
(tp_CS_DT/(tp_CS_DT+fn_CS_DT))*100,
(tp_CS_SVC/(tp_CS_SVC+fn_CS_SVC))*100,
(tp_CS_XGB/(tp_CS_XGB+fn_CS_XGB))*100,
(tp_CS_LGBM/(tp_CS_LGBM+fn_CS_LGBM))*100],
'Predicted TN%': [(tn_CS_LR/(tn_CS_LR+fp_CS_LR))*100,
(tn_CS_DT/(tn_CS_DT+fp_CS_DT))*100,
(tn_CS_SVC/(tn_CS_SVC+fp_CS_SVC))*100,
(tn_CS_XGB/(tn_CS_XGB+fp_CS_XGB))*100,
(tn_CS_LGBM/(tn_CS_LGBM+fp_CS_LGBM))*100],
'Predicted FP%': [(fp_CS_LR/(tn_CS_LR+fp_CS_LR))*100,
(fp_CS_DT/(tn_CS_DT+fp_CS_DT))*100,
(fp_CS_SVC/(tn_CS_SVC+fp_CS_SVC))*100,

```

```

(fp_CS_XGB/(tn_CS_XGB+fp_CS_XGB))*100,
(fp_CS_LGBM/(tn_CS_LGBM+fp_CS_LGBM))*100],
    'Predicted FN%': [(fn_CS_LR/(tp_CS_LR+fn_CS_LR))*100,
(fp_CS_DT/(tp_CS_DT+fn_CS_DT))*100, (fn_CS_RF/(tp_CS_RF+fn_CS_RF))*100,
(fn_CS_SVC/(tp_CS_SVC+fn_CS_SVC))*100,
(fn_CS_XGB/(tp_CS_XGB+fn_CS_XGB))*100,
(fn_CS_LGBM/(tp_CS_LGBM+fn_CS_LGBM))*100]
}
df_cs = pd.DataFrame(results_cs)

# Fifth DataFrame (Threshold Optimization)
results_th = {
    'Imbalance Technique': ['Threshold Optimization'] * 6,
    'Model': ['Logistic Regression', 'Decision Tree', 'Random Forest', 'SVC', 'XGBoost',
'LightGBM'],
    'F1-Score': [f1_Th_LR, f1_Th_DT, f1_Th_RF, f1_Th_SVC, f1_Th_XGB, f1_Th_LGBM],

    'Precision': [precision_Th_LR, precision_Th_DT, precision_Th_RF, precision_Th_SVC, precision_Th_XGB, precision_Th_LGBM],
    'recall' : [recall_Th_LR, recall_Th_DT, recall_Th_RF, recall_Th_SVC, recall_Th_XGB, recall_Th_LGBM],
    'ROC AUC': [roc_auc_Th_LR, roc_auc_Th_DT, roc_auc_Th_RF, roc_auc_Th_SVC, roc_auc_Th_XGB, roc_auc_Th_LGBM],
    'Predicted TP%': [(tp_Th_LR/(tp_Th_LR+fn_Th_LR))*100,
(tp_Th_DT/(tp_Th_DT+fn_Th_DT))*100, (tp_Th_RF/(tp_Th_RF+fn_Th_RF))*100,
(tp_Th_SVC/(tp_Th_SVC+fn_Th_SVC))*100, (tp_Th_XGB/(tp_Th_XGB+fn_Th_XGB))*100,
(tp_Th_LGBM/(tp_Th_LGBM+fn_Th_LGBM))*100],
    'Predicted TN%': [(tn_Th_LR/(tn_Th_LR+fp_Th_LR))*100,
(tn_Th_DT/(tn_Th_DT+fp_Th_DT))*100, (tn_Th_RF/(tn_Th_RF+fp_Th_RF))*100,
(tn_Th_SVC/(tn_Th_SVC+fp_Th_SVC))*100, (tn_Th_XGB/(tn_Th_XGB+fp_Th_XGB))*100,
(tn_Th_LGBM/(tn_Th_LGBM+fp_Th_LGBM))*100],
    'Predicted FP%': [(fp_Th_LR/(tn_Th_LR+fp_Th_LR))*100,
(fp_Th_DT/(tn_Th_DT+fp_Th_DT))*100, (fp_Th_RF/(tn_Th_RF+fp_Th_RF))*100,
(fp_Th_SVC/(tn_Th_SVC+fp_Th_SVC))*100, (fp_Th_XGB/(tn_Th_XGB+fp_Th_XGB))*100,
(fp_Th_LGBM/(tn_Th_LGBM+fp_Th_LGBM))*100],
    'Predicted FN%': [(fn_Th_LR/(tp_Th_LR+fn_Th_LR))*100,
(fn_Th_DT/(tp_Th_DT+fn_Th_DT))*100, (fn_Th_RF/(tp_Th_RF+fn_Th_RF))*100,
(fn_Th_SVC/(tp_Th_SVC+fn_Th_SVC))*100, (fn_Th_XGB/(tp_Th_XGB+fn_Th_XGB))*100,
(fn_Th_LGBM/(tp_Th_LGBM+fn_Th_LGBM))*100]
}
df_th = pd.DataFrame(results_th)

# Combine all
combined_df = pd.concat([df_o, df_rus, df_smote, df_th, df_cs], ignore_index=True)

# Add S.No.
combined_df.insert(0, 'S.No.', range(1, len(combined_df) + 1))

# Sorted df
sorted_df = combined_df.sort_values(by="F1-Score",
ascending=False).reset_index(drop=True)

# Display final table
print("\nCombined Comparison of Techniques:")

```

```
display(combined_df)
```

```
# In[ ]:
```

```
# # ROC_AUC Curve
```

```
# In[127]:
```

```
import matplotlib.pyplot as plt
import numpy as np
from sklearn.metrics import roc_curve, auc
import seaborn as sns
```

```
# Set up the plot style
plt.style.use('default')
sns.set_palette("husl")
plt.rcParams['figure.figsize'] = [10, 8]
plt.rcParams['font.size'] = 12
```

```
# Define colors for different models
```

```
model_colors = {
    'Logistic Regression': '#1f77b4',    # blue
    'Decision Tree': '#ff7f0e',          # orange
    'Random Forest': '#2ca02c',          # green
    'SVC': '#d62728',                    # red
    'XGBoost': '#9467bd',                 # purple
    'LightGBM': '#8c564b'                 # brown
}
```

```
# Define line styles for better distinction
```

```
model_line_styles = {
    'Logistic Regression': '-',
    'Decision Tree': '--',
    'Random Forest': '-.',
    'SVC': ':',
    'XGBoost': (0, (3, 1, 1, 1)),
    'LightGBM': (0, (5, 1))
}
```

```
# Define marker styles (optional, for extra distinction)
```

```
model_markers = {
    'Logistic Regression': 'o',
    'Decision Tree': 's',
    'Random Forest': '^',
    'SVC': 'D',
    'XGBoost': 'v',
    'LightGBM': '<'
}
```

```

# Get unique techniques
techniques = combined_df['Imbalance Technique'].unique()

# Create a figure with 5 subplots (one for each technique)
fig, axes = plt.subplots(2, 3, figsize=(20, 15))
axes = axes.flatten()

# Remove the extra subplot if we have 5 techniques
if len(techniques) == 5:
    fig.delaxes(axes[5])

# Plot ROC curves for each technique
for i, technique in enumerate(techniques):
    if i >= len(axes): # Safety check
        break

    ax = axes[i]
    technique_data = combined_df[combined_df['Imbalance Technique'] == technique]

    # Plot diagonal reference line
    ax.plot([0, 1], [0, 1], 'k--', lw=2, alpha=0.6, label='Random Classifier (AUC = 0.500)')

    # Plot each model's ROC curve for this technique
    for _, row in technique_data.iterrows():
        model_name = row['Model']
        roc_auc = row['ROC AUC']

        # REPLACE THIS WITH YOUR ACTUAL PROBABILITIES
        # Example: y_probs = y_probs_o_LR for Original Logistic Regression
        # y_test = your actual test labels

        # This is placeholder - replace with your actual probability arrays
        if f'y_probs_{technique.split()[0].lower()}_{model_name.split()[0].lower()}' in globals():
            y_probs = globals()[f'y_probs_{technique.split()[0].lower()}_{model_name.split()[0].lower()}']
            fpr, tpr, _ = roc_curve(y_test, y_probs)
            roc_auc_calculated = auc(fpr, tpr)
        else:
            # Create dummy ROC curve (replace with your actual data)
            fpr = np.linspace(0, 1, 100)
            base_tpr = np.sqrt(fpr) # Reasonable ROC shape approximation
            # Scale to match the reported AUC
            scaling_factor = roc_auc / auc(fpr, base_tpr)
            tpr = base_tpr * scaling_factor
            tpr = np.clip(tpr, 0, 1) # Ensure valid values
            roc_auc_calculated = auc(fpr, tpr)

    # Plot the ROC curve
    ax.plot(fpr, tpr,
            color=model_colors[model_name],
            linestyle=model_line_styles[model_name],
            #marker=model_markers[model_name],
            #markersize=4,
            lw=2.5,
            alpha=0.8,

```

```

        label=f'{model_name} (AUC = {roc_auc:.3f})')

# Customize the subplot
ax.set_xlim([0.0, 1.0])
ax.set_ylim([0.0, 1.05])
ax.set_xlabel('False Positive Rate', fontsize=11)
ax.set_ylabel('True Positive Rate', fontsize=11)
ax.set_title(f'ROC Curves - {technique}', fontsize=14, fontweight='bold')
ax.legend(loc="lower right", fontsize=9)
ax.grid(True, alpha=0.3, linestyle='--')
ax.set_aspect('equal', adjustable='box')

plt.tight_layout()
plt.suptitle('ROC Curves by Imbalance Handling Technique', fontsize=16, fontweight='bold',
y=1.02)
plt.show()

# Create a combined comparison plot
plt.figure(figsize=(14, 10))

# Plot diagonal reference line
plt.plot([0, 1], [0, 1], 'k--', lw=2, alpha=0.6, label='Random Classifier')

# Plot all techniques and models
for technique in techniques:
    technique_data = combined_df[combined_df['Imbalance Technique'] == technique]

    for _, row in technique_data.iterrows():
        model_name = row['Model']
        roc_auc = row['ROC AUC']

        # REPLACE WITH YOUR ACTUAL PROBABILITIES
        if f'y_probs_{technique.split()[0].lower()}_{model_name.split()[0].lower()}' in globals():
            y_probs = globals()[f'y_probs_{technique.split()[0].lower()}_{model_name.split()[0].lower()}']
            fpr, tpr, _ = roc_curve(y_test, y_probs)
        else:
            # Dummy curve (replace with actual data)
            fpr = np.linspace(0, 1, 100)
            base_tpr = np.sqrt(fpr)
            scaling_factor = roc_auc / auc(fpr, base_tpr)
            tpr = base_tpr * scaling_factor
            tpr = np.clip(tpr, 0, 1)

        plt.plot(fpr, tpr,
            color=model_colors[model_name],
            linestyle=model_line_styles[model_name],
            lw=2,
            alpha=0.7,
            label=f'{technique} - {model_name} (AUC = {roc_auc:.3f})')

plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate', fontsize=12)
plt.ylabel('True Positive Rate', fontsize=12)

```

```
plt.title('Combined ROC Curves - All Techniques and Models', fontsize=16, fontweight='bold')
plt.legend(loc="lower right", fontsize=8, ncol=2)
plt.grid(True, alpha=0.3, linestyle='--')
plt.tight_layout()
plt.show()
```

```
# Create AUC comparison bar chart
plt.figure(figsize=(16, 8))
```

```
# Prepare data for bar chart
technique_model_auc = []
for technique in techniques:
    technique_data = combined_df[combined_df['Imbalance Technique'] == technique]
    for _, row in technique_data.iterrows():
        technique_model_auc.append({
            'Technique': technique,
            'Model': row['Model'],
            'AUC': row['ROC AUC']
        })
```

```
auc_df = pd.DataFrame(technique_model_auc)
```

```
# Create grouped bar chart
x = np.arange(len(techniques))
width = 0.13 # Width of each bar
```

```
fig, ax = plt.subplots(figsize=(16, 8))
models = auc_df['Model'].unique()
```

```
for i, model in enumerate(models):
    model_aucs = auc_df[auc_df['Model'] == model]['AUC'].values
    ax.bar(x + i * width, model_aucs, width, label=model, color=model_colors[model],
           alpha=0.8)
```

```
ax.set_xlabel('Imbalance Handling Technique', fontsize=12)
ax.set_ylabel('ROC AUC Score', fontsize=12)
ax.set_title('ROC AUC Scores by Technique and Model', fontsize=16, fontweight='bold')
ax.set_xticks(x + width * (len(models) - 1) / 2)
ax.set_xticklabels(techniques, rotation=45, ha='right')
ax.legend(loc='upper left', bbox_to_anchor=(1, 1))
ax.grid(True, alpha=0.3, axis='y')
```

```
# Add value labels on bars
for i, technique in enumerate(techniques):
    technique_data = auc_df[auc_df['Technique'] == technique]
    for j, (_, row) in enumerate(technique_data.iterrows()):
        ax.text(i + j * width, row['AUC'] + 0.005, f'{row["AUC"]:.3f}',
                ha='center', va='bottom', fontsize=8, rotation=0)
```

```
plt.tight_layout()
plt.show()
```

```
# In[128]:
```

```

from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
import matplotlib.pyplot as plt

# Example: use Logistic Regression predictions
y_true = y_test
y_pred = y_lr_o # replace with predictions of your chosen model

# Compute confusion matrix
cm = confusion_matrix(y_true, y_pred, labels=[1,0])
# Note: order [1,0] ensures row=actual default/non-default

# Display confusion matrix (counts)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=["Default (1)", "Non-Default (0)"])
disp.plot(cmap=plt.cm.Blues)
plt.title("Confusion Matrix - Logistic Regression (Original Data)")
plt.show()

# ---- Calculate percentages ----
TN, FP, FN, TP = cm.ravel()

total = cm.sum()
TP_percent = (TP / total) * 100
TN_percent = (TN / total) * 100
FP_percent = (FP / total) * 100
FN_percent = (FN / total) * 100

print(f"TP%: {TP_percent:.2f}, TN%: {TN_percent:.2f}, FP%: {FP_percent:.2f}, FN%: {FN_percent:.2f}")

# # SHAP

# In[129]:

import shap

# LGBM with cost Sensitive Analysis

# In[130]:

# Get the selected feature names from your RFECV
LGBMselected_feature_names =
X_train_preprocessed_df.columns[LGBMrfecv.support_].tolist()

# Convert your selected training data to DataFrame
LGBM_X_train_selected_df = pd.DataFrame(LGBM_X_train_selected,
columns=LGBMselected_feature_names)

# Also convert your test data for consistency

```



```
LGBM_X_test_selected_df = pd.DataFrame(LGBM_X_test_selected,
columns=LGBMselected_feature_names)
```

```
# Common practice: use a sample of the training data as the background
explainer = shap.TreeExplainer(rf_weighted_LGBM)
```

```
# In[131]:
```

```
shap_values = explainer.shap_values(LGBM_X_test_selected)
```

```
# Global feature importance
shap_importance = np.abs(shap_values).mean(0)
```

```
# In[132]:
```

```
shap.summary_plot(shap_values, LGBM_X_test_selected_df, max_display=10,
plot_type='bar', title='LightGBM Feature Importance with Cost Sensitive Analysis',
alpha=1, feature_names=LGBMselected_feature_names, show=False)
```

```
# LightGBM with SMOTE
```

```
# In[133]:
```

```
# Common practice: use a sample of the training data as the background
explainer = shap.TreeExplainer(rf_smote_LGBM)
shap_values = explainer.shap_values(LGBM_X_test_selected)
```

```
# Global feature importance
shap_importance = np.abs(shap_values).mean(0)
```

```
shap.summary_plot(shap_values, LGBM_X_test_selected_df, max_display=10,
plot_type='bar', title='LightGBM Feature Importance with SMOTE',
alpha=1, feature_names=LGBMselected_feature_names, show=False)
```

```
# LGBM with Threshold Optimization
```

```
# In[134]:
```

```
# Common practice: use a sample of the training data as the background
explainer = shap.TreeExplainer(rf)
shap_values = explainer.shap_values(LGBM_X_test_selected)
```

```
# Global feature importance
shap_importance = np.abs(shap_values).mean(0)
```

```
shap.summary_plot(shap_values, LGBM_X_test_selected_df, max_display=10,
plot_type='bar', title='LightGBM Feature Importance with Threshold Optimization',
```

```

        alpha=1,feature_names=LGBMselected_feature_names, show=False)

# LGBM on Origina Imbalanced dataset

# In[135]:

# Common practice: use a sample of the training data as the background
explainer = shap.TreeExplainer(lgbm_o)
shap_values = explainer.shap_values(LGBM_X_test_selected)

# Global feature importance
shap_importance = np.abs(shap_values).mean(0)

shap.summary_plot(shap_values, LGBM_X_test_selected_df, max_display=10,
plot_type='bar', title='LightGBM Feature Importance on original imbalanced data',
alpha=1,feature_names=LGBMselected_feature_names, show=False)

# XGBoost with Threshold

# In[136]:

# Get the selected feature names from your RFECV
XGBselected_feature_names =
X_train_preprocessed_df.columns[XGBrfecv.support_].tolist()

# Convert your selected training data to DataFrame
XGB_X_train_selected_df = pd.DataFrame(XGB_X_train_selected,
columns=XGBselected_feature_names)

# Also convert your test data for consistency
XGB_X_test_selected_df = pd.DataFrame(XGB_X_test_selected,
columns=XGBselected_feature_names)

# In[137]:

# Common practice: use a sample of the training data as the background
explainer = shap.TreeExplainer(xgb_th)
shap_values = explainer.shap_values(XGB_X_test_selected)

# Global feature importance
shap_importance = np.abs(shap_values).mean(0)

shap.summary_plot(shap_values, XGB_X_test_selected_df, max_display=10,
plot_type='bar', title='XGBoost Feature Importance with Threshold Optimization',
alpha=1,feature_names=XGBselected_feature_names, show=False)

# Cost sensitive XGBoost

```

In[138]:

Common practice: use a sample of the training data as the background

```
explainer = shap.TreeExplainer(rf_weighted_XGB)
shap_values = explainer.shap_values(XGB_X_test_selected)
```

Global feature importance

```
shap_importance = np.abs(shap_values).mean(0)
```

```
shap.summary_plot(shap_values, XGB_X_test_selected_df, max_display=10,
plot_type='bar', title='XGBoost Feature Importance with Cost Sensitivity',
alpha=1, feature_names=XGBselected_feature_names, show=False)
```

In[139]:

Common practice: use a sample of the training data as the background

```
explainer = shap.TreeExplainer(rf_smote_XGB)
shap_values = explainer.shap_values(XGB_X_test_selected)
```

Global feature importance

```
shap_importance = np.abs(shap_values).mean(0)
```

```
shap.summary_plot(shap_values, XGB_X_test_selected_df, max_display=10,
plot_type='bar', title='XGBoost Feature Importance with SMOTE',
alpha=1, feature_names=XGBselected_feature_names, show=False)
```

Cost Sensitive RF

In[140]:

Get the selected feature names from your RFECV

```
RFselected_feature_names = X_train_preprocessed_df.columns[RFrfecv.support_].tolist()
```

Convert your selected training data to DataFrame

```
RF_X_train_selected_df = pd.DataFrame(RF_X_train_selected,
columns=RFselected_feature_names)
```

Also convert your test data for consistency

```
RF_X_test_selected_df = pd.DataFrame(RF_X_test_selected,
columns=RFselected_feature_names)
```

In[141]:

Common practice: use a sample of the training data as the background

```
explainer = shap.Explainer(rf_smote_RF)
shap_values = explainer.shap_values(RF_X_test_selected)
```

Global feature importance

```
shap_importance = np.abs(shap_values).mean(0)
```

```
shap.summary_plot(shap_values, RF_X_test_selected_df, max_display=10, plot_type='bar',  
title='RF Feature Importance with Cost Sensitivity',  
alpha=1,feature_names=RFselected_feature_names, show=False)
```

```
# In[ ]:
```

```
# In[ ]:
```

```
# In[ ]:
```